Munich University of Applied Sciences

Hochschule Muenchen

Department of Computer Science & Mathematics

Fakultaet fuer Informatik und Mathematik

Quality-Driven Next-Best-View Planning for Target-Conditioned Egocentric 3D Reconstruction

Qualitätsgetriebene Next-Best-View-Planung für zielkonditionierte egozentrische 3D-Rekonstruktion

Master's Thesis / Masterarbeit

for obtaining the academic degree Master of Science (M.Sc.)

submitted to
Munich University of Applied Sciences
Department of Computer Science & Mathematics

Submitted by:	Jan Duchscherer
Email:	j.duchscherer@hm.edu
Program:	Informatik - Master
Specialization:	Visual Computing and Machine Learning
Matriculation Number:	00807819

First Examiner:	Prof. Dr. Markus Friedrich
Second Examiner:	Prof. Dr. David Spieler
Supervisor:	Prof. Dr. Markus Friedrich
Submission Date:	7. Jan 2027

Transparency in the Use of AI Tools

Generative-AI tools supported literature-note organization, consistency checks, and language revision. The author selected the research questions, verified sources and technical claims, implemented and evaluated the system, and remains responsible for the submitted work.

Acknowledgements

Acknowledgements are omitted from this version.

Abstract

This thesis studies target-conditioned, quality-driven Next-Best-View (NBV) planning for egocentric 3D reconstruction in Aria Synthetic Environments (ASE). It defines target-specific Relative Reconstruction Improvement (RRI) as an oracle signal, constructs finite candidate and replay contracts, and separates privileged label generation from the actor-visible inputs available to a learned finite-horizon value model. The evaluation is designed to measure oracle-lookahead headroom and the fraction recovered by a learned policy under matched oracle re-evaluation. The evidence available for this version establishes the evaluation contract and implementation readiness, but does not contain confirmatory held-out policy outcomes; it therefore supports no claim that a learned policy improves over the specified baselines. The main remaining limitation is the absence of a validated, population-level rollout bundle with paired endpoint estimates and uncertainty.

Zusammenfassung

Diese Arbeit untersucht zielkonditionierte, qualitätsgetriebene Planung der nächsten besten Ansicht für die egozentrische 3D-Rekonstruktion in ASE. Sie definiert die zielspezifische RRI als Orakelsignal, legt Verträge für endliche Kandidatenmengen und Replay-Daten fest und trennt die privilegierte Erzeugung von Trainingssignalen von den für ein gelerntes Modell mit endlichem Horizont sichtbaren Eingaben. Die Evaluation soll den Spielraum einer vorausschauenden Orakelstrategie und den durch eine gelernte Strategie erreichten Anteil unter identischer Orakel-Neubewertung messen. Die für diese Fassung verfügbare Evidenz belegt den Evaluationsvertrag und die Implementierungsbereitschaft, enthält jedoch keine bestätigenden Ergebnisse auf zurückgehaltenen Daten. Daher wird keine Überlegenheit einer gelernten Strategie gegenüber den festgelegten Baselines behauptet. Die wesentliche verbleibende Einschränkung ist das Fehlen eines validierten Rollout-Datensatzes auf Populationsebene mit gepaarten Endpunktschätzungen und Unsicherheitsangaben.

Contents

Glossary and Abbreviations

Dataset

ADT **– Aria Digital Twin**: Project Aria dataset with real-world captures and digital-twin scene annotations. ADT is adjacent to ARIA-NBV's sim-to-real context; the current experiments focus on ASE and EFM3D/ATEK exports, while ADT remains a relevant transfer surface.

AEO **– Aria Everyday Objects**: Small-scale real-world Project Aria object dataset used by EFM3D for egocentric 3D perception evaluation. AEO is relevant as a possible sim-to-real check for ARIA-NBV ideas that are first developed on ASE mesh-supervised snippets.

ASE **– Aria Synthetic Environments**: Large-scale synthetic indoor dataset with simulated Project Aria sensor characteristics, egocentric trajectories, and scene annotations. ARIA-NBV uses ASE snippets and the public mesh-supervised subset to generate oracle RRI labels from ground-truth meshes and semi-dense point clouds.

CPF **– Central Pupil Frame**: Coordinate frame placed at the midpoint between the left and right eye boxes of Project Aria glasses. CPF is used by Project Aria for gaze-related quantities and should remain distinct from rig, camera, world, and PyTorch3D frames in ARIA-NBV docs.

MPS **– Machine Perception Services**: Project Aria processing services that derive pose, mapping, gaze, hand, and related perception outputs from sensor recordings. ARIA-NBV mainly depends on MPS-style pose and semi-dense mapping products as the current reconstruction state for RRI computation.

MTD **– Motion Trajectory Data**: Device poses over time, usually represented as a sequence of 6-DoF transformations. MTD supplies the logged egocentric trajectory used to define snippet state, current reconstruction context, and candidate-view reference poses.

MFCD **– Multi-Frame Camera Data**: Synchronized camera streams from multiple Project Aria cameras over a temporal window. MFCD provides the egocentric image evidence consumed by EVL and aligned with pose, calibration, and semi-dense point streams.

***MSDPD* – Multi-Semi-Dense Point Data**: Semi-dense 3D point observations generated by SLAM-style processing across a snippet or trajectory window. MSDPD is the sparse observed geometry that ARIA-NBV fuses with candidate point clouds and compares against ground-truth meshes for RRI labels.

***Project Aria* – Project Aria**: Meta's egocentric research-device and tooling ecosystem for calibrated, time-aligned multimodal sensing. ARIA-NBV treats Project Aria and its MPS-style products as the actor-visible sensing contract, while ASE meshes remain offline supervision and evaluation assets.

***SSL* – SceneScript Language**: Structured language representation for indoor scene layout using primitives such as walls, doors, windows, and objects. SceneScript is relevant to ARIA-NBV as a possible semantic/global planning layer and as one source of ASE scene-structure context.

***snippet* – Snippet**: Short synchronized temporal window of Aria sensor data used as one EVL/VIN input sample. A snippet typically contains RGB or grayscale streams, poses, calibration, semi-dense points, and scene metadata that EVL lifts into a voxel grid.

***VRS* – Virtual Reality Standard**: File format used by Project Aria tooling to store multi-modal sensor recordings efficiently. VRS is part of the upstream Project Aria ecosystem, while ARIA-NBV primarily consumes ASE/ATEK-derived tensorized snippets for current experiments.

Entity

***target* – Target of Interest**: Selected entity, object crop, point, region, or surface-deficit hypothesis whose reconstruction quality should be improved. Target-conditioned ARIA-NBV variants use this target as an explicit input to candidate scoring and planning instead of optimizing only scene-level RRI.

Geometry

***DoF* – Degrees of Freedom**: Number of independent pose parameters available to a camera, object, or action representation. ARIA-NBV distinguishes full 6-DoF pose estimation from reduced 5-DoF candidate or action spaces used for practical view planning.

***frustum* – Frustum**: Truncated pyramidal camera-visible volume bounded by near and far clipping planes plus lateral field-of-view planes. Candidate-view frusta define which scene surfaces can project into a camera image and are therefore central to visibility, rendering, and RRI diagnostics.

***LUF* – Left-Up-Forward**: Camera coordinate convention whose x axis points left, y axis points up, and z axis points forward. LUF is one of the coordinate-frame conventions that must stay explicit when moving between Project Aria cameras, PyTorch3D cameras, and ARIA-NBV candidate poses.

***OBB* – Oriented Bounding Box**: 3D bounding box with arbitrary orientation, used to represent object extent more tightly than an axis-aligned box. OBBs are a natural target encoding for entity-aware ARIA-NBV because they provide center, extent, orientation, semantic class, and confidence signals.

***6DoF* – Six Degrees of Freedom**: Pose parameterization with three translational and three rotational degrees of freedom. Project Aria poses, candidate cameras, and world-to-device transforms are usually represented as 6-DoF rigid-body transforms.

Metrics

***AUC* – Area Under Curve**: Aggregate score computed by integrating a metric curve over an acquisition, threshold, or ranking axis. AUC can summarize how quickly reconstruction quality, coverage, or ranking quality improves as views are acquired or candidate thresholds change.

***CD* – Chamfer Distance**: Historical bidirectional distance family used to compare reconstructed points against reference geometry. Thesis-facing ARIA-

NBV notation uses point-mesh error D with directional components $D_{P \to M}$ and $D_{M \to P}$; older seminar material may still call this CD.

***CR* – Coverage Ratio**: Fraction of a target surface, scene, or region treated as observed under a chosen visibility or distance threshold. Coverage ratio is a useful diagnostic baseline for NBV, but ARIA-NBV treats reconstruction-quality improvement through RRI as the preferred optimization target.

***GT* – Ground Truth**: Reference data or annotations treated as the trusted target for training, validation, or evaluation. In ARIA-NBV, ground truth usually refers to ASE meshes, object annotations, poses, or labels used to compute oracle supervision and diagnostic metrics.

***RRI* – Relative Reconstruction Improvement**: Metric quantifying the relative reconstruction-quality improvement obtained by adding a candidate observation to the current reconstruction. ARIA-NBV computes RRI by comparing reconstruction error before and after fusing candidate-view geometry, usually against an ASE ground-truth mesh for oracle supervision and evaluation.

***reward* – Target-RRI Reward**: Quality-only immediate reward for thesis-core target-aware rollouts. It is the target-specific point-mesh error reduction normalized by the rollout-root target error; scalar motion, rule, diversity, and validity penalties are later ablations.

***target RRI* – Target-Specific RRI**: RRI computed only on the ground-truth and reconstructed geometry associated with a selected target of interest. Target-specific RRI lets the thesis compare views by how much they improve a selected object or region, even when scene-level RRI would prefer large background surfaces.

Model

***EFM3D* – Egocentric Foundation Model 3D**: Egocentric 3D foundation-model stack used as the frozen spatial backbone for ARIA-NBV candidate scoring. The current project uses EFM3D and its EVL architecture to expose voxel occupancy, centerness, semantic, and OBB evidence for VIN-style RRI prediction.

***EVL* – Egocentric Voxel Lifting**: EFM3D architecture that lifts synchronized egocentric observations into a gravity-aligned 3D voxel feature volume. ARIA-NBV uses EVL head outputs, pre-head features, and OBB predictions as local actor-visible evidence and target support, not as the complete frozen scene representation.

***Q_H* – Finite-Horizon Q Function**: Target-conditioned finite-candidate value function trained from ASE oracle rollout traces. The first thesis implementation uses candidate-to-state query attention and emits one masked continuous-return value per valid candidate-table action, using DQN-style replayed Bellman targets, Double-DQN selector/evaluator backups, and offline support constraints.

***target-conditioned scorer* – Target-Conditioned Scorer**: VIN-style candidate scorer that receives scene state, a candidate view, and an encoding of the target of interest. The scorer predicts target-specific utility so view ranking can prioritize a selected entity or region instead of only optimizing scene-level RRI.

***VIN* – View Introspection Network**: Learned candidate-view scorer that predicts RRI or an ordinal RRI-derived utility without capturing the candidate view. ARIA-NBV adapts VIN-NBV by placing a lightweight RRI prediction head on top of frozen EVL features and candidate-pose evidence from ASE snippets.

Planning

***cost* – Acquisition Cost**: Budget consumed to acquire observations, measured by view count, path length, elapsed time, invalid-action rate, or a weighted combination. The thesis should first report cumulative root-normalized target gain, diagnostic target RRI, and acquisition cost separately, then use scalarized objectives only when the tradeoff is explicit.

***candidate* – Candidate View**: Proposed camera pose whose expected reconstruction utility is evaluated before selecting the next observation. ARIA-NBV samples candidate views around a reference pose or target, renders candidate depths from the mesh for oracle labels, and scores candidates with RRI or learned VIN predictions.

***transition* – Counterfactual Transition**: Deterministic or replayable rollout update that renders or retrieves only the selected candidate observation, backprojects it to candidate geometry, and fuses it into the accumulated reconstruction proxy.

***action set* – Finite Candidate Action Set**: Valid discrete action set for ARIA-NBV planning, obtained by masking a sampled candidate view table. The action is the candidate-table index, not a continuous pose; the selected pose is $q_t = q_{t,a_t}$. Invalid candidates are constraints, not low-quality actions.

***return* – Finite-Horizon Return**: Symbolic H-step return used by ARIA-NBV rollout evaluation and Q_H targets. The formula keeps γ explicit while the thesis-core comparison reports cumulative root-normalized target gain under equal acquisition budget.

***5DoF* – Five Degrees of Freedom**: Reduced camera-action parameterization commonly used when roll is fixed or otherwise constrained. ARIA-NBV uses 5-DoF language for candidate and planning abstractions where position and viewing direction matter while roll is not an independently optimized control dimension.

***CF+ state* – Geometry-Rich Counterfactual State**: Ablation actor state that extends the minimal counterfactual state with selected prior synthetic observations only: rendered depth, valid mask, backprojected point cloud, and derived local normals or support summaries.

***historic state* – Historic Snippet State**: Raw actor-visible state available along the logged Project Aria/ASE trajectory before counterfactual rollout synthesis. It includes captured camera streams, calibration, timestamps, trajectory/gravity, semidense points with support or uncertainty, EVL/EFM evidence, and detected or predicted OBBs; GT mesh and GT OBBs are excluded as actor inputs.

CF0 state **– Minimal Counterfactual Actor State**: Main thesis-core actor-visible counterfactual rollout state. It keeps the fused point proxy $\mathcal{P}_t$ as broad scene state, optional lifted image-foundation features attached to points, local root EVL evidence for target support, selected-view history, target descriptor, budget, candidate table, validity masks and reason codes, and current candidate-query features, without exposing all-candidate GT renders.

NBV **– Next-Best-View**: Problem of selecting the next sensor viewpoint to improve an active reconstruction or inspection objective under a limited acquisition budget. In ARIA-NBV, the preferred objective is reconstruction quality improvement rather than coverage alone, with candidate views ranked by oracle or learned RRI scores.

oracle state **– Oracle Rollout State**: Privileged state used only for ASE mesh-supervised label generation, upper-bound planning, and evaluation. It includes GT mesh and target crops, all-candidate renders, points, normals, mesh-face visibility, RRI and point-mesh error metrics, oracle scores, and GT OBBs.

offline state **– Persisted Offline Sample State**: Compact VIN offline-store state persisted for training and diagnostics. It contains the VinSnippetView payload, candidate table and cameras, candidate count, validity or count metadata, oracle metrics as labels, optional rendered depths, compact OBBs, trajectory metadata, and selected EVL numeric fields.

state **– Rollout State**: Family of rollout state variants used by ARIA-NBV. The actor-visible variants contain logged or replayed observations, reconstruction proxies, candidates, validity masks, target descriptors, history, and budget; the oracle variant adds privileged GT geometry and all-candidate labels for supervision and evaluation only.

***NBV MDP* – Target-Conditioned NBV MDP**: Finite-horizon Markov decision process used to define ARIA-NBV's target-conditioned rollout and fitted Q_H training contract. The actor observes only current reconstruction state, sampled candidate views, validity masks, target descriptors, history, and budget state; oracle meshes and GT crops are supervision and evaluation signals.

***mask* – Validity Mask**: Hard candidate-level feasibility mask used during candidate ranking, rollout generation, and Q_H training. Invalidity is represented separately from reconstruction quality and should not be encoded as the lowest RRI class.

Protocol

***GT-EVAL* – Ground-Truth Target Evaluation**: Main thesis protocol component restricting ground-truth OBBs and target mesh crops to oracle labels, target-RRI computation, matching checks, and evaluation. GT target crops are not actor-visible inputs in the main result.

***OBS-SEL* – Observed Target Selection**: Main thesis protocol component requiring target selection to use only actor-visible evidence such as predicted or tracked OBBs, class probabilities, confidence, projected area, and semidense or EVL support. Ground-truth target annotations are not visible to the selector in this protocol.

***PRED-Q* – Predicted-Target Q**: Main thesis protocol component requiring the scorer or fitted Q_H model to condition on predicted or observed target descriptors rather than ground-truth target inputs. It covers target-conditioned one-step scoring and finite-candidate Q_H selection.

Reconstruction

***MVS* – Multi-view Stereo**: Reconstruction family that estimates dense or semi-dense scene geometry from multiple calibrated views. MVS is part of the broader reconstruction context for NBV, although ARIA-NBV's current oracle labels are built from ASE meshes and Aria-style semi-dense point observations.

***PC* – Point Cloud**: Set of 3D points representing observed scene geometry. ARIA-NBV compares current and candidate-fused point clouds against ground-truth meshes to compute oracle RRI and surface-distance diagnostics.

***SLAM* – Simultaneous Localization and Mapping**: Method for estimating sensor motion while building a map of the surrounding scene. In ARIA-NBV, logged SLAM poses and semi-dense points form the current reconstruction state against which candidate-view RRI is evaluated.

***track* – Track**: Temporal sequence of corresponding image-feature detections across frames, usually carrying per-frame image coordinates, timestamps, and camera IDs. Tracks can be triangulated or optimized into 3D points, making them a bridge between image evidence and the semi-dense point clouds used by ARIA-NBV.

***VIO* – Visual-Inertial Odometry**: Pose-estimation method that combines visual measurements from cameras with inertial measurements from IMUs. Project Aria pose and mapping products build on visual-inertial estimation, and ARIA-NBV depends on those pose streams for snippet state and candidate frame contracts.

Representation

***Occupancy Grid* – Occupancy Grid**: Spatial grid whose cells encode whether space is occupied, free, unknown, or represented by a related occupancy probability. Occupancy-style voxel evidence appears in EVL outputs and VIN feature construction, where it helps summarize local 3D scene state for candidate scoring.

***3DGS* – 3D Gaussian Splatting**: Explicit radiance-field representation using optimized 3D Gaussian primitives for real-time novel-view synthesis. ARIA-NBV treats active 3DGS work as proposal, uncertainty, and simulator-bridge literature, not as a replacement for ASE mesh-supervised RRI.

Supervision

oracle RRI **– Oracle RRI**: RRI label computed with privileged ground-truth geometry, used for supervised training and evaluation. The current oracle renders candidate depth maps from ASE ground-truth meshes, backprojects candidate point clouds, fuses them with the current semi-dense reconstruction, and scores the resulting surface-distance improvement.

List of Symbols

Symbol	Meaning	Key
	Oracle and Reconstruction	
RRI	Relative Reconstruction Improvement score for a candidate view.	`oracle.rri`
$\mathcal{P}$	Actor-visible point set accumulated from observed or replayed views.	`oracle.points`
$\mathcal{P}_q$	Candidate-view point contribution used in oracle or counterfactual updates.	`oracle.points_q`
$\mathcal{Q}$	Finite candidate-view set considered by the planner.	`oracle.candidates`
$\mathcal{Q}_t$	Candidate-view set available at rollout step t.	`oracle.candidates_t`
$q_{t,i}$	Candidate pose i at rollout step t.	`oracle.candidate_qti`
$\boldsymbol{D}_q$	Oracle-rendered or candidate-specific depth map for a proposed view.	`oracle.depth_q`
$D_{P \to M}$	Point-to-mesh directional error component.	`oracle.acc`
$D_{M \to P}$	Mesh-to-point directional error component.	`oracle.comp`
	RRI Metrics	
D	Aggregate point-mesh reconstruction error used by RRI definitions.	`rri.cd_value`
	ASE Assets	
$\mathcal{M}^{\mathrm{GT}}$	Ground-truth scene mesh used for oracle labels and evaluation.	`ase.mesh`
$\mathcal{M}_e^{\mathrm{GT}}$	Ground-truth mesh crop for target e.	`ase.mesh_target`
$\mathcal{F}^{\mathrm{GT}}$	Ground-truth mesh faces used in point-mesh distance calculations.	`ase.faces`
$\boldsymbol{T}_{\mathrm{rig}}^{\mathrm{w}}(t)$	Project Aria rig pose in the world frame at time t.	`ase.traj`
$\boldsymbol{T}_{\mathrm{rig}}^{\mathrm{w}}(T)$	Final rig pose used as an anchor for gravity-aligned local fields.	`ase.traj_final`
	VIN Scorer	
r	Target or scene RRI value used as a model target.	`vin.rri`
$\hat{r}$	Predicted RRI value emitted by the learned scorer.	`vin.rri_hat`
$\mathcal{L}$	Training loss for scorer or value-model objectives.	`vin.loss`
m	Candidate validity indicator used by scorer diagnostics.	`vin.cand_valid`
	Entity and Target	
RRI_e	Target-specific Relative Reconstruction Improvement for entity e.	`entity.rri_e`

Symbol	Meaning	Key
$\boldsymbol{\varphi}_e$	Actor-visible target/entity descriptor used to condition target-specific value prediction.	`entity.target_desc`
	Planning and Rollout	
s	Abstract planning state.	`rl.s`
a	Action index selecting a candidate row.	`rl.a`
r	Scalar reward or immediate gain.	`rl.r`
G	Finite-horizon return accumulated across rollout steps.	`rl.G`
$\mathcal{M}_{\mathrm{NBV}}$	Target-conditioned finite-candidate NBV decision process.	`rl.mdp_nbv`
$\mathcal{A}(s_t)$	Valid action set available in state s_t.	`rl.action_set`
T	State-transition operator for selected candidate actions.	`rl.transition`
s_t^{hist}	Logged historic snippet state at rollout step t.	`rl.s_hist`
s_t^{off}	Persisted offline sample state used by training readers.	`rl.s_off`
s_t^{cf0}	Minimal actor-visible counterfactual state.	`rl.s_cf0`
$s_t^{\mathrm{cf+}}$	Geometry-enriched counterfactual successor state.	`rl.s_cf_geom`
s_t^{oracle}	Oracle-only state used for labels, upper bounds, and evaluation.	`rl.s_oracle`
r_t^e	Target-specific reward at rollout step t.	`rl.reward_target`
$G_t^{(H)}$	H-step finite-horizon return from rollout step t.	`rl.return_h`
Q_H	Finite-horizon candidate-value function.	`rl.qh`
γ	Return discount factor.	`rl.gamma`
H	Planning or rollout horizon length.	`rl.H`
$m_{t,i}$	Hard validity mask for candidate i at rollout step t.	`rl.validity_mask`
$\rho_{t,i}$	Invalid-action reason code for candidate i at rollout step t.	`rl.invalid_reason`
e_t	Selected target entity at rollout step t.	`rl.target`
b_t	Remaining acquisition budget at rollout step t.	`rl.budget`
$C(\tau)$	Acquisition cost of a selected trajectory.	`rl.acquisition_cost`
	Shape and Size	
N_q	Number of candidate views in a candidate table.	`shape.Nq`
K	Number of ordinal bins or discrete levels when used in scorer losses.	`shape.K`
	scene	
$\boldsymbol{M}_t^{\mathrm{ray}}$	Sparse actor-visible ray-aware memory separating occupied, free, and unknown support at step t.	`scene.ray_memory_t`
	Observation and Actor State	

Symbol	Meaning	Key
$\boldsymbol{F}_t^{\text{DINO@pt}}$	Visibility-gated logged DINO feature bank attached to semidense or fused world points.	`obs.dino_point_bank_t`
$\boldsymbol{X}_t^{\text{pt}}$	Point-token set combining geometry, support, uncertainty, history, and optional compressed logged descriptors.	`obs.point_tokens_t`
	scene	
$\boldsymbol{\Phi}_t^{\text{scene}}$	Composite actor-visible scene-memory descriptor queried by the finite-candidate value model.	`scene.scene_memory_t`
$\boldsymbol{E}_0^{\text{EVL-local}}$	Root local EVL evidence field used for target/candidate support reads.	`scene.evl_local`
$\omega_{t,i}^{\text{EVL}}$	Candidate or target-query fraction supported by the root EVL voxel extent.	`scene.evl_support_frac`
$\boldsymbol{g}_{t,i}^{\text{EVL}}$	Pooled local EVL evidence token for a target, candidate frustum, or target-candidate intersection query.	`scene.evl_support_token`
$\boldsymbol{g}_e^{\text{tgt}}$	Pooled target-support descriptor for the selected target hypothesis.	`scene.target_support_pool`
$\boldsymbol{g}_{t,i}^{\text{fr}}$	Candidate-frustum support descriptor pooled from actor-visible scene memory.	`scene.frustum_support_pool`
$\boldsymbol{g}_{t,e,i}^{\text{cap}}$	Pooled descriptor over the intersection between target support and candidate frustum support.	`scene.target_frustum_pool`
$\boldsymbol{g}_{t,i}^{\text{ray}}$	Candidate-conditioned ray query over occupied, free, unknown, hit, target-support, and uncertainty summaries.	`scene.ray_query_ti`
RenderQuery	Abstract query operator that summarizes scene memory in a candidate camera/frustum without introducing fresh counterfactual RGB features.	`scene.render_query`
	spatial	
r_t	Reference pose for candidate-relative descriptor construction at decision step t.	`spatial.ref_pose`
$\boldsymbol{T}_{r_t,i}^{\text{rel}}$	Relative transform from the reference pose r_t to candidate camera i.	`spatial.ref_candidate_transform`
$\boldsymbol{h}_{t,i}^{\text{pose}}$	Relative/local candidate pose descriptor derived from a reference pose rather than raw world coordinates.	`spatial.candidate_pose_feat`
$\boldsymbol{h}_{t,e\|i}^{\text{rel}}$	Candidate-target relation descriptor in the candidate/query local frame.	`spatial.candidate_target_rel_feat`
$\boldsymbol{e}_{a\|i}^{\text{rel}}$	Query-local relative positional embedding for target, history, support, or candidate relations.	`spatial.relation_rpe`
	model	

Symbol	Meaning	Key
$\boldsymbol{h}_e^{\mathrm{tgt}}$	Learned selected-target token built from the actor-visible target descriptor and scene support.	`model.target_token`
$\boldsymbol{x}_{t,i}$	Per-candidate row feature assembled from pose, relation, support, validity, provenance, and history descriptors.	`model.candidate_row`

List of Figures

List of Tables

1 Introduction

Active reconstruction couples sensing and inference: after a fixed acquisition budget, reconstruction quality depends not only on the surface estimator but also on which feasible viewpoints were acquired. Classical active perception therefore treats sensing actions as part of perception, and view-planning work formalizes the recurring generate–score–select loop for three-dimensional inspection [AWB88, Baj88, SRR03]. This thesis studies a deliberately bounded instance of that problem: target-conditioned view selection from a finite admissible candidate table for egocentric indoor reconstruction.

VIN-NBV provides the closest objective precedent by supervising candidate ranking with RRI, computed from the reduction in point–mesh reconstruction error after adding a query view [Fra+25]. ARIA-NBV retains this quality-driven principle but changes the task and planning question. The utility is target-specific rather than scene-wide, and non-myopic value is interpreted only if bounded oracle lookahead improves endpoint quality over one-step selection under the same finite candidate support. Project Aria, ASE, and Egocentric Foundation Model 3D (EFM3D) provide the calibrated egocentric, mesh-supervised, and local actor-visible substrate for this study [Eng+23, Met25a, Str+24].

The current data-generation tasks are defined directly from geometry-valid GT oriented bounding boxes. They provide target identity, target crops, and oracle labels, but they do not implement actor-visible target discovery or identity matching. Consequently, deployable-input claims require a separate observed- or predicted-target protocol; the current target tasks support oracle supervision and controlled evaluation only.

Problem statement. Given a GT-defined target task, actor-visible reconstruction evidence, a finite candidate table with hard validity constraints, and a fixed acquisition budget, determine whether bounded oracle lookahead improves endpoint target reconstruction over one-step oracle-greedy selection and, only if such headroom exists, whether a learned policy using non-privileged inputs recovers a meaningful fraction of it under the same oracle evaluation.

The thesis separates supervision from decision-time information. Ground-truth geometry defines current target tasks, target crops, oracle labels, and evaluation; it is not an ordinary actor input. Invalid candidates are outside the admissible action set rather than examples with low utility. One-step oracle greedy is an immediate-reward comparator over the evaluated valid candidate table, while oracle lookahead is an upper reference only within the fixed candidate generator, horizon, branch factor, target pool, and validity regime. These restrictions make negative results interpretable without turning a bounded experiment into a universal statement about view planning.

1.1 Research Questions

RQ1 — Objective and endpoint contract. Do target-cropped point–mesh error, root-normalized immediate gain, and fixed-budget endpoint gain provide a repeatable objective for comparing target-conditioned view-selection policies?

RQ2 — Offline finite-candidate planning. Under identical target tasks, roots, candidate support, validity rules, and acquisition budgets, does bounded oracle lookahead exhibit positive endpoint-quality headroom over one-step oracle-greedy selection, and can an offline finite-horizon policy recover a meaningful fraction of that headroom?

RQ3 — Actor-visible representation. Which observed or predicted target descriptor and actor-visible state representation are sufficient for learned one-step and finite-horizon candidate scoring without privileged target geometry or all-candidate oracle renders at decision time?

RQ4 — Support and scale. How do target-task coverage, candidate support, hard validity, rollout diversity, and scene-level scale constrain the reliability and generality of the conclusions for RQ1–RQ3?

Together, these questions delimit the evaluated contribution; online control, continuous actions, and real-device deployment remain outside it.

2 Foundations

ARIA-NBV combines ASE trajectories and GT geometry, local EFM3D evidence, oracle RRI labels, and VIN-style finite-candidate scoring. ASE supplies calibrated Aria-like streams and aligned semi-dense maps; EVL supplies frozen local voxel evidence rather than a complete long-horizon scene memory [Met25a, Met25b, Str+24]. GT OBBs currently define data-generation target tasks, whereas observed or predicted target discovery remains outside the implemented actor path.

2.1 Related Work

Classical active perception establishes that sensing actions and inference must be designed jointly, while three-dimensional view planning supplies the finite generate–score–select abstraction used here [AWB88, Baj88, SRR03]. Receding-horizon and submodular formulations explain why lookahead and greedy comparators can be informative under specific utility assumptions [Bir+16, Lau+20]. They do not establish such guarantees for target-specific RRI, whose value depends on the evaluated target, current reconstruction support, and admissible candidates. This thesis therefore tests non-myopic headroom empirically under fixed support rather than assuming it from generic coverage objectives.

Quality-driven NBV replaces geometric coverage proxies with reconstruction-quality supervision. VIN-NBV is the closest precedent because it ranks sampled candidates using oracle RRI rather than visibility alone, while PB-NBV contributes efficient finite projection-based proposals [Fra+25, Jia+25]. Neither establishes target-cropped finite-horizon value under an Aria-native actor/oracle separation. ARIA-NBV consequently uses VIN-style one-step scoring as a myopic control and asks separately whether bounded lookahead improves endpoint target quality.

Target- and object-aware NBV shows why scene-wide utility can conceal the reconstruction needs of one entity. Object-centric NBV motivates conditioning view utility on a designated object, while OANBV and ObjViewBench contribute visibility, feasibility, occlusion, and target-difficulty analyses [Hu+26, Jeo+26, Pan+26]. The current ARIA-NBV data generator does not solve actor-visible target discovery: it samples geometry-valid GT OBB rows as target tasks and evaluates their target crops. Observed or predicted target records are therefore a separate requirement for learned-policy claims, not an implemented consequence of target-specific supervision.

Project Aria defines the calibrated multimodal egocentric sensing regime, ASE supplies synthetic trajectories and GT geometry, and EFM3D/EVL supplies local actor-visible spatial evidence and object hypotheses [Eng+23, Met25a, Str+24]. This stack enables controlled counterfactual labels but does not itself define an NBV objective or make oracle information actor-visible. Meshes, GT boxes, and all-candidate renders remain supervision and evaluation assets; candidate selection must ultimately use logged observations, accumulated selected-view evidence, and an observed or predicted target record.

Offline value learning and set models provide controls for the finite-action problem. Double DQN separates successor-action selection from evaluation, while CQL and BCQ illustrate the risk of unsupported actions in offline learning [FMP19, HGS15, Kum+20]. Deep Sets and Set Transformer provide permutation-aware baselines for an unordered candidate table [Lee+19, Zah+17]. These methods do not replace the domain contract: invalid actions remain hard-masked, candidate support is reported, and every selected trajectory is re-evaluated by endpoint target quality rather than accepted on predicted value alone.

Geometric representation learning supplies a hierarchy of inductive biases rather than one mandatory backbone. Query-local relative encoding can remove nuisance global coordinates while preserving target, frustum, and motion geometry [Zho+23]. PointNeXt, Point Transformer, KPConv, and sparse convolutions offer coordinate-bearing encoders when explicit point or ray-memory summaries become insufficient [CGS19, Qia+22, Tho+19, Wu+24]. EGNN, SE(3)-Transformer, and GATr provide stronger equivariant alternatives [Bre+23, Fuc+20, SHW21], but exact equivariance is not assumed beneficial for a gravity-aligned, metric, egocentric task; it remains an ablation against simpler local-frame features.

Broader scene memories address a different limitation from candidate interaction. GenNBV distinguishes occupied, free, and unknown grid state, Hestia records directional visibility, SceneScript extracts broad layout and object structure from Aria-style semidense input, and object-aware 3D Gaussian Splatting couples renderable primitives with target membership [Ave+24, Che+24, Jeo+26, Lu+26]. ARIA-NBV uses these works to define representation hypotheses and diagnostics; it does not import their rewards, action spaces, or performance claims into the finite-candidate target-RRI experiment.

2.2 Geometric Learning and Candidate-Set Theory

The learned object in ARIA-NBV is a typed finite-candidate decision map: given a target record, selected history, partial actor-visible geometry, and an unordered table of feasible candidate poses, assign one score to each candidate for selection under later oracle re-evaluation. The model must respect the symmetries and information boundary of this map before architectural capacity is interpreted [Bro+21].

The first required structure is candidate-row permutation equivariance. Reordering $\mathcal{Q}_t$ may reorder per-candidate Q_H outputs, but it must not change the value assigned to the same physical candidate. Deep Sets supplies a pooled symmetry baseline, while Set Transformer supplies candidate interaction without assigning semantic meaning to row order [Lee+19, Zah+17]. Both retain a row-level scoring path because the policy selects a candidate row.

For any permutation matrix Π acting on candidate rows, the value model should satisfy

$$f_\theta(\Pi \boldsymbol{X}_t, \Pi \boldsymbol{m}_t) = \Pi f_\theta(\boldsymbol{X}_t, \boldsymbol{m}_t)$$

where $\boldsymbol{X}_t = \{\boldsymbol{x}_{t,i}\}_{i=1}^{N_q}$ stores per-candidate actor-visible descriptors and $\boldsymbol{m}_t$ stores validity. Invalid and padded rows are outside the admissible action set, not low-utility examples. Selection is therefore

$$\mathcal{A}_t = \{i \in \{1, ..., N_q\} : m_{t,i} = 1\}, \quad a_t^\theta = \operatorname*{argmax}_{i \in \mathcal{A}_t} f_{\theta,i}(\boldsymbol{X}_t, \boldsymbol{m}_t)$$

The second required structure is frame discipline. Candidate, target, current-camera, and history poses use reference- or query-local relative features such as $\boldsymbol{T}^{r}_{c_q}$ and continuous rotation encodings [Zho+19, Zho+23]. This prevents an arbitrary world origin or row order from becoming a shortcut while preserving gravity alignment and camera-frustum geometry. It is a controlled coordinate choice, not a claim of full global SE(3) invariance.

The third required structure is explicit selected-view history. A camera pose also defines a viewing direction from which target-local surfaces may or may not have been observed. A compact history on $\mathbb{S}^2$ can therefore distinguish directional novelty from generic pose distance [e3n25]. Coverage, overlap, and uncertainty remain diagnostic features unless paired oracle evaluation shows that they improve target-specific RRI [GML22, JLD24].

Object	**Required structure**	**Testable consequence**
Candidate rows	Permutation-equivariant per-row value map.	Row-shuffle tests must satisfy $f_{\theta(\Pi X, \Pi m)} = \Pi f_{\theta(X,m)}$ for every selection score.
Invalid and padded rows	Separate feasibility head, trained on every non-padding row; RRI/value loss only on rows certified valid. Hard action masking is the initial policy, with a later calibrated soft feasibility gate.	Invalid rows cannot change valid-row scores except through explicit valid-count or support features.
Candidate-target geometry	Target-local and candidate-local relative pose features.	World-frame origin, yaw convention, and display-only camera transforms cannot become shortcuts.
Selected-view history	Directional memory on $\mathbb{S}^2$.	Novelty tests separate already-observed target directions from generic pose distance.
Actor/oracle boundary	Typed provenance for target descriptors, labels, and support features.	GT-defined tasks, meshes, crops, and all-candidate renders supervise data generation and evaluation only.

Table 1: Minimum geometric-learning contract for the finite-candidate value model.

The implemented hard mask defines $m_{t,i} = 0$ as exclusion from the action set, not as a numerical return. An invalid row's RRI is undefined: it must be neither forced to zero nor assigned a synthetic negative target, because zero can be a legitimate valid-row RRI and either choice would conflate feasibility with quality.

Implementation: planned **Evidence: pending**

A separate feasibility head is planned. It is not part of the current scorer or training module.

Literature: [GE19, Liu+19]

Promotion gate: implement feasibility head and evaluate held-out calibration without weakening the hard mask

Development source: aria_nbv/aria_nbv/rollouts/zarr_store.py; planned finite-horizon scorer

For each non-padding row, the planned model emits a feasibility probability and an RRI/value estimate. Binary cross-entropy supplies negative evidence on invalid rows, while the RRI loss is evaluated only when $m_{t,i} = 1$. This factorization treats feasibility as a reject decision rather than an artificial low-RRI target.

The hard mask remains the deployment safety constraint. A soft feasibility score is only a training and ranking ablation: after held-out feasibility calibration, compare a score proportional to predicted validity times the conditional valid-row value against the hard-mask baseline, while retaining the hard mask wherever it is available. The abstention literature likewise models rejection as a separate outcome rather than a continuous-regression target [Liu+19]. The curriculum hypothesis is therefore to train feasibility from the first batch, warm up the value head only on oracle-valid rows, and then introduce the calibrated soft-gate ablation [Ben+09]. It must never reintroduce an RRI target or RRI gradient for invalid rows.

Research TODO: Compare hard masking against a calibrated feasibility-times-value ranking only after feasibility calibration is measured. The hard mask remains the safety constraint in every comparison.

Source: invalid-row supervision hypothesis

Gate: feasibility calibration and policy ablation

These requirements become replay-field and model acceptance tests in Chapter Section 4.4. Row shuffles must permute outputs, invalid-row perturbations must leave valid scores unchanged, and equivalent local-frame encodings must preserve physical candidate identity. Architectural comparisons are interpretable only after those contracts hold.

The resulting foundation is narrow: active perception motivates action-conditioned sensing, VIN-NBV supplies the quality-driven one-step precedent, ASE and EFM3D define the logged egocentric evidence boundary, and finite-action learning supplies mask and support controls. Coverage and uncertainty remain diagnostics rather than substitutes for target-specific reconstruction quality.

3 Oracle and Data Generation

This chapter defines the non-deployable pipeline that turns logged ASE snippets into supervised target-conditioned NBV tasks. It separates actor state from privileged data-generation state, specifies target-task and candidate construction, defines hard invalidity and target-specific RRI, and records the resulting selected counterfactual chains. The learned method in Chapter 4 may consume only the actor-side projection of these artifacts; GT geometry, counterfactual renders, labels, and oracle search remain instruction, supervision, or evaluation assets.

3.1 State and Visibility Boundary

The thesis studies a masked finite-horizon candidate-decision process through offline, mesh-supervised counterfactual replay. Logged egocentric observations and frozen EVL evidence define the actor substrate, whereas ASE meshes, annotations, synthetic target instructions, rendered counterfactual geometry, and labels define privileged data generation [Fra+25, Met25a, Str+24]. An oracle product may supervise a loss or define an upper bound, but it may not enter the learned Q_H actor unless the experiment is explicitly labelled privileged.

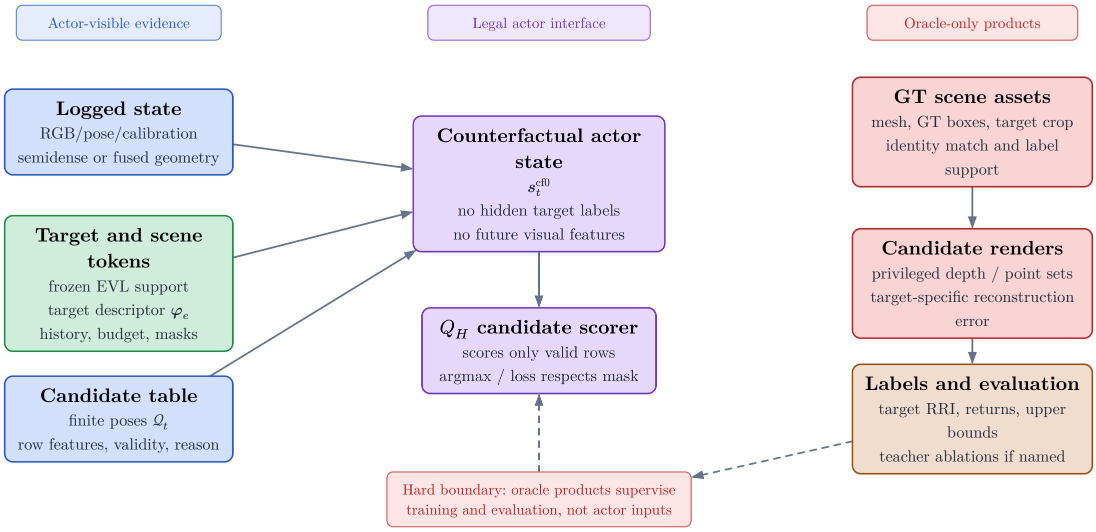

Figure 1: Actor and oracle boundary. Legal Q_H inputs are accumulated actor geometry, frozen EVL evidence, an explicitly declared target instruction, selected-view history, remaining budget, candidates, masks, and reason codes. GT geometry, target crops, dense counterfactual renders, labels, and endpoint evaluation remain privileged.

The logged historic state contains the recorded trajectory evidence,

$$s_t^{\mathrm{hist}} = \left(\boldsymbol{I}^{\mathrm{rgb}}, \boldsymbol{X}, \mathcal{P}^{\mathrm{semi}}, \boldsymbol{F}_v, e_t, b_t\right) \in \mathcal{S}^{\mathrm{hist}}$$

the counterfactual actor state adds only evidence acquired along the selected synthetic history,

$$s_t^{\mathrm{cf0}} = \left(\boldsymbol{F}_v^{\mathrm{root}}, \mathcal{P}_t, \mathcal{Q}_t, m_{t,i}, \rho_{t,i}, e_t, b_t\right) \in \mathcal{S}^{\mathrm{cf0}}$$

and the oracle state adds privileged geometry and labels outside the actor input graph:

$$s_t^{\mathrm{oracle}} = \left(s_t^{\mathrm{cf+}}, \mathcal{M}^{\mathrm{GT}}, \mathcal{M}_e^{\mathrm{GT}}, \boldsymbol{D}_q, \mathcal{P}_q, \mathrm{RRI}\right) \in \mathcal{S}^{\mathrm{oracle}}$$

$$\mathcal{M}_{\mathrm{NBV}} = \left(\mathcal{S}^{\mathrm{hist}}, \mathcal{S}^{\mathrm{cf0}}, \mathcal{S}^{\mathrm{oracle}}, \{\mathcal{A}_t\}, T, r_t^e, \gamma, H\right)$$

This separation is necessary because the logged snippet contains modalities that cannot be regenerated after a synthetic action. The counterfactual state therefore keeps root features fixed and updates only the fused geometry proxy, selected-view history, budget metadata, finite candidates, masks, and reason codes. In the current oracle protocol, the target descriptor supplied to candidate generation is a privileged task instruction derived from the selected GT box; it is not evidence that an actor-visible detector found or localized the target. A deployable experiment must replace that instruction with an observation-derived descriptor and preserve the same crop-free actor interface.

Invalidity is a constraint rather than a reward value. Geometry-invalid candidates receive a hard action mask and a persisted candidate reason code before policy selection, stochastic normalization, loss construction, and bootstrap maximization. A geometrically feasible candidate may still have low or negative target gain. Failure of the separate oracle evaluation does not create a candidate reason code: depending on the configured recipe, the affected row or table is skipped, or its `oracle_label_mask` and `q_train_mask` are cleared and the oracle failure is reported separately.

3.2 Data Generation and Target-Specific RRI Labels

An oracle target task fixes the entity identity and GT evaluation crop. A deployable target descriptor would instead be predicted from actor-visible observations. The current rollout generator has only the first contract: it projects a selected GT box into a compact instruction for candidate generation. Consequently, the present target descriptor denotes the requested object and its geometry; it does not establish actor-visible target discovery.

The general descriptor notation remains

$$\boldsymbol{\varphi}_e = \mathrm{Enc}_{\mathrm{tgt}}\left(\hat{\boldsymbol{B}}_e, \hat{\boldsymbol{y}}_e, \hat{\pi}_e, A_e^{\mathrm{proj}}, n_e^{\mathrm{semi}}, n_e^{\mathrm{EVL}}, \omega_e^{\mathrm{EVL}}, \ell_e^{\mathrm{src}}, \boldsymbol{T}_{r_t,e}, \boldsymbol{T}_{c_t,e}\right)$$

The implemented oracle sampler populates only GT identity, class, confidence, pose, extent, and reference-relative geometry. Projected area, semidense support, EVL support, and proposal-match scores are not measured by this path and must not be interpreted from their schema placeholders.

The finite action interface consists of a candidate table $\mathcal{Q}_t$, hard validity mask $\boldsymbol{m}_t$, and invalid-reason vector $\boldsymbol{\rho}_t$. The admissible set is $\mathcal{A}_t = \left\{i : m_{t,i} = 1\right\}$. Invalid rows remain logged for coverage and failure analysis but lie outside policy argmax, sampling, loss targets, and bootstrap maximization. Low RRI is a valid low-utility outcome; it is never an encoding for infeasibility.

For target e, the oracle computes a target-cropped point–mesh error Δ_t^e. Candidate selection and Q_H supervision use root-normalized target gain; state-relative RRI is retained as a diagnostic. Fixed-budget endpoint gain is the intended policy estimand, but the current replay store records cumulative selected-chain gains rather than an independent post-horizon endpoint reconstruction for every policy. Confirmatory policy comparisons must therefore add matched oracle endpoint re-evaluation or an explicitly persisted endpoint record.

3.2.1 Oracle Label Provenance

The one-step scene-level oracle from the seminar work supplies the rendering and point–mesh substrate [Fra+25, Met25c]. The target-specific pipeline renders valid candidates from the GT mesh, backprojects their depth, fuses candidate points with the privileged current evaluation cloud, crops both points and mesh to the selected target box, and evaluates the resulting point–mesh error. In the current protocol the root evaluation cloud is reconstructed from ASE GT depth, so both the crop and the root metric are oracle-only.

Camera geometry is therefore useful here only as part of a relation: a logged trajectory, a target crop, a finite candidate set, or a rendered oracle query. The calibrated camera boundary is not a contribution by itself. The consolidated candidate-support figure in Figure 3 uses camera glyphs only to expose those relations; native `CameraTW` unprojection and world-from-camera `PoseTW` remain the geometry owners behind the render path.

3.2.2 Target Selection

The implemented sampler does not match actor proposals to GT objects. It enumerates the non-padding GT OBB rows in a snippet, accepts rows with finite positive geometry, and applies seeded uniform sampling without replacement up to the configured per-snippet cap. Thus the stored `matched` status currently means geometry-valid GT task row, not successful proposal-to-identity association. IoU, ambiguity-gap, visibility, and support thresholds are absent from this admission rule.

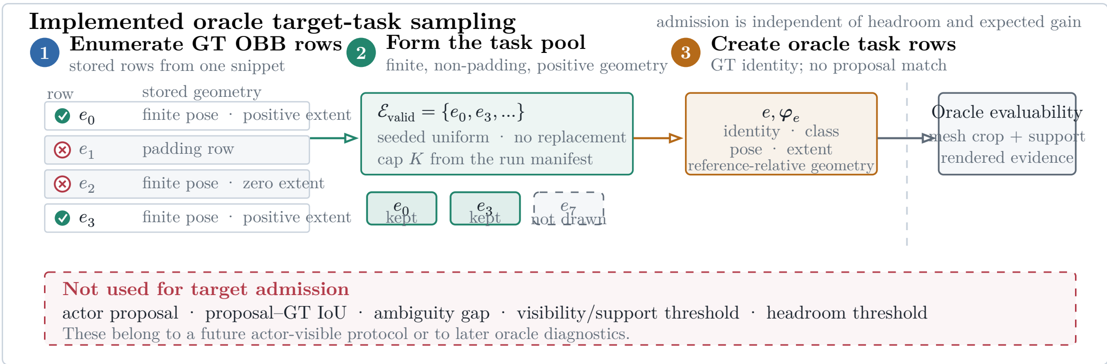

Figure 2: Implemented oracle target-task sampler. Geometry-valid GT OBB rows form the task pool, and seeded uniform sampling without replacement applies the manifest-defined cap. Rollout scoring later decides whether the selected crop is evaluable. No actor proposal, IoU match, visibility gate, or support gate is used.

The sampler bounds rollout cost without pre-filtering on headroom. Candidate scoring may subsequently invalidate the task when its mesh crop, current support, or rendered evidence is unusable; otherwise near-solved and negative-gain targets remain scientifically informative. A later actor-visible protocol must introduce proposal identity and observation-quality diagnostics as a separate selection stage rather than retroactively interpreting the present oracle fields as measurements.

3.2.3 Candidate View Generation

Candidate generation turns one oracle instruction into a finite action table. Every quantitative choice—source population, target cap, shell size, family weights, motion limits, pruning thresholds, rollout recipes, renderer settings, and retention policy—belongs to the resolved run manifest and report bundle. The Methods text defines semantics only; it does not designate one mutable TOML profile as canonical.

At rollout step t, candidate generation constructs a full shell

$$\mathcal{Q}_t = \{q_{t,i}\}_{i=1}^{N_q}, \quad N_q = 60$$

with a fixed provenance component $k(i)$ per row. The core mixture contains forward-local, target-bearing, and lateral-bypass motion; the diversity challenger may additionally allocate mass to local refinement and revisit/backtrack components. The resolved manifest, rather than prose, determines which components and counts apply to a run.

For each row, the raw direction is sampled in the reference rig frame. The realistic profile uses the forward-biased Power Spherical distribution from the current implementation [DAT20]:

$$\boldsymbol{u}_i \sim \mathrm{PS}(\boldsymbol{e}_z, \kappa), \quad p(\boldsymbol{u}) = c_\kappa \big(1 + \boldsymbol{e}_z^T \boldsymbol{u}\big)^\kappa, \quad \kappa = 8$$

The draw is mapped into configured azimuth and elevation caps without rejection. With $\psi = \operatorname{atan2}(u_x, u_z)$ and $u_y = \sin\theta$, the cap transform is

$$\psi' = \psi \frac{\Delta_\psi}{2\pi}, \quad y' = \sin\theta_{\min} + \frac{u_y + 1}{2} \cdot (\sin\theta_{\max} - \sin\theta_{\min})$$

$$\boldsymbol{d}_i^0 = \left\| \left(\sqrt{1 - y'^2} \sin\psi', y', \sqrt{1 - y'^2} \cos\psi' \right) \right\|$$

The sampler draws a capped direction in the reference rig frame and reinterprets it as egocentric forward motion, target-bearing motion, lateral bypass, local refinement, or backtracking according to component provenance. Target-looking families orient their optical axis toward the oracle instruction. In this chapter that point is privileged; calling the generator target-conditioned does not make the point actor-visible.

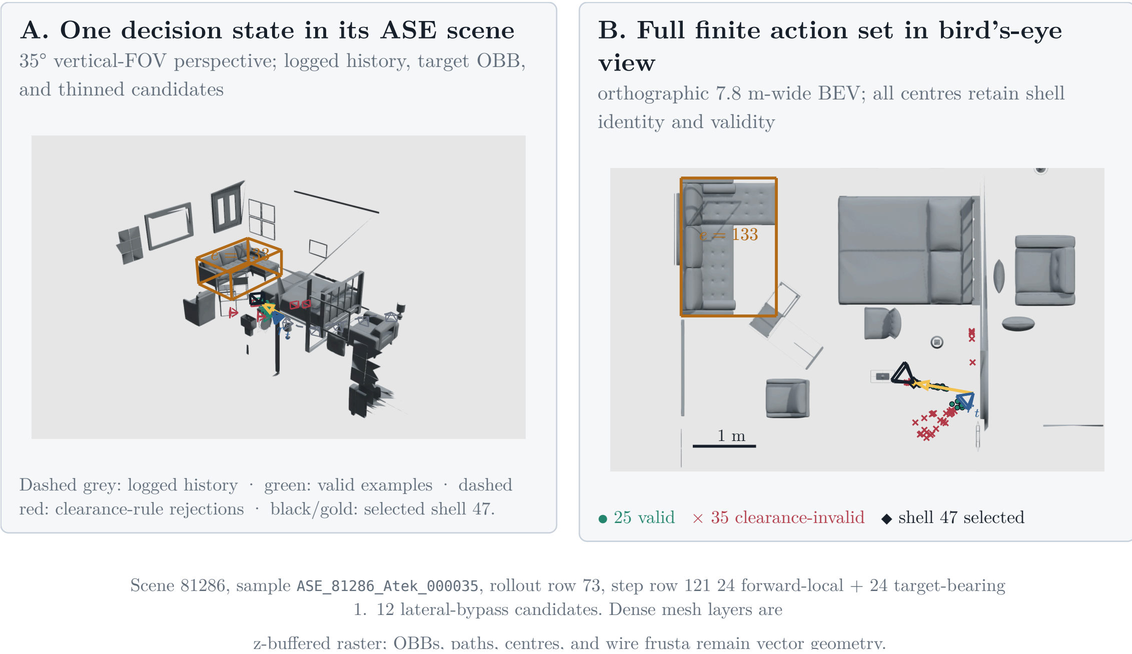

Figure 3: One pinned finite-candidate decision state from ASE scene 81286, sample `ASE_81286_Atek_000035`, rollout row 73, and step row 121. Panel A uses a 35-degree vertical-FOV perspective view to place logged camera history, root state, target OBB, selected path, and a deterministically thinned set of wire frusta in the real scene. Panel B uses a 7.8-metre-wide orthographic bird's-eye view and retains all 60 candidate centres: 25 rows are admissible, 35 are hard-rejected by the clearance rule, and oracle-greedy selects shell 47. Dense scene geometry is z-buffered; OBBs, paths, centres, and camera glyphs remain vector overlays. Full eye, look-at, up, clipping, and resolution parameters are recorded in the figure JSON. This is an auditable contract example, not a policy-performance result.

The three core position families then reinterpret this capped direction. Let $\boldsymbol{f} = \boldsymbol{e}_z$ be the rig-forward unit vector, $\boldsymbol{b}_e$ the supplied target bearing in the reference frame, $\boldsymbol{l}_e = \|\boldsymbol{e}_y \times \boldsymbol{b}_e\|$ the horizontal lateral direction, and $\boldsymbol{e}_y$ the world-up direction expressed in the sampling frame. The family directions are:

$$\boldsymbol{d}_i^{\text{forward}} = \left\|\boldsymbol{f} + \alpha_f(\boldsymbol{d}_i^0 - (\boldsymbol{d}_i^0 \cdot \boldsymbol{f})\boldsymbol{f})\right\|, \quad \alpha_f = 0.45$$

$$\boldsymbol{d}_i^{\text{target}} = \left\|\boldsymbol{b}_e + \alpha_t(\boldsymbol{d}_i^0 - (\boldsymbol{d}_i^0 \cdot \boldsymbol{b}_e)\boldsymbol{b}_e)\right\|, \quad \alpha_t = 0.4$$

$$\boldsymbol{d}_i^{\text{bypass}} = \left\|0.55\boldsymbol{b}_e + 0.85\,\text{sign}(d_{i,x}^0)\boldsymbol{l}_e + \text{clip}\left(d_{i,y}^0, -0.35, 0.35\right)\boldsymbol{e}_y\right\|$$

Finally, the sampler draws a radius and transforms the reference-frame offset into world coordinates:

$$r_i \sim \mathcal{U}(0.25, 1.1), \quad \boldsymbol{c}_i^w = \boldsymbol{T}_r^w \left(r_i \boldsymbol{d}_i^{k(i)} \right)$$

`forward_local` keeps the reference rig orientation. Target-looking families orient the camera toward the supplied target center $\boldsymbol{p}_e$:

$$\boldsymbol{z}_i^w = \|\boldsymbol{p}_e - \boldsymbol{c}_i^w\|, \quad \boldsymbol{y}_i^w = \left\|\boldsymbol{e}_y - \left(\boldsymbol{e}_y^T \boldsymbol{z}_i^w\right)\boldsymbol{z}_i^w\right\|, \quad \boldsymbol{x}_i^w = \boldsymbol{y}_i^w \times \boldsymbol{z}_i^w$$

These equations describe the sampler, not an optimal proposal distribution. `forward_local` preserves egocentric continuity, `target_bearing_local` moves along the target ray, and `lateral_target_bypass` introduces side-step views; the optional challenger families test smaller corrections and reversals. Candidate-profile utility must be judged after pruning. A store in which target-aware families rarely survive cannot support a target-conditioned planning claim.

Pruning converts the full shell into a compact valid-action table. A row remains valid only if it lies in the snippet occupancy support, stays clear of the GT mesh, avoids straight-line path collision, and satisfies local egocentric motion limits:

$$\|\boldsymbol{o}_i\|_2 \leq 1.0 \text{ m}, \quad |\Delta h_i| \leq 0.25 \text{ m}, \quad \max\left(0, -o_{i,z}\right) \leq 0.25 \text{ m}, \quad \Delta\psi_i \leq 70 \text{ deg}$$

The full shell is retained with position, strategy, mixture, sampling probability, rule masks, diagnostics, and invalid-reason bitsets. Panel B of Figure 3 makes the distinction concrete for one stored table: invalid rows remain inspectable but sit outside the admissible set. Invalid candidates cannot enter Q_H selection, stochastic normalization, or loss targets. Conversely, a feasible row with weak target support or low expected gain remains a valid low-utility example rather than receiving an invalid-reason code. A manifest-defined root-support threshold may reject an entire rollout task when too few actions remain:

$$N_{\text{valid}} \geq \max\left(12, \text{ceil}\left(0.25 N_q\right)\right)$$

This threshold is a data-support guard, not a reward threshold. Preflight reporting must retain rejected-root counts and failure reasons by candidate family.

3.2.4 Rollout Branch Sampling and Dataset Impact

Rollout recipes select and retain finite chains from the valid action table. The implemented families are uniform valid sampling, one-step oracle greedy selection, bounded oracle lookahead, and temperature-softmax sampling. Their horizon, branch factor, beam width, temperature, and seed are resolved parameters. These recipes generate replay diversity and bounded references; they do not constitute a learned policy.

Bounded oracle lookahead can select a different first action from one-step greedy because it ranks retained finite-horizon chains rather than immediate gain alone (Figure 4). The persisted artifact contains these selected or beam-retained chains and their full per-step candidate shells; it is not an exhaustive materialization of the counterfactual action tree.

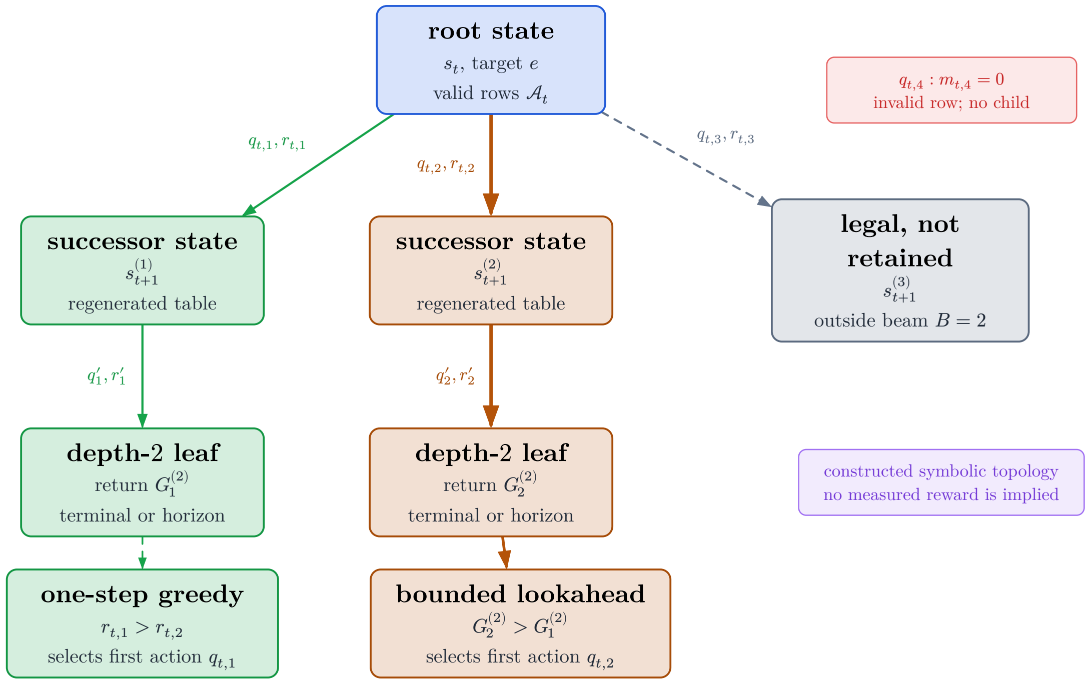

Figure 4: Constructed symbolic topology of the bounded oracle-lookahead reference. Nodes are counterfactual states and edge labels are candidate actions with immediate rewards. Only valid first-action rows produce children; the beam retains a bounded prefix set, and invalid rows receive no branch. The inequalities illustrate how the selected first action can differ from one-step greedy without inventing measured reward values.

For stochastic branches, let s_i be the finite oracle score of valid row i. The robust logit used for temperature-softmax is

$$\ell_i = \frac{s_i - \mathrm{median}(s)}{\mathrm{IQR}(s)\tau}, \quad P(i \mid m_i = 1) = \frac{\exp(\ell_i)}{\sum_{j:m_j=1}} \exp\big(\ell_j\big)$$

The target source defines the supervised task, the candidate mixture defines learnable action support, validity defines admissibility, and the recipe defines which chains enter replay. Reporting must therefore cover target-task coverage, valid fanout and invalid reasons by family, selected-family diversity, gain distributions, and retention cost before attributing policy behavior to non-myopic planning. The paired train-only pilots are bandwidth and candidate-profile probes; they are non-confirmatory and provide no held-out policy-performance evidence.

3.2.5 Target-Specific RRI

Let $C_e(\mathcal{P}_t)$ denote the oracle-only crop of accumulated evaluation points to the selected target region. The target error adapts the VIN-NBV objective to this crop [Fra+25]: point-to-mesh accuracy plus mesh-to-point completeness.

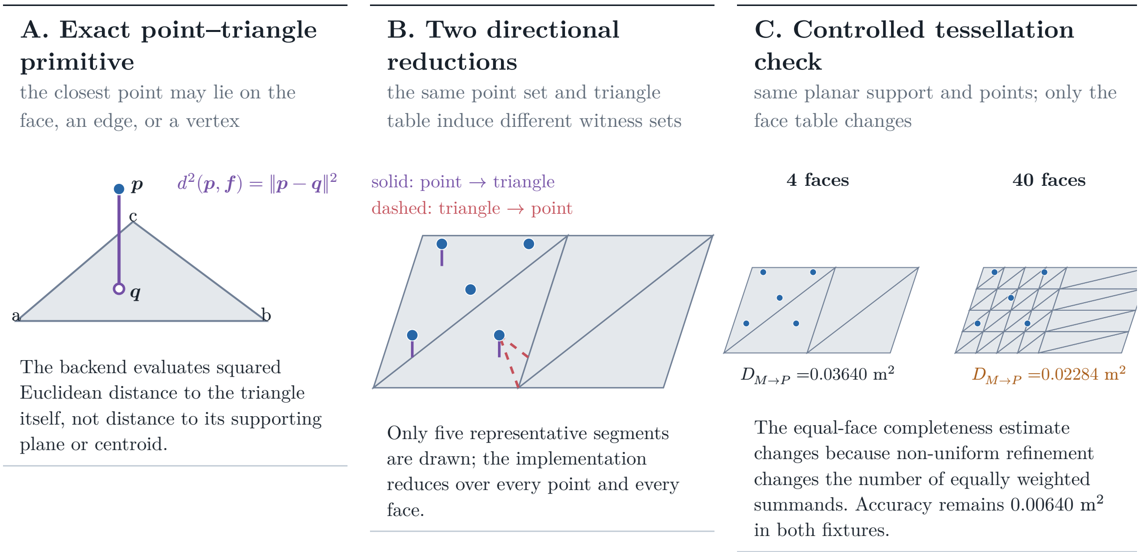

Controlled validity fixture computed by the repository's PyTorch3D metric and exact Trimesh closest-point witnesses; it is not an ASE performance result.

Figure 5: Point–mesh metric contract and controlled validity fixture. Panel A isolates the exact point-to-triangle primitive. Panel B distinguishes the point-to-face accuracy and face-to-point completeness reductions using computed closest-point witnesses. Panel C holds the planar support and point set fixed while changing only the triangle table; the measured equal-face completeness term changes from 0.03640 to 0.02284 m^2. These values are generated by the repository's PyTorch3D metric on a synthetic fixture and demonstrate a tessellation-sensitivity mechanism; they are not an ASE performance result.

$$\Delta_t^e = d\big(C_e(\mathcal{P}_t), \mathcal{M}_e^{\mathrm{GT}}\big) = D_{P\to M,t}^e + D_{M\to P,t}^e$$

The implementation crops mesh faces when any vertex lies inside the oriented target OBB and evaluates the configured point–mesh scorer. Geometry-invalid candidates are removed by the hard action mask and retain persisted candidate reason codes. Empty mesh crops, insufficient current support, or unusable renders instead invalidate the separate oracle evaluation: the recipe either skips the affected row or table, or clears `oracle_label_mask` and `q_train_mask`, while reporting the oracle failure independently rather than assigning a candidate reason code. The controlled fixture in Figure 5 demonstrates that the current equal-face completeness reduction is not invariant to non-uniform retessellation, even when the planar support and point set are fixed. Its sensitivity to root and candidate point sampling remains unmeasured. Confirmatory use therefore requires the sensitivity tests and metric-validity gate specified in Chapter 5; freezing one mesh does not by itself establish a representation-independent oracle.

The immediate training reward adapts VIN-NBV's reconstruction-improvement idea to a target crop and normalizes by the root target error rather than the current error [Fra+25]. This makes equal-horizon rollouts additive against a common root baseline:

$$r_t^e = \frac{\Delta_t^e - \Delta_{t+1}^e}{\Delta_0^e + \varepsilon}$$

The finite-horizon return is the discounted sum of those target rewards along a selected counterfactual branch. It is a training target for Q_H, not a claim that the deployed system has an online continuous-control policy:

$$G_{t,e}^{(h)} = \sum_{k=0}^{\min(h,b_t)-1} \gamma^k r_{t+k}^e$$

Endpoint gain is the primary fixed-budget comparison metric because it measures the target quality after the same number of acquisitions for each policy:

$$J_e^{(H)} = \frac{\Delta_0^e - \Delta_H^e}{\Delta_0^e + \varepsilon}$$

The log-gain variant is retained only as an internal scale-sensitivity ablation:

$$J_{e,\log}^{(H)} = \log(\Delta_0^e + \varepsilon) - \log(\Delta_H^e + \varepsilon)$$

$J_e^{(H)}$ is the fixed-budget evaluation metric, $G_t^{(H)}$ is the learning return, and $J_{e,\log}^{(H)}$ is a sensitivity diagnostic. Root-normalized gains telescope to endpoint gain when the discount is unity and geometry, horizon, and acquisition budget are matched. State-relative one-step RRI remains a VIN-compatible diagnostic. The resolved manifest must freeze the discount, clipping, target cap, crop policy, and all evaluation-geometry parameters for each reported experiment.

3.3 Replay Stores and Diagnostic Evidence

The pipeline materializes two stores with different ownership. Immutable `vin_offline/` caches logged snippet evidence and expensive one-step oracle products. Standalone `rollouts.zarr/` references those source rows and stores target tasks, retained rollout chains, per-step finite candidate shells, masks, reason codes, selected-action successor evidence, and a derived Q_H view. This separation prevents counterfactual experiments from mutating the source substrate and prevents one-step candidate labels from being mistaken for multi-step replay.

The source store is immutable because it defines the logged actor substrate. Its manifest records source identity, split, schema, materialized blocks, and provenance. A source row may contain frozen EVL/VIN fields, one-step candidate poses and labels, compact rendering payloads, and semidense geometry/history. These fields support scorer training and rollout construction, but they do not themselves define target-conditioned finite-horizon data. Exact directory names, array groups, chunking, and codec choices are versioned reproducibility metadata rather than scientific entities, so they are not reproduced as directory-tree figures in the main text.

The rollout store is normalized around replay identity. One source can produce multiple target rows, each target can produce multiple recipe and retained-chain rows, each chain expands into selected steps, and each step owns a full candidate shell. The store therefore captures selected or beam-retained chain evidence, not the complete counterfactual search tree. Padded `q_h/` arrays are a validated cache over the factual row tables. Invalid candidate rows remain inspectable, but action, training, and bootstrap masks exclude them.

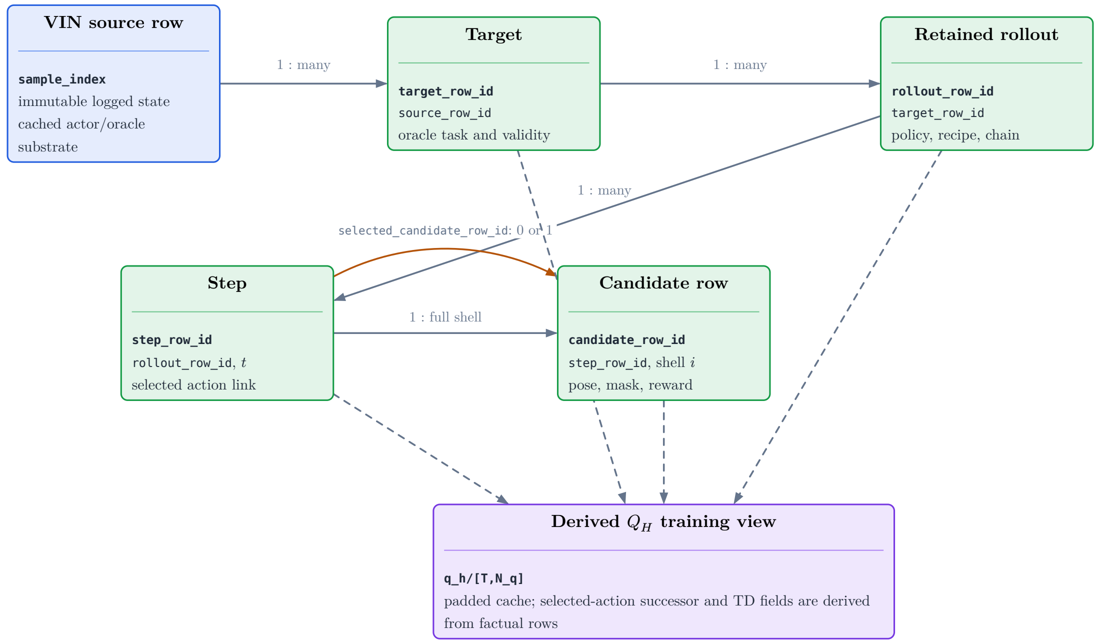

Figure 6: Normalized lineage of the implemented replay evidence. An immutable VIN source row may define several oracle target tasks; each target may produce several retained policy chains; each chain contains ordered steps; and each step owns one full candidate shell. The step row may identify one chosen candidate through `selected_candidate_row_id`; selected-action successor and TD fields are then constructed in the derived `q_h/` join rather than persisted as a separate transition table.

Evidence family	Interpretation contract
Sources and targets	Manifest-backed task coverage; not proof of actor-visible target discovery.
Candidates and invalidity	Full-shell support with hard action, training, and bootstrap masks.
Retained chains and steps	Recipe-selected evidence; not a persisted exhaustive search tree.
Selected depth	Chosen-action successor history; not all-candidate or endpoint evidence.
Q_H view	Derived training cache whose rewards and masks must agree with factual rows.

Table 2: Interpretation contract for rollout-store audits. Numeric values are rendered from the resolved report bundle in the experiment and reproducibility sections.

Selected-depth persistence stores only the depth raster and calibration for the chosen action at each retained step. It is sufficient to reconstruct selected history without duplicating dense all-candidate renders, but it is not an independently scored endpoint artifact. Likewise, rollout rows summarize final cumulative selected-chain metrics; they do not preserve every rejected branch or a policy-neutral endpoint reconstruction. These limitations must be resolved by matched endpoint re-evaluation before confirmatory policy comparison.

Scientific reporting is generated from the same inspection frames used by the diagnostic application. The main text needs target-task coverage, candidate validity and invalid reasons, family survival and selection, selected-history sanity, gain distributions, and runtime/storage summaries. Exact schema columns, compression, chunking, hashes, and cluster invocation belong in the reproducibility appendix and must be read from the resolved manifest and report bundle. Development bandwidth pilots are train-only feasibility checks; their counts and throughput may size later jobs but cannot support held-out reconstruction or policy claims.

4 Method

ARIA-NBV is formulated as target-conditioned selection from a finite candidate table. The implemented substrate generates and evaluates masked multi-step pose rollouts, records selected transitions, and exposes a dense training view for a planned finite-horizon value model. The current learned control is a myopic VIN scorer. Consequently, this chapter separates implemented data and replay contracts from the planned target-conditioned candidate-to-state Transformer; it does not report an untrained Q_H model as a completed method.

The chapter first identifies the actor-visible state and the coverage limitation of local EFM3D evidence. It then specifies the persisted and planned descriptors, the finite-action and replay semantics, geometric acceptance tests, a ranked representation and architecture design space, and the canonical bounded Q_H model. Descriptive status blocks distinguish implementation maturity from evidence maturity throughout.

4.1 Actor State and Representation Boundary

4.1.1 Implemented carrier and information boundary

Implementation: implemented **Evidence: pending**

The current carrier, replay masks, and target-task provenance are implemented and covered by schema tests. Frozen scientific validation remains pending. This status does not imply that the planned target-conditioned finite-horizon scorer is implemented.

Literature: [Str+24]

Promotion gate: retain actor/oracle provenance checks in every reader

Development source: aria_nbv/aria_nbv/targets/descriptor.py; aria_nbv/aria_nbv/vin/candidate_scorer.py; aria_nbv/aria_nbv/rollouts/replay/state.py; aria_nbv/tests/rollouts/test_zarr_store.py

The implemented one-step scorer consumes an actor-visible snippet view, a row-aligned candidate table, a reference rig pose, calibrated candidate cameras, and optionally cached EVL output. Oracle RRI labels enter the loss after prediction and are not scorer inputs. The V0 target descriptor contains semantic identity, pose, positive metric extents, and reference-relative pose in an actor-safe shape, but its current values are derived from privileged GT target tasks. A deployable claim therefore still requires observed or predicted target fields with explicit provenance.

The replay transition stores the candidate result, selected row, selection policy, root-to-selected pose chain, remaining budget, and successor-table linkage. Candidate regeneration changes pose, history, budget, and the finite action table. It does not update a learned persistent scene representation after every selected action. Stored selected-depth maps are privileged GT-mesh counterfactual evidence unless a named experiment establishes a deployable observation source.

Three boundaries are invariant. Invalidity is a hard mask with versioned reason codes, not a small value. GT meshes, target crops, all-candidate renders, associations, and returns remain oracle or audit fields. Candidate rows remain tied to stable shell identities and documented coordinate frames.

4.1.2 Task-sufficient actor state

Implementation: planned **Evidence: pending**

The thesis adopts a representation-independent readout contract. Particular scene carriers are promoted only when they preserve the information boundary and improve target-conditioned decisions.

Literature: [Bro+21, Zah+17]

Promotion gate: reader implementation, source-dropout tests, and held-out target-RRI ranking

Development source: docs/contents/theory/efm3d_scene_embeddings.qmd; docs/contents/theory/candidate_view_dependence.qmd

ARIA-NBV does not require one universal reconstruction format. It requires an actor-visible state from which a model can compare a target e and candidate $q_{t,i}$ at decision step t and requested horizon H:

$$\boldsymbol{h}_{t,e,i} = \mathrm{Read}\big(\boldsymbol{\Phi}_t^{\mathrm{scene}}, \boldsymbol{h}_e^{\mathrm{tgt}}, q_{t,i}, \boldsymbol{H}_t, t, H\big)$$

Here $\boldsymbol{\Phi}_t^{\mathrm{scene}}$ denotes persistent scene evidence, $\boldsymbol{h}_e^{\mathrm{tgt}}$ identifies the task, and $\boldsymbol{H}_t$ records selected-view history. A useful representation must preserve distinctions that can change target-specific return, distinguish missing evidence from predicted free space, and admit a causal update after selection. This requirement is narrower than a globally dense or photorealistic reconstruction and stronger than a pooled scene vector with no spatial support or provenance.

The desired memory separates observed surface, observed free space, unknown space, uncertainty, support count, recency, and directional history. Candidate queries should read the same memory in target-local, candidate-local, and ray-relative coordinates. A common world translation or yaw convention must not change the physical candidate values, while gravity, scale, target extent, camera direction, occlusion, and temporal order remain observable task variables. Exact global SE(3) invariance is therefore neither required nor claimed.

4.1.3 Local EFM3D evidence and the coverage gap

Implementation: partial **Evidence: pending**

EFM3D fields, their voxel pose, and finite extent are persisted. A frozen target/candidate coverage artifact and their sufficiency as the only long-horizon scene representation remain pending.

Literature: [Met25b, Str+24]

Promotion gate: report target and candidate support against the persisted voxel pose and extent

Development source: docs/literature/tex-src/arXiv-EFM3D/method.tex, Sec. Egocentric Voxel Lifting, lines 2–44; docs/contents/literature/efm3d.qmd; docs/contents/theory/efm3d_scene_embeddings.qmd; aria_nbv/aria_nbv/data_handling/offline/writer.py

EFM3D encodes logged RGB frames before lifting selected evidence into a finite gravity-aligned voxel field anchored to a snippet pose. The stored voxel transform and extent define the support of lifted features and dense heads. They must not be interpreted as a global volume containing every region observed by the trajectory. A target or candidate outside this support can remain physically valid; it has missing local-EVL evidence rather than an ordinary zero feature or automatically invalid action.

The coverage limitation motivates a layered interface rather than repeated inference at hypothetical candidate poses. Local EVL features provide high-quality Aria-native evidence where supported. Broader semidense or fused point carriers can preserve observed surfaces beyond that volume. Sparse ray-aware state can distinguish occupied, free, and unknown regions and can be updated from selected geometry. Logged appearance descriptors can be attached to points only after visibility and source lineage are established. None of these carriers creates RGB, DINO, detector, or EVL evidence at an unvisited counterfactual pose.

4.1.4 Ranked representation design space

Implementation: exploratory **Evidence: pending**

The table is a ranked design space, not a claim that every carrier will be implemented. The canonical model consumes a typed scene-memory interface so representations can be exchanged without redefining candidate or target semantics.

Literature: [Ave+24, Che+24, Fra+25, Jeo+26]

Promotion gate: promote only after matched support, leakage, runtime, and oracle-policy ablations

Development source: docs/literature/tex-src/arXiv-GenNBV/3-Method.tex, Secs. Formulation and Generalizable State Embedding, lines 10–49; docs/literature/tex-src/arXiv-Hestia/sec/3_method.tex, Sec. Methods, lines 14–58; docs/literature/tex-src/arXiv-scene-script/sections/structured_scene_language.tex; docs/contents/literature/efm3d.qmd

Carrier	Status	Inductive benefit	Cost and promotion gate
Persisted EVL and snippet evidence	active baseline	Aria-native local features, OBB hypotheses, calibrated cameras, and explicit extent.	Limited spatial support; retain only with coverage metadata.
Semidense or fused point memory	first broad-memory control	Observed surfaces follow trajectory extent and support target/frustum pooling.	No free/unknown state; test density and visibility sensitivity.
Sparse ray-aware memory	planned primary extension	Separates surface, free, unknown, support, uncertainty, and causal updates.	Requires deterministic fusion and counterfactual-source masks.
Target-centred EVL re-lifting	adaptation ablation	Reuses logged frame features when the target lies outside the root field.	Domain shift in the 3D neck; compare against simple logged-feature pooling.
DINO-on-point	appearance ablation	Extends logged appearance to observed points beyond the EVL grid.	Needs visibility gating, compression, and missing-descriptor masks.
TSDF/SDF or sparse encoder	geometry/encoder ablation	Compact metric geometry or learned coordinate-bearing tokens.	Must preserve observation weights and unknown-space semantics.
Object-aware 3DGS	renderable-memory ablation	Candidate rendering, soft target membership, and primitive uncertainty.	Per-scene optimization and mask supervision change the state contract.
SceneScript	global-context control	Broad ASE-aligned layout and object hypotheses from semidense input.	Not an ATEK drop-in and does not provide causal free/unknown updates.

Table 3: Ranked scene-representation design space. Status describes its role in the thesis method, not empirical superiority.

Representation comparisons begin with the persisted carrier and explicit support summaries. Learned point, sparse, equivariant, or renderable encoders are introduced only when a simpler carrier demonstrably limits target-RRI ranking or finite-horizon recovery. This ordering keeps representation capacity separate from target identity, candidate support, and replay validity.

4.2 Descriptor and Encoding Protocol

4.2.1 Persisted replay interface

Implementation: implemented **Evidence: pending**

The factual replay schema and dense `q_h/` view are implemented and unit-tested. Frozen scientific replay evidence remains pending. The schema is a storage contract; readability from the store does not make a field actor-visible.

Promotion gate: preserve row identity, masks, and provenance in every tensor reader

Development source: aria_nbv/aria_nbv/rollouts/zarr_store.py; aria_nbv/aria_nbv/targets/descriptor.py; aria_nbv/tests/rollouts/test_zarr_store.py

The rollout store preserves source, target, rollout, step, candidate, diagnostic, and lineage tables. Its derived `q_h/` arrays provide a padded state–candidate view without changing factual identities or labels. Target pose and extents in the V0 task remain privileged GT-derived instructions. Candidate geometry, selected history, remaining budget, and hard masks are decision inputs; target gains, GT associations, mesh diagnostics, crops, and candidate renders remain label or audit fields.

Carrier	Persisted content	Learning role
Target	identity, class, pose, extents, reference-relative pose, source and validity provenance	privileged V0 instruction; learned actors need observed or predicted equivalents
Candidate	stable row identities, world/root-relative pose, masks, reasons, sampler provenance, support fields	finite action row; privileged diagnostics remain source-gated
Selected chain	selected row, shell index, step order, policy, seed, successor link, terminal state	history and temporal-difference linkage
Oracle labels	target RRI, target root gain, errors, optional crops and selected depth	loss, evaluation, and audit only

Table 4: Implemented replay carriers and admissible learning roles.

4.2.2 Planned model-input contract

Implementation: planned **Evidence: pending**

The canonical reader will construct typed target, scene, history, budget, time, horizon, candidate, relation, mask, and provenance tensors. It will not expose oracle labels to the actor graph.

Literature: [Bro+21, Zho+23]

Promotion gate: dedicated `q_h/` reader, positive-width target path, and leakage tests

Development source: docs/literature/tex-src/arXiv-QCNet/main.tex, Sec. Query-Centric Scene Encoder, lines 159–161; aria_nbv/aria_nbv/vin/models/target_finite_horizon.py; docs/contents/theory/candidate_view_dependence.qmd

The intended input for target e at step t is

$$\mathcal{I}_{t,e} = \left(\boldsymbol{h}_e^{\text{tgt}}, \boldsymbol{\Phi}_t^{\text{scene}}, \boldsymbol{H}_t, \boldsymbol{b}_t, t, H, \left\{\boldsymbol{x}_{t,i}, \boldsymbol{e}_{a|i}^{\text{rel}}, m_{t,i}, \boldsymbol{\rho}_{t,i}\right\}_{i=1}^{N_q}\right)$$

The current candidate shell is shared across targets. Multiple target tokens can therefore be evaluated in parallel by adding a target axis and masking unavailable target–candidate pairs; candidate and scene encodings that are target-independent are computed once. The requested horizon and current step remain explicit because otherwise identical geometry can have different value near termination. A horizon mask prevents a short-horizon query from attending to unavailable future transitions.

The target descriptor separates identity, geometry, observed support, confidence, and source:

$$\boldsymbol{\varphi}_e = \text{Enc}_{\text{tgt}}\left(\hat{\boldsymbol{B}}_e, \hat{\boldsymbol{y}}_e, \hat{\pi}_e, A_e^{\text{proj}}, n_e^{\text{semi}}, n_e^{\text{EVL}}, \omega_e^{\text{EVL}}, \ell_e^{\text{src}}, \boldsymbol{T}_{r_t,e}, \boldsymbol{T}_{c_t,e}\right)$$

The learned token combines that descriptor with target-local scene support:

$$\boldsymbol{h}_e^{\text{tgt}} = \text{MLP}_{\text{tgt}}(\text{concat}(\boldsymbol{\varphi}_e, \boldsymbol{g}_e^{\text{tgt}}))$$

In the current V0 data, target geometry is GT-derived. The same tensor interface becomes actor-admissible only when populated by an observed or predicted target track. Identity IoU and ambiguity-gap equations are deliberately not part of the present Method because the corresponding actor-visible matcher is not implemented.

4.2.3 Relative candidate geometry

Implementation: partial **Evidence: pending**

World and root-relative poses are persisted and candidate-query pooling exists. The complete target-relative finite-horizon descriptor is planned.

Literature: [Li+21, Zho+19, Zho+23]

Promotion gate: frame-transform, row-shuffle, and held-out descriptor ablations

Development source: docs/literature/tex-src/arXiv-QCNet/main.tex, Sec. Query-Centric Scene Encoder, lines 159–161; aria_nbv/aria_nbv/data_handling/offline/batch.py; aria_nbv/aria_nbv/vin/modules/pooling.py

Canonical rigid transforms remain in the store. The reader derives a candidate pose relative to the current decision reference rather than learning the arbitrary world origin:

$$\boldsymbol{T}^{\text{rel}}_{r_t,i} = \boldsymbol{T}^{-1}_{w,r_t}\boldsymbol{T}_{w,c_{t,i}}, \quad \boldsymbol{\delta}^{p}_{r_t,i} = \boldsymbol{R}^{\top}_{r_t}\left(\boldsymbol{c}_{t,i} - \boldsymbol{c}_{r_t}\right), \quad \boldsymbol{R}_{r_t,i} = \boldsymbol{R}^{\top}_{r_t}\boldsymbol{R}_{t,i}$$

$$\boldsymbol{h}^{\text{pose}}_{t,i}\left(q_{t,i}; r_t\right) = \operatorname{concat}\left(\boldsymbol{\delta}^{p}_{r_t,i}, \operatorname{R6D}\left(\boldsymbol{R}_{r_t,i}\right), \left\|\boldsymbol{\delta}^{p}_{r_t,i}\right\|_2, \operatorname{atan2}\left(\delta^{y}_{r_t,i}, \delta^{x}_{r_t,i}\right), \Delta h_{t,i}, \boldsymbol{u}^{\text{up/frustum}}_{t,i}\right)$$

The target relation is encoded separately so candidate self-motion cannot be confused with target conditioning:

$$\boldsymbol{h}^{\text{rel}}_{t,e|i}\left(q_{t,i}, e\right) = \operatorname{concat}\left(\boldsymbol{\delta}^{p}_{e|i}, \left\|\boldsymbol{\delta}^{p}_{e|i}\right\|_2, \cos\theta^{\text{opt}}_{t,e,i}, \beta^{\text{elev}}_{t,e,i}, \lambda^{\text{obb}}_{t,e,i}\right)$$

Continuous rotation features, metric range, bearing, elevation, height, and frustum variables preserve physical task structure. Query-local relative positional encoding is a later ablation for candidate–target, candidate–history, or candidate–candidate interactions:

$$\boldsymbol{\delta}^{\text{p}}_{a|i} = \boldsymbol{R}^{\top}_{t,i}\left(\boldsymbol{p}_a - \boldsymbol{c}_{t,i}\right), \quad \boldsymbol{\delta}^{\text{R}}_{a|i} = \boldsymbol{R}^{\top}_{t,i}\boldsymbol{R}_a$$

$$\boldsymbol{\eta}_{a|i} = \operatorname{concat}\left(\boldsymbol{\delta}^{\text{p}}_{a|i}, \left\|\boldsymbol{\delta}^{\text{p}}_{a|i}\right\|_2, \operatorname{enc}_R\left(\boldsymbol{\delta}^{\text{R}}_{a|i}\right)\right), \quad \boldsymbol{e}^{\text{rel}}_{a|i} = \psi_{\text{rel}}\left(\mathcal{F}\left(\boldsymbol{\eta}_{a|i}\right)\right), \quad a \in \{j, e, k\}$$

This adopts QCNet's query-centric geometric discipline, not its trajectory decoder or streaming claim [Zho+23].

4.2.4 Candidate rows and directional history

Implementation: planned **Evidence: pending**

The minimal planned directional feature pairs a signed target-local first moment with the second moment. This avoids identifying opposite approach directions. Spherical harmonics remain a richer optional ablation.

Literature: [e3n25, Lu+26]

Promotion gate: directional-novelty fixture and matched architecture ablation

Development source: docs/literature/tex-src/arXiv-Hestia/sec/3_method.tex, Sec. Methods, lines 44–58; docs/contents/theory/candidate_view_dependence.qmd

The row token keeps self pose, target relation, support, validity, provenance, selected history, step, horizon, and budget as typed blocks:

$$\boldsymbol{x}_{t,i} = \mathrm{concat}\Big(\boldsymbol{h}_{t,i}^{\text{pose+rel}}, \boldsymbol{h}_{t,i}^{\text{geom}}, \boldsymbol{h}_{t,i}^{\text{valid}}, \boldsymbol{h}_{t,i}^{\text{prov}}, \boldsymbol{H}_t, \mathrm{Emb}(t), \mathrm{Emb}(H), \boldsymbol{b}_t\Big)$$

Viewing history is not reducible to camera distance. For a target-local point or cell, selected camera directions define a signed first moment together with the second-moment memory

$$\boldsymbol{\mu}_t^{\text{dir}}(\boldsymbol{v}) = \frac{\sum_{k<t} w_k(\boldsymbol{v})\boldsymbol{d}_k(\boldsymbol{v})}{\sum_{k<t} w_k(\boldsymbol{v}) + \varepsilon}, \quad \boldsymbol{M}^{\text{dir}}(\boldsymbol{v}) = \frac{\sum_{k<t} w_k(\boldsymbol{v})\boldsymbol{d}_k(\boldsymbol{v})(\boldsymbol{d}_k(\boldsymbol{v}))^\top}{\sum_{k<t} w_k(\boldsymbol{v}) + \varepsilon}$$

The signed first moment distinguishes opposite approach directions, whereas the second moment records concentration along viewing axes. Their query-relative projections form directional features without committing to a full spherical-harmonic field. Low-order spherical harmonics are promoted only if these moments improve over pose history and still leave systematic directional errors.

The first reader may omit optional appearance, ray, and directional blocks through explicit missing-modality masks. Missing support is never encoded as an ordinary all-zero observation, and sampler family remains provenance rather than a geometric shortcut.

4.3 Finite Candidate and Replay Contract

Implementation: implemented **Evidence: pending**

Finite candidate tables, hard masks, lineage, selected transitions, and the derived `q_h/` view are implemented and schema-tested. Frozen scientific store evidence remains pending.

Promotion gate: preserve deterministic shell identity and selected-transition validation

Development source: aria_nbv/aria_nbv/rollouts/replay/types.py; aria_nbv/aria_nbv/rollouts/replay/engine.py; aria_nbv/aria_nbv/rollouts/zarr_store.py; aria_nbv/tests/rollouts/test_zarr_store.py

At step t, candidate generation returns a finite full-shell table $\mathcal{Q}_t$ with a hard-valid mask $\boldsymbol{m}_t$ and versioned invalid-reason bitsets. Scores are stored compactly only for hard-valid rows and are bound back to stable shell indices before selection. The admissible action set is

$$\mathcal{Q}_t = \{q_{t,i}\}_{i=1}^{N_q}, \quad \mathcal{A}_t = \{i \in \{1, ..., N_q\} : m_{t,i} = 1\}$$

Invalid rows remain available for diagnostics and dense replay, but cannot be selected. A row enters the training mask only when it is actor-selectable and has a finite oracle target. Invalid rows have false masks and undefined labels; scene RRI is never substituted for a missing target-specific label.

4.3.1 Implemented replay transition

Rollout expansion records the full candidate table, selected valid and shell indices, policy scores and probabilities, selection policy, and random seed. The implemented transition is

$$(x_{t+1}, \boldsymbol{H}_{t+1}, b_{t+1}, \mathcal{Q}_{t+1}) = \mathrm{Step}\left(x_t, \boldsymbol{H}_t, b_t, q_{t,a_t}, \xi_t\right)$$

where x_t is the current reference pose, $\boldsymbol{H}_t$ the selected-pose history, b_t the remaining budget, and ξ_t the deterministic generation context. The next candidate table is regenerated around the selected pose under the same target task, history constraints, and versioned generator configuration. This transition changes pose, history, budget, lineage, and action support. It does not imply that the actor has received a new RGB observation or recomputed EFM3D field.

The persisted factual tables retain source and target identity, lineage hashes, step order, selected rows, masks, reasons, sampler provenance, rewards, and support diagnostics. The derived `q_h/` view right-pads states, stores one-step and target-root-gain labels, and exposes only the factual selected transition needed for a temporal-difference backup. Readers can rebuild this view at a different discount without mutating factual replay.

4.3.2 Planned scene-state transition

Implementation: planned **Evidence: pending**

A deployable successor-state encoder must update only evidence produced by the selected observation. Current GT-mesh selected depth is privileged counterfactual evidence.

Literature: [Che+24, Lu+26]

Promotion gate: deterministic fusion, source masks, and no-future-observation tests

Development source: docs/contents/theory/efm3d_scene_embeddings.qmd; aria_nbv/aria_nbv/oracle/evidence.py

The existing geometry-level counterfactual uses a set union of retained points,

$$\mathcal{P}_{t+1} = \mathcal{P}_t \cup \mathcal{P}_{q_t}$$

but a planning memory must additionally preserve whether space is observed occupied, observed free, or still unknown. The proposed sparse update is

$$\boldsymbol{M}^{\text{ray}}_{t+1} = \text{Fuse}\big(\boldsymbol{M}^{\text{ray}}_{t}, \boldsymbol{P}^{\text{selected-obs}}_{t+1}(a_t), \mathcal{R}^{\text{selected-obs}}_{t+1}(a_t)\big), \quad a_t \in \mathcal{A}_t$$

Here both evidence terms are indexed by the selected action a_t; no unselected candidate render enters the successor state. The source may be a deployable sensor observation or an explicitly privileged selected-depth counterfactual, and that role is stored. The update may add selected surface evidence, carve observed free space, update support and uncertainty, and refresh target-local directional memory. It must not attach RGB, DINO, detector, or EVL descriptors to counterfactual geometry unless a corresponding actor-visible observation exists. Counterfactual-only cells instead carry explicit source and missing-modality masks.

This split prevents two common errors: treating pose-only rollout as complete visual simulation, and treating privileged selected-depth geometry as a deployable sensor stream. The same replay schema supports both the current oracle experiment and later causal scene-memory ablations because source role is part of the state contract.

4.4 Geometric and Mask Acceptance Tests

Implementation: partial **Evidence: pending**

Replay and current VIN data surfaces cover part of the contract. The unimplemented finite-horizon model must pass the complete suite before architectural comparisons are interpretable.

Literature: [Bro+21, Lee+19, Zah+17]

Promotion gate: end-to-end finite-horizon permutation, mask, duplicate, and frame tests

Development source: aria_nbv/aria_nbv/vin/types/prediction.py; aria_nbv/aria_nbv/vin/modules/heads.py; aria_nbv/tests/lightning/test_vin_batch_collate.py; aria_nbv/tests/rollouts/test_zarr_store.py

Candidate order carries no task meaning. For a per-candidate scorer f_θ, jointly permuting row-aligned inputs by Π must permute outputs by the same amount:

$$f_\theta(\Pi \boldsymbol{X}_t, \Pi \boldsymbol{m}_t) = \Pi f_\theta(\boldsymbol{X}_t, \boldsymbol{m}_t)$$

The current VIN path uses row-wise maps, shared-token queries, and symmetric normalization without candidate-index embeddings. Existing tests verify row-aligned collation and metric alignment, not an end-to-end finite-horizon equivariance claim.

Equivariance alone does not guarantee invalid-row isolation. Holding valid rows and the mask fixed while changing only invalid-row contents must satisfy

$$\mathrm{Mask}(\boldsymbol{X}_t, \boldsymbol{m}_t) = \mathrm{Mask}(\boldsymbol{X}'_t, \boldsymbol{m}_t) \quad \Rightarrow \quad \mathrm{Mask}(f_\theta(\boldsymbol{X}_t, \boldsymbol{m}_t), \boldsymbol{m}_t) = \mathrm{Mask}(f_\theta(\boldsymbol{X}'_t, \boldsymbol{m}_t), \boldsymbol{m}_t)$$

The valid-action mask gates selection, softmax, and successor maximization. The narrower training mask also requires a finite oracle target. Padding, invalid rows, and their undefined labels must not alter valid outputs or gradients. Duplicate-row and valid-count tests are required because attention normalization or per-set centering can otherwise change the absolute value of an unchanged physical candidate.

Coordinate handling is deliberately weaker than exact SE(3) equivariance. World poses remain available for reproducibility, while model inputs use root-, target-, or candidate-relative geometry. Gravity, scale, height, yaw, camera direction, target orientation, motion limits, and frustum geometry remain physical variables. A global origin convention must not become a shortcut [Bro+21].

Actor/oracle provenance is the final acceptance condition. Target gains, GT associations, mesh distances, rendered depth, and target crops may supervise or audit a model but may not enter its actor graph unless the experiment is explicitly privileged. Source-dropout tests must show that removing an unavailable optional carrier changes only its masked branch.

4.4.1 Conservative architecture ladder

Implementation: planned **Evidence: pending**

The ladder orders hypotheses; it does not assign equal status to every model family. Candidate-to-state attention is the canonical planned core.

Literature: [Bre+23, Fuc+20, Lee+19, SHW21, Zah+17, Zho+23]

Promotion gate: promote a level only after all lower controls pass the same oracle-evaluated protocol

Development source: docs/literature/tex-src/arXiv-Deep-Sets/nips_2017.tex; docs/literature/tex-src/arXiv-Set-Transformer/03_main.tex, Sec. Set Transformer, lines 2–51; docs/literature/tex-src/arXiv-QCNet/main.tex, Sec. Query-Centric Scene Encoder, lines 159–161; docs/literature/tex-src/arXiv-SE3-Transformer/EA4PC.tex, Sec. Introduction, lines 119–143; docs/contents/theory/candidate_view_dependence.qmd

Level	Model	Scientific role
A0	Independent masked MLP	Locks reader, target/candidate descriptors, masks, and one-step or finite-horizon label learnability.
A1	Candidate-to-state cross-attention	Canonical planned model: each candidate independently queries shared target, scene, history, budget, step, and horizon tokens.
A2	DeepSets context	Tests whether unordered summaries of valid candidates add information without pairwise attention.
A3	Masked Set Transformer	Tests candidate-candidate interaction while preserving row equivariance and mask isolation.
A4	Query-local relation bias	Tests QCNet-style target, history, and candidate relations without importing its forecasting decoder.
A5	Directional/support memory	Tests target-local view novelty, target-frustum support, overlap, and uncertainty as features.
A6	Uncentred residual Q_H	Tests finite-horizon recovery over a calibrated myopic control in continuous return units.
A7+	Point, sparse, recurrent, or exact-equivariant encoder	Escalates only if compact scene descriptors demonstrably bottleneck ranking or recovery.

Table 5: Architecture ladder for attributable geometric-learning claims.

Candidate-to-candidate attention is therefore not required by the core task. It may improve relative policy context or diversity, but unrelated sampled rows must not silently redefine the absolute value of candidate $q_{t,i}$. Exact equivariant layers are similarly scoped to support encoders or candidate graphs after local-frame scalar controls. This order favors reusable scene and target encodings: target-independent scene tokens are computed once, target tokens once per target, and candidate queries independently per candidate before optional set interaction.

4.5 Target-Conditioned Finite-Horizon Value Model

4.5.1 Implemented control and explicit scaffold

Implementation: partial **Evidence: pending**

The one-step VIN/CORAL scorer and finite-horizon configuration scaffold exist. Frozen matched-control evidence is pending. Positive-width target conditioning and the finite-horizon network are not implemented.

Literature: [CMR19, Fra+25]

Promotion gate: retain the current one-step scorer as the matched myopic control

Development source: aria_nbv/aria_nbv/vin/models/target_myopic.py; aria_nbv/aria_nbv/vin/models/target_finite_horizon.py; aria_nbv/aria_nbv/lightning/_candidate_scorer_contract.py

The current VINv3 scorer predicts per-candidate ordinal one-step RRI evidence from scene and candidate geometry without a learned target token. Its target-conditioned configuration accepts only target-descriptor width zero. The finite-horizon configuration records horizon, discount, and candidate-token width, but model construction raises `NotImplementedError`, and the current Lightning contract rejects the configuration because its CORAL objective does not consume rollout returns, hard action masks, or selected-transition links.

This boundary makes the implemented VIN path a necessary control rather than evidence for the proposed planner. Learned-policy claims require a dedicated `q_h/` reader, finite-horizon checkpoint, frozen split and configuration, and oracle re-evaluation of selected trajectories.

4.5.2 Canonical planned transformer

Implementation: planned **Evidence: pending**

The canonical model is a minimal target-conditioned candidate-to-state Transformer. Candidate-candidate interaction, exact equivariance, and richer scene encoders are ordered ablations rather than prerequisites.

Literature: [HGS15, Lee+19, Zah+17, Zho+23]

Promotion gate: A0 reader/MLP control, A1 implementation, invariant tests, and held-out oracle policy comparison

Development source: docs/literature/tex-src/arXiv-Set-Transformer/03_main.tex, Sec. Set Transformer, lines 2–51; docs/literature/tex-src/arXiv-QCNet/main.tex, Sec. Query-Centric Scene Encoder, lines 159–161; aria_nbv/aria_nbv/vin/models/target_finite_horizon.py; docs/contents/theory/candidate_view_dependence.qmd

For target e, the return over at most H selections is

$$G_{t,e}^{(h)} = \sum_{k=0}^{\min(h,b_t)-1} \gamma^k r_{t+k}^e$$

and the learned quantity is

$$Q_{h,e}(s_t, i) = \mathbb{E}\Big[G_{t,e}^{(h)} \mid s_t, a_t = i\Big], \quad i \in \mathcal{A}_t, \quad 1 \le h \le b_t, \quad Q_{0,e}(s,i) = 0$$

One candidate row is one query. Target, scene, selected history, remaining budget, current step, and requested horizon form shared state tokens. The canonical interaction is

$$\boldsymbol{u}_{t,i} = \mathrm{CrossAttn}_\theta\big(\boldsymbol{x}_{t,i}, \{\boldsymbol{h}_e^{\mathrm{tgt}}, \boldsymbol{\Phi}_t^{\mathrm{scene}}, \boldsymbol{H}_t, \mathrm{Emb}(t), \mathrm{Emb}(H), \boldsymbol{b}_t\}\big)$$

followed by a shared per-row value head. The valid-action mask is applied to attention where rows are keys, to selection, and to every bootstrap maximum. The training mask gates supervised losses. No candidate-index embedding is permitted. Multiple target tokens may share scene and candidate encodings and be evaluated in parallel under a target–candidate mask, which preserves the same network for single-target and multi-target inference.

Time and horizon conditioning prevent an otherwise identical candidate from receiving the same value at different remaining budgets. They do not reveal future observations: a state at step t contains only logged or selected evidence available through t. A requested horizon longer than the remaining rollout support is masked or terminated rather than padded with invented transitions.

The first implementation sequence is A0 followed by A1 from Table 5. A0 establishes whether the replay target is learnable without attention. A1 then tests whether reusable target and scene tokens explain additional value. Candidate-candidate attention is introduced only if fixed-state candidate queries leave a measurable error correlated with valid-set context.

4.5.3 One-step base and finite-horizon residual

Implementation: exploratory **Evidence: pending**

The residual decomposition is a testable hypothesis, not a consequence of object-centric NBV literature. Direct continuous Q_H remains the required control.

Literature: [CMR19, HGS15]

Promotion gate: target-specific label audit and direct-versus-residual ablation

Development source: docs/literature/tex-src/arXiv-VIN-NBV/sec/3_methods.tex, ordinal RRI paragraph, line 125; docs/contents/theory/candidate_view_dependence.qmd

The hypothesis separates calibrated immediate utility from downstream effects:

$$b_{\psi,i} = f_{\psi}^{\text{1-step}}\left(s_t^{\text{cf0}}, \boldsymbol{\varphi}_e, q_{t,i}\right), \quad \delta_{\theta,t,e,i}^{h} = g_{\theta}\left(\mathcal{I}_{t,e}, q_{t,i}, h\right), \quad Q_{h,\theta,e,i} = b_{\psi,i} + \delta_{\theta,t,e,i}^{h}$$

The base $b_{\psi,i}$ represents one-step target gain. The residual captures candidate regeneration, selected-history geometry, occlusion, support, and remaining horizon. It is not exactly mean-centred within each candidate table because duplicate or unrelated rows would then change absolute temporal-difference targets. Magnitude regularization may be tested without redefining the value field:

$$Q_{h,\theta,e,i} = \hat{r}_{\psi}^{e}\left(s_t^{\text{cf0}}, \boldsymbol{\varphi}_e, q_{t,i}\right) + \delta_{\theta,t,e,i}^{h}\left(\mathcal{I}_{t,e}, q_{t,i}\right), \quad \mathcal{L}_{\delta} = \lambda_{\delta} \frac{1}{|\mathcal{A}_t|} \sum_{j \in \mathcal{A}_t} \left(\delta_{\theta,t,e,j}^{h}\right)^2$$

CORAL remains a motivated one-step ranking and calibration interface,

$$\begin{aligned} p_{t,i,k}^{\text{CORAL}} &= \sigma\left(o_{t,i,k}^{\text{CORAL}}\right), \quad k = 0, \ldots, K-2 \\ \pi_{t,i,k}^{\text{CORAL}} &= p_{t,i,k-1}^{\text{CORAL}} - p_{t,i,k}^{\text{CORAL}}, \quad p_{t,i,-1}^{\text{CORAL}} = 1, \quad p_{t,i,K-1}^{\text{CORAL}} = 0 \\ \hat{r}_{\psi}^{e}\left(s_t^{\text{cf0}}, \boldsymbol{\varphi}_e, q_{t,i}\right) &= \sum_{k=0}^{K-1} \pi_{t,i,k}^{\text{CORAL}} u_k \end{aligned}$$

but additive finite-horizon returns are learned in continuous units. The direct continuous Q_H head, the residual head with a frozen base, and the residual head with slow or joint fine-tuning form the controlled comparison.

Research TODO: Compare direct continuous value prediction against the uncentred residual only after the target-specific label distribution, one-step calibration, and positive oracle-lookahead headroom are established.

Source: finite-horizon model contract

Gate: A0/A1 learning and headroom report

Open decision: Freeze, slow-fine-tune, and joint-training variants are all admissible until validation evidence selects one; the final thesis must record the chosen base-network update rule and checkpoint provenance.

Source: residual training hypothesis

Gate: model-selection protocol freeze

The scientific endpoint is not training loss. For every trained policy, selected trajectories are regenerated under the documented candidate contract and re-scored by the same target-specific oracle used for the baselines. The model succeeds only if it recovers a prespecified fraction of positive oracle-lookahead headroom on held-out scenes without violating mask, provenance, or support constraints.

5 Experimental Design

The experiments separate feasibility evidence from policy evidence. Feasibility requires a frozen scene-level population, reproducible target and candidate support, valid replay linkage, and measured oracle throughput. Policy inference additionally requires a held-out scene split, matched acquisition budgets, independent endpoint oracle evaluation, and an implemented actor-visible decision rule. The current train-only bandwidth pilots address only the first category: incomplete runs and resource failures characterize the generation system, but they cannot support comparisons between policies.

5.1 Study Population and Evidence Gates

The source population comprises ASE/ATEK snippet windows from scenes for which the configured GT mesh and object-box table resolve. A frozen manifest assigns entire scenes to train, validation, or test before model selection; no scene may cross these boundaries through another snippet. Each reported run records the manifest hash and the exact counts of scenes, snippets, admitted target tasks, rollout chains, transitions, and retained candidate rows. A capped train-only pilot is therefore a throughput and support probe, not a sample from which held-out policy performance can be estimated.

The present data generator defines oracle target tasks by seeded sampling from geometry-valid GT OBB rows. Its task-coverage report therefore describes the available GT pool, sampled tasks, classes, scenes, and later oracle-evaluation failures. It does not measure proposal matching, IoU ambiguity, projected visibility, or actor-observation support. Those quantities belong to a future observed-target selector required for deployable-input claims. Until that selector exists, GT target geometry may define labels and bounded oracle references but cannot be presented as actor-visible input. The privileged-supervision boundary in Figure 1 applies equally to render-derived evidence: dense GT candidate depth may produce labels, returns, or explicitly named privileged ablations, but it is not a legal actor input. A separate teacher policy, distillation path, or current-belief renderer remains a hypothesis until implemented and evaluated, so none is promoted to a main-text process figure.

The first policy gate is an actor-visible myopic scorer over the same finite candidate table intended for Q_H. The scorer must expose one value per candidate, respect the hard action mask, and be assessed by candidate ranking, calibration, and oracle-rescored selected actions. The existing scene-level VIN scorer is historical substrate; it is not a target-conditioned control until the observed target descriptor is wired into the model and evaluated on a frozen held-out split.

The second policy gate estimates whether bounded oracle lookahead has headroom over one-step oracle greedy:

$$\Delta_{\text{look}} = J_e^{(H)}(\pi_{\text{oracle-look}}) - J_e^{(H)}(\pi_{\text{oracle-1}})$$

Only if the preregistered analysis classifies this headroom as meaningful is Q_H evaluated for recovery from offline rollout traces:

$$\eta_Q = \frac{J_e^{(H)}(\pi_Q) - J_e^{(H)}(\pi_{\text{learned-1}})}{J_e^{(H)}(\pi_{\text{oracle-look}}) - J_e^{(H)}(\pi_{\text{learned-1}}) + \varepsilon}$$

Success is measured by matched endpoint oracle evaluation, not predicted values or training loss. If lookahead has no meaningful headroom, the result is scoped to the frozen split, target protocol, candidate generator, horizon, branch factor, and validity regime. If headroom exists but the learned model does not recover it, target observability, action support, replay coverage, reward construction, and model capacity remain separate candidate explanations.

Claim	Primary evidence	Decision rule
Population	scene-split manifest and coverage bundle	Inference is restricted to the frozen held-out scene population.
Task protocol	GT pool, sampled tasks, classes, and oracle failures	Current evidence is oracle-task coverage, not observed-target matching.
Myopic control	ranking, calibration, and oracle-rescored selections	Actor-visible target conditioning must be implemented before comparison.
Planning headroom	Δ_{look} and recovered fraction η_Q	Q_H is evaluated only after a meaningful-headroom gate.

Table 6: Objective-to-evidence matrix.

5.2 Replay Eligibility and Learning Gate

The factual rollout tables determine which records may supervise each learning problem. `valid_action_mask` defines actor-selectable actions. The stricter `q_train_mask` additionally requires a valid target/GT label state and finite target-root-gain and diagnostic target-RRI labels. Masked rows remain available for support and failure analysis but cannot enter action selection, supervised loss, or bootstrap maximization.

All eligible candidate rows can support one-step supervision. Finite-horizon temporal-difference supervision is narrower: the factual selected action must have a stored reward, terminal flag, and discount. A nonterminal transition additionally requires a valid successor step identifier whose state exposes a reproducible candidate table and hard mask; a terminal transition instead carries no successor id and has zero bootstrap discount. The derived `q_h/` arrays align these fields on a padded state–candidate view; they do not create labels for unobserved transitions. The replay pipeline in Figure 7 visualizes this distinction.

Data product	Purpose	Minimum evidence
target tasks	define supervised entities	GT pool, sampled tasks, class/scene coverage, and oracle failures
candidate shells	define finite action support	family counts, valid fraction, invalid reasons, and selected-family coverage
`q_train_mask`	admit one-step labels	actor-selectable rows with valid target state and finite target labels
selected transitions	admit TD linkage	reward, terminal flag, discount, and—when nonterminal—successor id and mask

Table 7: Required support evidence. Each narrower learning surface inherits the validity requirements above it.

Gate	Available evidence	Required before inference
Oracle data	target labels, masks, lineage, and replay validation	held-out coverage and matched endpoint evaluation
Myopic control	scene-level VIN substrate	actor-visible target-conditioned scorer and frozen checkpoint
Q_H	selected-transition tensors and TD linkage	implemented learner, frozen checkpoint, and oracle-rescored policy
Policy claim	train-only feasibility pilots	completed held-out paired comparison under equal budget

Table 8: Learning-readiness gates. Data-contract evidence does not substitute for an implemented and evaluated policy.

Implementation: planned **Evidence: pending**

The replay tensors required for a masked Double-Q learner exist; the learner and policy evidence do not.

Literature: [HGS15]

Promotion gate: finite-horizon reader, learner, checkpoint, and held-out oracle re-evaluation

Development source: aria_nbv/aria_nbv/rollouts/zarr_store.py; aria_nbv/aria_nbv/vin/models/target_finite_horizon.py

The intended finite-candidate value model decodes actions only over valid candidate tokens:

$$\boldsymbol{u}_{t,i} = \mathrm{Enc}_\theta(\mathcal{I}_{t,e}, q_{t,i})$$

$$a_t^\theta = \operatorname*{argmax}_{i:m_{t,i}=1} Q_{h,\theta,e}(s_t, i)$$

The planned first backup is fitted masked Double-Q [HGS15]. Here $d_t = 1$ at horizon termination, budget termination, or when no valid successor action exists. The boundary value $Q_{0,e} = 0$ makes a one-step target non-bootstrapping by definition:

$$B_t^{(h,e)} = \begin{cases} Q_{h-1,\theta^-,e}\big(s_{t+1}, \operatorname{argmax}_{i:m_{t+1,i}=1} Q_{h-1,\theta,e}(s_{t+1}, i)\big) & \text{if } \; h > 1 \\ d_t = 0 & \\ \sum_i m_{t+1,i} > 0 & \\ 0 & \text{otherwise} \end{cases}$$

$$y_t^{(h,e)} = r_t^e + \gamma B_t^{(h,e)}$$

For an implemented learner, the replay dataset $\mathcal{D}$ contains selected-action transition rows with state, action, immediate target reward, successor state, validity masks, and terminal flags:

$$\mathcal{L}_Q(\theta) = \frac{\sum_{(s_t,e,h,a_t,r_t^e,s_{t+1},\boldsymbol{m}_{t+1},d_t)\in\mathcal{D}} m_{t,a_t}^{\text{train}} \Big(Q_{h,\theta,e}(s_t, a_t) - y_t^{(h,e)}\Big)^2}{\sum_{(s_t,e,h,a_t,r_t^e,s_{t+1},\boldsymbol{m}_{t+1},d_t)\in\mathcal{D}} m_{t,a_t}^{\text{train}} + \varepsilon}$$

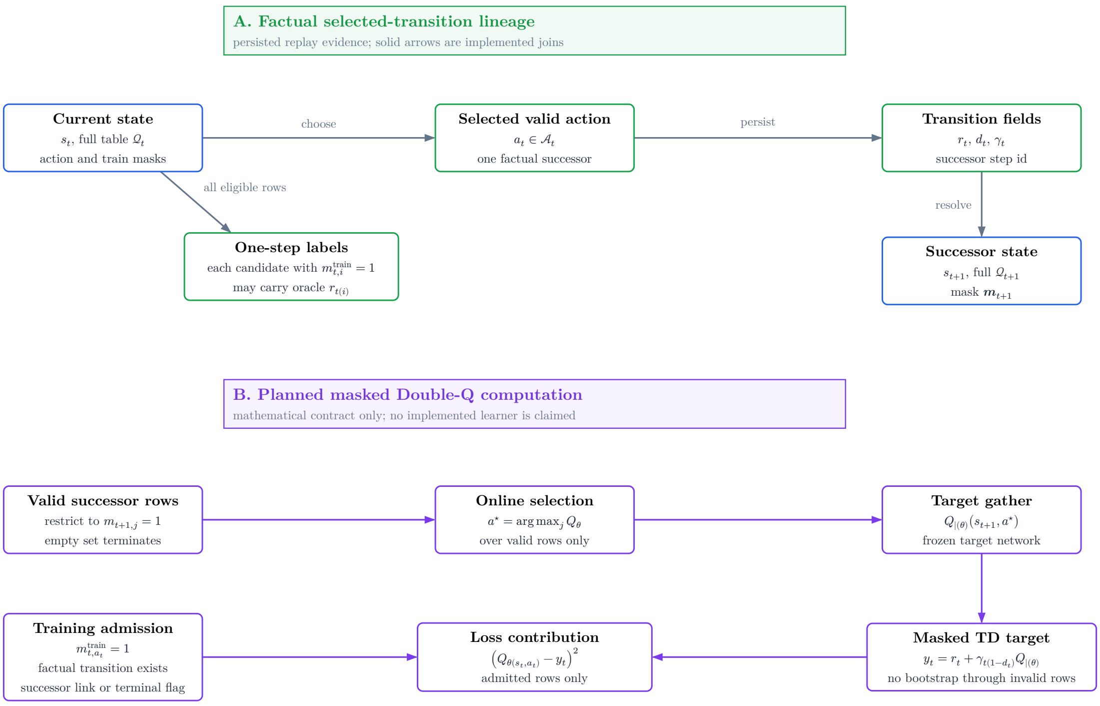

Figure 7: Factual replay lineage and planned masked Double-Q computation. Panel A contains only implemented evidence: every candidate row admitted by its row-wise `q_train_mask` entry may carry a one-step oracle label, whereas TD admission additionally requires that the selected candidate entry is train-valid and that factual transition fields identify either a valid successor or a terminal transition. Panel B states the proposed masked backup mathematically and is not evidence that a learner, residual head, or policy evaluation already exists.

The repository currently implements the replay and mask contract but not the finite-horizon learner. Consequently, rollout audits may establish data readiness and oracle headroom, whereas Q_H policy performance remains outside the supported evidence until a learner and matched held-out evaluation exist.

5.3 Policy Comparison and Statistical Protocol

Implementation: partial **Evidence: pending**

The aggregation and artifact-driven reporting seam exists, but confirmatory paired policy evidence does not.

Promotion gate: frozen held-out manifest, completed paired endpoints, and validated report bundle

Development source: aria_nbv/aria_nbv/rollouts/reporting.py; docs/typst/thesis/experiment_data.typ

The experimental unit is the scene. Every comparison is paired by scene, target task, root state, candidate-generation distribution, hard validity regime, and acquisition horizon. Planner depth may differ because it defines the decision rule, but every policy receives the same acquisition budget. Repeated snippets, targets, and seeds are aggregated within scene before the primary comparison so that densely sampled scenes do not dominate the estimand.

The primary estimand is the paired per-scene difference in fixed-budget endpoint target-quality gain. The centralized analysis artifact freezes the scene aggregation function, uncertainty procedure, interval level, comparison set, and minimum meaningful effect before final aggregation. Secondary summaries—cumulative root-normalized gain, invalid-action rate, runtime, path length, and task/candidate coverage—diagnose mechanism and feasibility but do not replace the endpoint estimand. The current replay store alone cannot estimate this endpoint comparison because it does not persist an independent post-budget reconstruction for every policy; confirmatory analysis therefore requires matched endpoint oracle re-evaluation or an equivalent frozen endpoint record.

Support and failure strata are fixed before policy inspection. The report distinguishes source rows without an admitted target task, admitted tasks whose oracle evaluation fails, roots without a valid action, trajectories that terminate before the shared budget, and completed paired trajectories. A policy that cannot continue remains in the comparison with no further gain; it is not silently dropped. These strata delimit the analysis population and expose whether a policy difference is instead a support or systems failure.

Policy	Decision information	Acquisition budget	Scientific role
π_{rand}	actor-visible	H	valid-action lower reference
$\pi_{\text{learned-1}}$	actor-visible	H	learned myopic control
$\pi_{\text{oracle-1}}$	oracle immediate reward	H	one-step oracle-greedy comparator
$\pi_{\text{oracle-look}}$	bounded oracle lookahead	H	conditional finite-support upper reference
π_Q	actor-visible	H	planned learned finite-horizon recovery

Table 9: Matched policy comparison. All rows acquire the same number of views; oracle access and planner depth define the decision rule rather than the acquisition budget.

The first inferential comparison is bounded oracle lookahead against one-step oracle greedy. Repeated oracle evaluation estimates the metric noise floor, and the analysis artifact freezes the meaningful-headroom rule before learned-policy inspection. Recovered-headroom ratios are reported only when their denominator passes that gate. The learned comparison then contrasts Q_H with the actor-visible myopic control under the same endpoint evaluation and acquisition budget; because neither learned target-conditioned control is presently complete, this comparison remains prospective.

Validation TODO: Populate the evidence chain in order: metric repeatability, candidate and target support, oracle-lookahead headroom, myopic-control calibration, finite-horizon recovery, then representation and architecture ablations. Missing upstream evidence blocks downstream claims rather than becoming a zero result.

Source: thesis objective-to-evidence contract

Gate: artifact-backed Results bundle

Validation TODO: Run row-shuffle, mask-isolation, duplicate-row, valid-count, frame-transform, target-source-dropout, and horizon-boundary tests for every model admitted to the policy comparison.

Source: geometric acceptance contract

Gate: architecture validity report

Open decision: Freeze the scene aggregation, interval procedure, interval level, comparison family, and meaningful-headroom threshold in the resolved analysis manifest rather than in thesis prose.

Source: artifact-driven reporting plan

Gate: confirmatory analysis freeze

5.4 Outcome Logic

Meaningful, stable lookahead headroom establishes only that the frozen finite-candidate setup contains exploitable non-myopic structure. No meaningful headroom is a setup-specific negative result, not evidence that target-aware planning is universally myopic. An unstable oracle metric or insufficient paired support blocks downstream policy claims. Likewise, train-only bandwidth pilots and failed generation attempts, including memory-exhaustion failures, inform compute configuration and throughput planning only; incomplete artifacts cannot establish headroom or policy superiority.

6 Results

The thesis report bundle is loaded through the strict schema checked in `experiment_data.typ`. Its declared evidence class is *pilot*. Schema validity establishes only that provenance, missingness, and table contracts are readable; it does not turn a development fixture or an incomplete pilot into scientific evidence.

Result	Required evidence	Status
Population and targets	population facts with denominators in every validated store	not available
Candidate support	valid/total candidates and zero-action failures in every validated store	not available
Paired policy effect	effect, interval, scene count, and headroom gate in every validated store	not available
Resource feasibility	wall time, peak GPU memory, and storage in every validated store	not available

Table 10: Evidence availability in the loaded report bundle. "Not available" means that no thesis result is inferred from fixture values, train-only attempts, or an incomplete store.

6.1 Candidate and Store Feasibility

The available artifacts show that the finite-candidate rollout path reaches mesh rendering, target-specific oracle scoring, and selected-action replay on training sources. They therefore support an implementation-readiness claim only. A CUDA out-of-memory failure in an unbatched candidate render and later memory-bounded attempts identify rendering as a scale gate; neither establishes rollout throughput, storage cost, candidate-family support, or policy quality for the intended study population.

6.2 Target-Task Coverage

The implemented pipeline samples geometry-valid GT target tasks and records target and validity fields in the rollout schema. The loaded bundle does not contain a validated population and exclusion table for the study. Scene coverage, eligible-target coverage, invalid-reason frequencies, and their denominators are consequently unavailable as results.

6.3 Policy Comparison

The primary estimand is not estimable from the current evidence because no validated held-out bundle contains the matched policy outcomes and aggregation inputs. Oracle repeatability, oracle-lookahead headroom, learned one-step performance, finite-horizon recovery, uncertainty intervals, and sensitivity analyses are therefore not reported.

6.4 Runtime and Storage Gate

The runtime and storage claim requires completed validated stores whose resolved manifests, failure records, and resource summaries refer to the same run. Until that evidence exists for both rollout profiles, the large-scale data-generation gate remains pending. The current attempts justify memory-bounded rendering and retained failure provenance, but no extrapolation to cluster throughput or dataset volume.

7 Discussion

The current evidence supports a narrow implementation conclusion. ARIA-NBV has an executable finite-candidate path that separates actor-visible policy inputs from oracle-only target geometry, applies invalidity as a hard decision constraint, evaluates target-specific reconstruction change offline, and exposes the resulting store through one typed reporting seam. This establishes that the proposed experiment can be represented and audited; it does not establish that one candidate family, rollout policy, representation, or learned model performs better than another.

The rollout attempts also expose a systems limitation. Counterfactual scoring repeatedly renders a large mesh for a candidate set, so renderer memory and latency constrain feasible branch factor and rollout volume before statistical efficiency becomes relevant. An out-of-memory observation is evidence for batching and resource measurement, not for or against the candidate distribution or planning objective. Completed validated stores are required before throughput, storage, or failure rates are generalized.

The scientific limitations are more restrictive. Current attempts use training sources, no held-out paired policy table supports the endpoint estimand, metric repeatability and meaningful oracle-lookahead headroom remain unestablished, and the target-conditioned finite-horizon comparison is not implemented. The present evidence therefore cannot distinguish negligible non-myopic structure from inadequate candidate support, target observability, replay coverage, representation support, or model failure.

7.1 Evidence-conditioned architecture bridges

Implementation: exploratory **Evidence: pending**

Alternative representations and policies remain part of the thesis design registry, but they are promoted only as responses to diagnosed bottlenecks.

Promotion gate: a measured limitation that the proposed bridge directly tests

Development source: Table 3; Table 5

If the local EFM3D field is the limiting variable, broad semidense memory, sparse occupied/free/unknown state, target-centred re-lifting, or logged appearance features provide progressively stronger representation tests. If compact descriptors lose relevant spatial structure, point, sparse, or equivariant encoders become justified. If independent candidate queries leave errors correlated with the valid candidate set, DeepSets or masked candidate interaction becomes relevant. If selected-view history fails specifically by approach direction, signed first- plus second-moment memory, followed only if needed by spherical harmonics, becomes the targeted ablation. A renderable 3DGS state is appropriate only when candidate rendering or soft instance membership is itself the measured missing capability.

Continuous and hierarchical policies change the action contract rather than merely increasing model capacity. A target-then-pose policy could first select or mask a target token and then condition pose queries on target, horizon, and step; multiple targets can share scene encodings and be evaluated in parallel. This remains post-core work because continuous pose generation requires separate feasibility, safety, support, and simulator-realism evidence. The finite-candidate model is retained as the calibrated control even if a continuous policy is later introduced.

Sequence or recurrent models are similarly conditional. They become scientifically motivated when errors persist after explicit history, time, horizon, and scene-memory conditioning, not simply because the dataset contains trajectories. Looping or recurrent refinement must preserve causal masks and candidate-row equivariance and must be compared against the same fixed-budget endpoint oracle evaluation.

Research TODO: After the primary evidence chain is complete, promote only bridges tied to a measured failure: EVL support, target observability, directional history, valid-set interaction, long-horizon credit, or action-support mismatch.

Source: method design-space registry

Gate: artifact-backed failure analysis

The next evidential step remains procedural: complete validated stores under resolved manifests, freeze the eligible population and exclusions, establish oracle stability and headroom, train the matched controls, and only then interpret architectural differences. The registry above preserves concrete follow-up hypotheses without presenting them as validated conclusions.

8 Conclusion

This thesis defines a leakage-auditable experiment for target-conditioned finite-candidate next-best-view planning. Its present contribution is the separation of actor-visible state from oracle supervision, the target-specific reconstruction objective, hard validity and replay contracts, and an artifact-driven reporting seam that keeps provenance and missingness attached to later results.

The available evidence does not answer whether bounded oracle lookahead improves fixed-budget target reconstruction or whether a learned finite-horizon policy recovers such headroom. The current training-source rollout attempts establish pipeline reachability and reveal a renderer resource gate; the development report fixture establishes the data contract only. Neither supports a held-out policy claim, a population-level effect, or a scale estimate.

The final scientific conclusion is therefore conditional on evidence that is not yet available. A stable oracle metric and positive paired lookahead effect would define measurable headroom for the evaluated finite support. Negligible headroom would be a setup-specific negative result. Unstable oracle evaluation would block planning claims, and stable headroom without learned recovery would remain a learned-control failure with several unresolved mechanisms. These outcomes delimit the thesis without extending it to continuous control, online reinforcement learning, or real-device deployment.

Bibliography

[Baj88] R. Bajcsy, “Active Perception,” *Proceedings of the IEEE*, vol. 76, no. 8, pp. 966–1005, 1988, doi: 10.1109/5.5968.

[AWB88] J. Aloimonos, I. Weiss, and A. Bandyopadhyay, “Active Vision,” *International Journal of Computer Vision*, vol. 1, no. 4, pp. 333–356, 1988, doi: 10.1007/BF00133571.

[SRR03] W. R. Scott, G. Roth, and J.-F. Rivest, “View Planning for Automated Three-Dimensional Object Reconstruction and Inspection,” *ACM Computing Surveys*, vol. 35, no. 1, pp. 64–96, 2003, doi: 10.1145/641865.641868.

[Fra+25] N. Frahm *et al.*, “VIN-NBV: A View Introspection Network for Next-Best-View Selection.” [Online]. Available: https://arxiv.org/abs/2505.06219

[Eng+23] J. Engel *et al.*, “Project Aria: A New Tool for Egocentric Multi-Modal AI Research.” [Online]. Available: https://arxiv.org/abs/2308.13561

[Met25] Meta Platforms Inc., “Aria Synthetic Environments Dataset.” [Online]. Available: https://facebookresearch.github.io/projectaria_tools/docs/open_datasets/aria_synthetic_environments_dataset

[Str+24] J. Straub, D. DeTone, T. Shen, N. Yang, C. Sweeney, and R. Newcombe, “EFM3D: A Benchmark for Measuring Progress Towards 3D Egocentric Foundation Models.” [Online]. Available: https://arxiv.org/abs/2406.10224

[Met25] Meta Platforms Inc., “Egocentric Voxel Lifting (EVL) documentation.” [Online]. Available: https://facebookresearch.github.io/projectaria_tools/docs/open_models/evl

[Bir+16] A. Bircher, M. Kamel, K. Alexis, H. Oleynikova, and R. Siegwart, “Receding Horizon “Next-Best-View” Planner for 3D Exploration,” in *2016 IEEE International Conference on Robotics and Automation*, 2016, pp. 1462–1468. doi: 10.1109/ICRA.2016.7487281.

[Lau+20] M. Lauri, J. Pajarinen, J. Peters, and S. Frintrop, "Multi-Sensor Next-Best-View Planning as Matroid-Constrained Submodular Maximization," *IEEE Robotics and Automation Letters*, vol. 5, no. 4, pp. 5323–5330, 2020, doi: 10.1109/LRA.2020.3007445.

[Jia+25] Z. Jia, Y. Li, Q. Hao, and S. Zhang, "PB-NBV: Efficient Projection-Based Next-Best-View Planning Framework for Reconstruction of Unknown Objects," *IEEE Robotics and Automation Letters*, vol. 10, no. 7, pp. 7444–7451, 2025, doi: 10.1109/LRA.2025.3573631.

[Jeo+26] S. Jeong, E. Lee, J. Kim, and A. Kim, "Informative Object-centric Next Best View for Object-aware 3D Gaussian Splatting in Cluttered Scenes." [Online]. Available: https://arxiv.org/abs/2602.08266

[Hu+26] B. Hu *et al.*, "OA-NBV: Occlusion-Aware Next-Best-View Planning for Human-Centered Active Perception on Mobile Robots." [Online]. Available: https://arxiv.org/abs/2603.11072

[Pan+26] S. Pan, H. Hu, X. Huang, B. Wingender, and M. Bennewitz, "ObjView-Bench: Rethinking Difficulty and Deployment for Object-Centric View Planning." [Online]. Available: https://arxiv.org/abs/2605.10707

[HGS15] H. van Hasselt, A. Guez, and D. Silver, "Deep Reinforcement Learning with Double Q-learning." [Online]. Available: https://arxiv.org/abs/1509.06461

[Kum+20] A. Kumar, A. Zhou, G. Tucker, and S. Levine, "Conservative Q-Learning for Offline Reinforcement Learning," in *Advances in Neural Information Processing Systems*, 2020, pp. 1179–1191. [Online]. Available: https://papers.nips.cc/paper/2020/hash/0d2b2061826a5df3221116a5085a6052-Abstract.html

[FMP19] S. Fujimoto, D. Meger, and D. Precup, "Off-Policy Deep Reinforcement Learning without Exploration," in *Proceedings of the 36th International Conference on Machine Learning*, in Proceedings of Machine Learning Research, vol. 97. PMLR, 2019, pp. 2052–2062.

[Online]. Available: https://proceedings.mlr.press/v97/fujimoto19a.html

[Zah+17] M. Zaheer, S. Kottur, S. Ravanbakhsh, B. Poczos, R. Salakhutdinov, and A. J. Smola, "Deep Sets," in *Advances in Neural Information Processing Systems*, 2017. [Online]. Available: https://papers.nips.cc/paper_files/paper/2017/hash/f22e4747da1aa27e363d86d40ff442fe-Abstract.html

[Lee+19] J. Lee, Y. Lee, J. Kim, A. R. Kosiorek, S. Choi, and Y. W. Teh, "Set Transformer: A Framework for Attention-based Permutation-Invariant Neural Networks," in *Proceedings of the 36th International Conference on Machine Learning*, in Proceedings of Machine Learning Research, vol. 97. PMLR, 2019, pp. 3744–3753. [Online]. Available: https://proceedings.mlr.press/v97/lee19d.html

[Zho+23] Z. Zhou, J. Wang, Y.-H. Li, and Y.-K. Huang, "Query-Centric Trajectory Prediction," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.

[Qia+22] G. Qian *et al.*, "PointNeXt: Revisiting PointNet++ with Improved Training and Scaling Strategies." [Online]. Available: https://arxiv.org/abs/2206.04670

[Wu+24] X. Wu *et al.*, "Point Transformer V3: Simpler, Faster, Stronger," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024. [Online]. Available: https://arxiv.org/abs/2312.10035

[Tho+19] H. Thomas, C. R. Qi, J.-E. Deschaud, B. Marcotegui, F. Goulette, and L. J. Guibas, "KPConv: Flexible and Deformable Convolution for Point Clouds," in *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2019. [Online]. Available: https://openaccess.thecvf.com/content_ICCV_2019/html/Thomas_KPConv_Flexible_and_Deformable_Convolution_for_Point_Clouds_ICCV_2019_paper.html

[CGS19] C. Choy, J. Gwak, and S. Savarese, “4D Spatio-Temporal ConvNets: Minkowski Convolutional Neural Networks,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019. [Online]. Available: https://arxiv.org/abs/1904.08755

[SHW21] V. G. Satorras, E. Hoogeboom, and M. Welling, “E(n) Equivariant Graph Neural Networks,” in *Proceedings of the 38th International Conference on Machine Learning*, in Proceedings of Machine Learning Research, vol. 139. PMLR, 2021, pp. 9323–9332. [Online]. Available: https://proceedings.mlr.press/v139/satorras21a.html

[Fuc+20] F. B. Fuchs, D. E. Worrall, V. Fischer, and M. Welling, “SE(3)-Transformers: 3D Roto-Translation Equivariant Attention Networks,” in *Advances in Neural Information Processing Systems*, 2020. [Online]. Available: https://arxiv.org/abs/2006.10503

[Bre+23] J. Brehmer, P. de Haan, S. Behrends, and T. Cohen, “Geometric Algebra Transformer,” in *Advances in Neural Information Processing Systems*, 2023. [Online]. Available: https://arxiv.org/abs/2305.18415

[Che+24] X. Chen, Q. Li, T. Wang, T. Xue, and J. Pang, “GenNBV: Generalizable Next-Best-View Policy for Active 3D Reconstruction,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 16436–16445. [Online]. Available: https://openaccess.thecvf.com/content/CVPR2024/html/Chen_GenNBV_Generalizable_Next-Best-View_Policy_for_Active_3D_Reconstruction_CVPR_2024_paper.html

[Lu+26] C.-Y. Lu *et al.*, “Hestia: Voxel-Face-Aware Hierarchical Next-Best-View Acquisition for Efficient 3D Reconstruction,” in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2026. [Online]. Available: https://openaccess.thecvf.com/content/WACV2026/papers/Lu_Hestia_Voxel-Face-Aware_Hierarchical_

Next-Best-View_Acquisition_for_Efficient_3D_Reconstruction_WACV_2026_paper.pdf

[Ave+24] A. Avetisyan *et al.*, "SceneScript: Reconstructing Scenes With An Autoregressive Structured Language Model." [Online]. Available: https://arxiv.org/abs/2403.13064

[Bro+21] M. M. Bronstein, J. Bruna, T. Cohen, and P. Veličković, "Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges." [Online]. Available: https://arxiv.org/abs/2104.13478

[Zho+19] Y. Zhou, C. Barnes, J. Lu, J. Yang, and H. Li, "On the Continuity of Rotation Representations in Neural Networks," *arXiv preprint arXiv:1812.07035*, 2019, [Online]. Available: https://arxiv.org/abs/1812.07035

[e3n25] e3nn contributors, "e3nn spherical harmonics documentation." [Online]. Available: https://docs.e3nn.org/en/stable/api/o3/o3_sh.html

[GML22] A. Guédon, P. Monasse, and V. Lepetit, "SCONE: Surface Coverage Optimization in Unknown Environments by Volumetric Integration," in *Advances in Neural Information Processing Systems*, 2022. [Online]. Available: https://arxiv.org/abs/2208.10449

[JLD24] W. Jiang, B. Lei, and K. Daniilidis, "FisherRF: Active View Selection and Uncertainty Quantification for Radiance Fields using Fisher Information." [Online]. Available: https://arxiv.org/abs/2311.17874

[GE19] Y. Geifman and R. El-Yaniv, "SelectiveNet: A Deep Neural Network with an Integrated Reject Option," in *Proceedings of the 36th International Conference on Machine Learning (ICML)*, in Proceedings of Machine Learning Research, vol. 97. PMLR, 2019, pp. 2151–2159. [Online]. Available: https://proceedings.mlr.press/v97/geifman19a.html

[Liu+19] Z. Liu, Z. Wang, P. P. Liang, R. Salakhutdinov, L.-P. Morency, and M. Ueda, "Deep Gamblers: Learning to Abstain with Portfolio Theory," in *Advances in Neural Information Processing Systems*, 2019, pp. 10622–

10632. [Online]. Available: https://proceedings.neurips.cc/paper/2019/hash/0c4b1eeb45c90b52bfb9d07943d855ab-Abstract.html

[Ben+09] Y. Bengio, J. Louradour, R. Collobert, and J. Weston, “Curriculum Learning,” in *Proceedings of the 26th International Conference on Machine Learning (ICML)*, 2009. [Online]. Available: https://proceedings.mlr.press/v5/bengio09a.html

[Met25] Meta Platforms Inc., “PyTorch3D Cameras + Rendering Documentation, Tutorial: Render a Textured Mesh.” [Online]. Available: https://pytorch3d.readthedocs.io/en/latest/modules/renderer/cameras.html,%20https://pytorch3d.org/docs/renderer,%20https://pytorch3d.org/docs/renderer_getting_started,%20https://pytorch3d.org/docs/cameras,%20https://pytorch3d.org/tutorials/render_textured_meshes

[DAT20] N. De Cao, W. Aziz, and I. Titov, “The Power Spherical distribution.” [Online]. Available: https://arxiv.org/abs/2006.04437

[Li+21] Y. Li, S. Si, G. Li, C.-J. Hsieh, and S. Bengio, “Learnable Fourier Features for Multi-Dimensional Spatial Positional Encoding.” [Online]. Available: https://arxiv.org/abs/2106.02795

[CMR19] W. Cao, V. Mirjalili, and S. Raschka, “Rank consistent ordinal regression for neural networks with application to age estimation.” [Online]. Available: https://arxiv.org/abs/1901.07884

A Reproducibility Record

The report bundle is the single numerical interface between validated rollout artifacts and this manuscript. Its schema version, evidence status, store manifests, typed parameter rows, failure records, and sidecar hashes preserve provenance without duplicating analysis logic in Typst. The document loader rejects schema or status mismatches before any table is rendered. Compact labels and digest prefixes are used below; complete store identifiers, paths, and hashes remain in the machine-readable bundle.

Evidence class	Store	Manifest digest	Validated
Development fixture	store S1	not-a-real-m…	yes

Table 11: Loaded development-fixture provenance. The fixture verifies the document interface only and is not scientific evidence.

The bundled development fixture contains synthetic values solely to test typed loading and missingness. Its numerical rows are deliberately not rendered or interpreted.

B Development Diary

Archived source note: This marked diary is a development-only record. It is excluded from submission mode; only evidence-backed definitions, protocols, and results graduate into the thesis body.

Source: thesis mode contract

B.1 Decisions Retained

The thesis keeps the privileged oracle and actor-visible policy as separate contracts, evaluates finite candidate sets under a shared validity mask, and treats target-specific reconstruction improvement as the common comparison signal. Rollout stores retain selected transitions and provenance rather than becoming a second implementation of the oracle or training pipeline.

Open decision: Freeze the scene-disjoint analysis population and the minimum evidence scale only when the resolved manifests and exclusion ledger are available.

Source: experimental-design contract

Gate: confirmatory bundle freeze

B.2 Failures That Changed the Plan

Early rollout attempts reached the mesh-rendering path but exceeded available accelerator memory when candidate views were rendered without a bounded batch. This failure narrowed the immediate claim to implementation readiness and made peak memory, throughput, failure provenance, and completed-store validation explicit scale gates.

Validation TODO: Require completed stores, matching manifests, resource summaries, and recorded failures before extrapolating rollout bandwidth or storage demand.

Source: rollout attempt artifacts

Gate: cluster-readiness evidence

B.3 Rejected Scope

Continuous actions, online simulators, additional scene representations, and alternative reinforcement-learning families are excluded from the core study while the finite-candidate target-conditioned comparison remains unvalidated. They would change both the action contract and the supervision regime, so adding them now would weaken rather than extend the primary experiment.

Remove or rewrite before submission: Delete bridge material that cannot be tied to a concrete limitation or a matched follow-up experiment after the confirmatory analysis is complete.

Source: scope review

Gate: discussion freeze

B.4 Directional-Memory Hypothesis

Research TODO: Treat target-centred directional memory as an ablation hypothesis, not an actor feature or contribution. The representation must be implemented, mask-audited, and compared against pose distance and overlap before any predictive claim enters the manuscript.

Source: RQ3 representation hypothesis

Gate: paired held-out ablation

A. Target-centred directions on $\mathbb{S}^2$

actual logged camera centres, expressed in the target-object frame

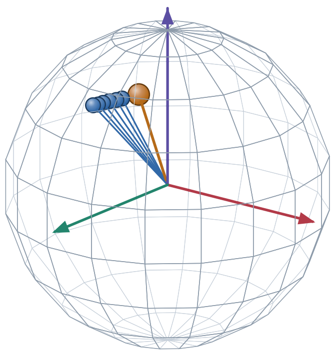

● logged history direction ⊕ example query

Scene 81286, sample ASE_81286_Atek_000024, target instance 128 frames 0, 2, 4, 6, 8, 10 and query frame 15.

B. Complete-domain view and prospective compression

Mollweide is equal-area; no smooth field or SH lobe is implied

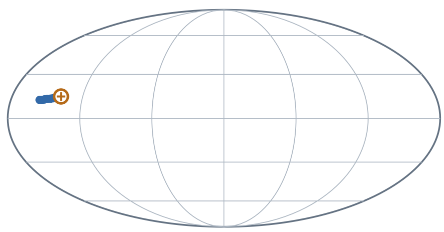

$+x_e$ at map centre · $+z_e$ north · wrap seam at $\pm 180°$

$$M_{\mathrm{dir}(v)} = \textstyle\sum_{k<t} w_{k(v)} \boldsymbol{d}_{k(v)} \boldsymbol{d}_{k(v)}^\top$$

$$\nu_{i(v)} = 1 - \frac{\boldsymbol{d}_i^\top \boldsymbol{M}_{\mathrm{dir}} \boldsymbol{d}_i}{\mathrm{tr}\, \boldsymbol{M}_{\mathrm{dir}} + \varepsilon}$$

This fixture uses uniform weights and a logged future pose only to make the second-moment query inspectable. It does not claim that the learner currently stores $\boldsymbol{M}_{\mathrm{dir}}$ or that this novelty predicts target RRI.

HYPOTHESIS / ABLATION MATERIAL. Keep outside submission claims until the representation exists and paired evaluation tests whether it adds information beyond pose distance and overlap.

Figure 8: Development-only directional-memory fixture. The sphere and Mollweide panels show the same logged ASE camera directions in a target-object frame; the orange query is another logged pose used to inspect one prospective second-moment novelty definition. No smooth field, spherical-harmonic representation, or learned benefit is claimed.

B.5 Next Evidence Gates

The next admissible evidence is a validated report bundle that fixes the study population, records target and candidate exclusions, supplies paired endpoint outcomes with uncertainty, and joins runtime and storage statistics to the same manifests. Only then can the manuscript replace the present availability statements with policy comparisons.

Implementation TODO: Export the confirmatory report bundle, compile both thesis modes against it, and inspect the resulting tables and failure cases before freezing claims.

Source: reporting contract

Gate: claim freeze

Ehrenwoertliche Erklaerung

Ich versichere, dass ich diese Masterarbeit selbststaendig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe.

Muenchen, 7. Jan 2027 Jan Duchscherer