



Munich University of Applied Sciences

Hochschule Muenchen

Department of Computer Science & Mathematics

Fakultaet fuer Informatik und Mathematik

Quality-Driven Next-Best-View Planning for Target-Conditioned Egocentric 3D Reconstruction

Qualitätsgetriebene Next-Best-View-Planung für zielkonditionierte egozentrische 3D-Rekonstruktion

Master's Thesis / Masterarbeit

for obtaining the academic degree Master of Science (M.Sc.)

submitted to
Munich University of Applied Sciences
Department of Computer Science & Mathematics

Submitted by: Jan Duchscherer
Email: j.duchscherer@hm.edu
Program: Informatik - Master
Specialization: Visual Computing and Machine Learning
Matriculation Number: 00807819

First Examiner: Prof. Dr. Markus Friedrich
Second Examiner: Prof. Dr. David Spieler
Supervisor: Prof. Dr. Markus Friedrich
Submission Date: 7. Jan 2027

Transparency in the Use of AI Tools

Generative-AI tools supported literature-note organization, consistency checks, and language revision. The author selected the research questions, verified sources and technical claims, implemented and evaluated the system, and remains responsible for the submitted work.

Acknowledgements

Acknowledgements are omitted from this version.

Abstract

Active perception couples sensing and reconstruction: what can be reconstructed depends on what a sensing action makes visible. Under a limited acquisition budget, next-best-view planning must therefore choose which feasible observation should come next. The literature reviewed in Chapter 2 treats scene coverage, one-step reconstruction-quality ranking, target-aware information criteria, and sequential coverage separately, but leaves their conjunction for target-specific egocentric reconstruction unresolved. This thesis studies that conjunction as finite candidate selection under hard validity constraints in Aria Synthetic Environments (ASE). It adapts Relative Reconstruction Improvement (RRI) to a target-cropped point-mesh objective, separates privileged task construction and oracle evaluation from the deployable actor path, and marks GT-derived target or selected-depth inputs as non-deployable controls. The primary experiment first tests whether bounded oracle lookahead provides endpoint-quality headroom over one-step oracle greedy; a learned policy is evaluated only if that headroom passes a prespecified meaningful-effect and uncertainty rule. The current implementation supports target-specific oracle scoring and selected-action replay, with rendering memory limiting scale. Actor-visible target matching, metric repeatability, a validated held-out population, headroom, and paired policy outcomes remain unestablished. The present contribution is therefore an auditable method and experimental design, not evidence of policy superiority or deployment readiness.

Zusammenfassung

Aktive Wahrnehmung beruht auf einer geometrischen Prämisse: Was rekonstruiert werden kann, hängt davon ab, was eine Sensorhandlung sichtbar macht. Bei begrenztem Aufnahmebudget muss die Next-Best-View-Planung daher entscheiden, welche zulässige Beobachtung als Nächstes erfolgen soll. Die in Kapitel 2 ausgewertete Literatur behandelt Szenenabdeckung, einstufige Rekonstruktionsqualität, zielbezogene Informationsmaße und sequenzielle Abdeckung getrennt, lässt deren Verbindung für die zielspezifische egozentrische Rekonstruktion jedoch offen. Diese Arbeit untersucht diese Verbindung als Auswahl aus einer endlichen Kandidatenmenge unter harten Zulässigkeitsbedingungen in ASE. Sie überträgt die RRI auf ein zielbeschnittenes Punkt–Mesh-Maß, trennt privilegierte Aufgabenerzeugung und Orakelevaluation vom einsatzfähigen Akteurpfad und kennzeichnet GT-basierte Ziel- oder Tiefeneingaben als nicht einsetzbare Kontrollen. Das primäre Experiment prüft zunächst, ob eine beschränkte Orakelvorausschau gegenüber einer einstufigen gierigen Orakelstrategie einen Vorteil bei der Endpunktqualität bietet; eine gelernte Strategie wird nur bewertet, wenn dieser Vorteil ein vorab festgelegtes Relevanz- und Unsicherheitskriterium erfüllt. Die aktuelle Implementierung unterstützt zielspezifische Orakelbewertung und Replay ausgewählter Aktionen, wobei der Speicherbedarf des Renderers die Skalierung begrenzt. Nicht belegt sind bislang eine beobachtungs-basierte Zielzuordnung, die Wiederholbarkeit des Qualitätsmaßes, eine validierte zurückgehaltene Studienpopulation, ein Vorteil der Orakelvorausschau und gepaarte Strategieergebnisse. Der gegenwärtige Beitrag ist daher ein prüfbarer Methoden- und Versuchsaufbau, kein Nachweis für Strategieüberlegenheit oder Einsatzreife.

Contents

List of Figures	27
List of Tables	30
1 Introduction	1
1.1 Aim and bounded problem	2
1.2 Contributions and current evidence	3
1.3 Research Questions	3
1.3.1 Research aim and evaluation logic	3
1.3.2 RQ1 — Target-specific objective and endpoint outcome	4
1.3.3 RQ2 — Bounded lookahead and learned gap closure	4
1.3.4 Conditions for interpretation	5
1.3.4.1 RQ3 — Actor-visible target and information state	5
1.3.4.2 RQ4 — Candidate, replay, and population support	6
1.3.5 Scope extensions	6
1.3.5.1 RQ5 — Online discrete bridge	6
1.3.5.2 RQ6 — Continuous or simulator-backed control	6
1.3.6 Shared evidence constraints	7
1.3.7 Research-to-evidence map	7
1.4 Thesis structure	7
2 Foundations and Related Work	8
2.1 Active Perception and the Next-Best-View Problem	8
2.2 View Utility and Target-Conditioned Reconstruction	9
2.3 Candidate-View Support and Motion Feasibility	11
2.4 Sequential Decision-Making under Partial Observability	12
2.5 Egocentric and Geometric Representations	14
2.6 Related-Work Synthesis and Research Gap	16
3 Oracle and Data Generation	19

3.1	State as an Information Contract	19
3.1.1	Logged Information State	20
3.1.2	Causal Successor State	21
3.1.3	Oracle Augmentation	22
3.1.4	Visibility, Feasibility, and Utility	23
3.2	Task and Action Construction	24
3.2.1	Target Selection	25
3.2.2	Observed-Target Admission	25
3.2.3	Candidate View Generation	26
3.2.4	Candidate-Support and Geometric-Coverage Diagnostics	29
3.2.5	Rollout Branch Sampling and Dataset Impact	32
3.3	Measurement and Outcomes	34
3.3.1	Operational Reconstruction Error	34
3.3.2	Reliability under the Frozen Protocol	35
3.3.3	Construct Validity	37
3.3.4	RRI, Training Reward, and Endpoint Gain	37
3.4	Dataset and Storage Semantics	39
3.4.1	Immutable Sources and Causal Replay	40
3.4.2	Missingness and Derived Views	41
4	Method	43
4.1	Current and Target Actor State	45
4.1.1	Sufficiency and the local-evidence limit	46
4.1.2	Candidate realizations of the state contract	47
4.2	Target, Candidate, and History Encoding	49
4.2.1	Logical data roles	49
4.2.2	Target-conditioned query	50
4.2.3	Relative candidate geometry	50
4.2.4	Strictly causal pose history	51
4.3	Finite Action and Causal Replay	53
4.3.1	Supervision and audit contracts	54

4.3.2	Current replay transition	55
4.3.3	Scientific target transition	56
4.4	Selected Interaction and Acceptance Conditions	57
4.4.1	Architectural feature-integration ladder	58
4.5	Target-Conditioned Finite-Horizon Value	61
4.5.1	Conditional scorer	61
4.5.2	Scalar requested horizon	62
4.5.3	Direct continuous objective	63
4.5.4	Decoder and estimator design space	63
4.5.5	Fitted-Q recursion	64
4.5.6	Why exact horizon two is decisive	65
5	Experimental Design	67
5.1	Population, Estimands, and Gate Order	68
5.1.1	Candidate-Support Estimands and Aggregation	68
5.2	Replay Admission for the Learning Gates	70
5.3	Matched Policies and Failure Attribution	72
6	Results	75
6.1	Measurement Validity	76
6.2	Population and Action Support	76
6.3	Actor-Visible Protocol	77
6.4	Oracle Headroom	77
6.5	Actor-Visible One-Step Value	77
6.6	Exact Two-Step Recursion	78
6.7	Endpoint Recovery	78
6.8	Resource Feasibility	78
7	Discussion	79
7.1	Interpreting future outcomes	79
7.2	Systems and external-validity boundaries	80
7.3	Evidence-conditioned architecture bridges	81
8	Conclusion	83

Development roadmap	85
TL;DR	85
Critical path	86
Milestones	86
Blockers	87
Promotion queue	87
Maintenance contract	87
HM/FK07 release gate	88
Freeze and submission	88
9 Development Pilots	89
9.1 Candidate-Family Support	89
9.2 Target-Frame $\mathcal{S}^2$ Diagnostics	90
Development Register: Method Alternatives	93
State and history alternatives	93
Interaction alternatives	94
Objective and estimator alternatives	94
Ranked scientific-core work	95
Bibliography	97
A Reproducibility Record	102
B Development Diary	103
B.1 Decisions Retained	103
B.2 Failures That Changed the Plan	103
B.3 Rejected Scope	103
B.4 Directional-Memory Hypothesis	104
B.5 Next Evidence Gates	105

Glossary and Abbreviations

Dataset

***ADT* – Aria Digital Twin:** Project Aria dataset with real-world captures and digital-twin scene annotations. ADT is adjacent to ARIA-NBV's sim-to-real context; the current experiments focus on ASE and EFM3D/ATEK exports, while ADT remains a relevant transfer surface.

***AEO* – Aria Everyday Objects:** Small-scale real-world Project Aria object dataset used by EFM3D for egocentric 3D perception evaluation. AEO is relevant as a possible sim-to-real check for ARIA-NBV ideas that are first developed on ASE mesh-supervised snippets.

***ASE* – Aria Synthetic Environments:** Large-scale synthetic indoor dataset with simulated Project Aria sensor characteristics, egocentric trajectories, and scene annotations. ARIA-NBV uses ASE snippets and the public mesh-supervised subset to generate oracle RRI labels from ground-truth meshes and semi-dense point clouds.

***CPF* – Central Pupil Frame:** Coordinate frame placed at the midpoint between the left and right eye boxes of Project Aria glasses. CPF is used by Project Aria for gaze-related quantities and should remain distinct from rig, camera, world, and PyTorch3D frames in ARIA-NBV docs.

***MPS* – Machine Perception Services:** Project Aria processing services that derive pose, mapping, gaze, hand, and related perception outputs from sensor recordings. ARIA-NBV mainly depends on MPS-style pose and semi-dense mapping products as the current reconstruction state for RRI computation.

***MTD* – Motion Trajectory Data:** Device poses over time, usually represented as a sequence of 6-DoF transformations. MTD supplies the logged egocentric trajectory used to define snippet state, current reconstruction context, and candidate-view reference poses.

***MFCD* – Multi-Frame Camera Data:** Synchronized camera streams from multiple Project Aria cameras over a temporal window. MFCD provides the

egocentric image evidence consumed by EVL and aligned with pose, calibration, and semi-dense point streams.

MSDPD – Multi-Semi-Dense Point Data: Semi-dense 3D point observations generated by SLAM-style processing across a snippet or trajectory window. MSDPD is the sparse observed geometry that ARIA-NBV fuses with candidate point clouds and compares against ground-truth meshes for RRI labels.

Project Aria – Project Aria: Meta's egocentric research-device and tooling ecosystem for calibrated, time-aligned multimodal sensing. ARIA-NBV treats Project Aria and its MPS-style products as the actor-visible sensing contract, while ASE meshes remain offline supervision and evaluation assets.

SSL – SceneScript Language: Structured language representation for indoor scene layout using primitives such as walls, doors, windows, and objects. SceneScript is relevant to ARIA-NBV as a possible semantic/global planning layer and as one source of ASE scene-structure context.

snippet – Snippet: Short synchronized temporal window of Aria sensor data used as one EVL/VIN input sample. A snippet typically contains RGB or grayscale streams, poses, calibration, semi-dense points, and scene metadata that EVL lifts into a voxel grid.

VRS – Virtual Reality Standard: File format used by Project Aria tooling to store multi-modal sensor recordings efficiently. VRS is part of the upstream Project Aria ecosystem, while ARIA-NBV primarily consumes ASE/ATEK-derived tensorized snippets for current experiments.

Entity

target – Target of Interest: Selected entity, object crop, point, region, or surface-deficit hypothesis whose reconstruction quality should be improved. Target-conditioned ARIA-NBV variants use this target as an explicit input to candidate scoring and planning instead of optimizing only scene-level RRI.

Geometry

***DoF* – Degrees of Freedom:** Number of independent pose parameters available to a camera, object, or action representation. ARIA-NBV distinguishes full 6-DoF pose estimation from reduced 5-DoF candidate or action spaces used for practical view planning.

***frustum* – Frustum:** Truncated pyramidal camera-visible volume bounded by near and far clipping planes plus lateral field-of-view planes. Candidate-view frusta define which scene surfaces can project into a camera image and are therefore central to visibility, rendering, and RRI diagnostics.

***LUF* – Left-Up-Forward:** Camera coordinate convention whose x axis points left, y axis points up, and z axis points forward. LUF is one of the coordinate-frame conventions that must stay explicit when moving between Project Aria cameras, PyTorch3D cameras, and ARIA-NBV candidate poses.

***OBB* – Oriented Bounding Box:** 3D bounding box with arbitrary orientation, used to represent object extent more tightly than an axis-aligned box. OBBs are a natural target encoding for entity-aware ARIA-NBV because they provide center, extent, orientation, semantic class, and confidence signals.

***6DoF* – Six Degrees of Freedom:** Pose parameterization with three translational and three rotational degrees of freedom. Project Aria poses, candidate cameras, and world-to-device transforms are usually represented as 6-DoF rigid-body transforms.

Metrics

***AUC* – Area Under Curve:** Aggregate score computed by integrating a metric curve over an acquisition, threshold, or ranking axis. AUC can summarize how quickly reconstruction quality, coverage, or ranking quality improves as views are acquired or candidate thresholds change.

***CD* – Chamfer Distance:** Historical bidirectional distance family used to compare reconstructed points against reference geometry. Thesis-facing ARIA-

NBV notation uses point-mesh error D with directional components $D_{P \rightarrow M}$ and $D_{M \rightarrow P}$; older seminar material may still call this CD.

CR – Coverage Ratio: Fraction of a target surface, scene, or region treated as observed under a chosen visibility or distance threshold. Coverage ratio is a useful diagnostic baseline for NBV, but ARIA-NBV treats reconstruction-quality improvement through RRI as the preferred optimization target.

GT – Ground Truth: Reference data or annotations treated as the trusted target for training, validation, or evaluation. In ARIA-NBV, ground truth usually refers to ASE meshes, object annotations, poses, or labels used to compute oracle supervision and diagnostic metrics.

RRI – Relative Reconstruction Improvement: Metric quantifying the relative reconstruction-quality improvement obtained by adding a candidate observation to the current reconstruction. ARIA-NBV computes RRI by comparing reconstruction error before and after fusing candidate-view geometry, usually against an ASE ground-truth mesh for oracle supervision and evaluation.

reward – Target Root-Gain Reward: Quality-only immediate reward for thesis-core target-aware rollouts. It is the target-specific point-mesh error reduction normalized by the rollout-root target error; scalar motion, rule, diversity, and validity penalties are later ablations.

target RRI – Target-Specific RRI: RRI computed only on the ground-truth and reconstructed geometry associated with a selected target of interest. Target-specific RRI lets the thesis compare views by how much they improve a selected object or region, even when scene-level RRI would prefer large background surfaces.

Model

Q-CORAL – CORAL Q-Value Decoder: Ordinal Q-decoder ablation that discretizes continuous fitted-Q targets into ordered classes and decodes

cumulative threshold outputs to a scalar conditional value through fixed return representatives.

EFM3D – Egocentric Foundation Model 3D: Egocentric 3D foundation-model stack used as the frozen spatial backbone for ARIA-NBV candidate scoring. The current project uses EFM3D and its EVL architecture to expose voxel occupancy, centerness, semantic, and OBB evidence for VIN-style RRI prediction.

EVL – Egocentric Voxel Lifting: EFM3D architecture that lifts synchronized egocentric observations into a gravity-aligned 3D voxel feature volume. ARIA-NBV uses EVL head outputs, pre-head features, and OBB predictions as local actor-visible evidence and target support, not as the complete frozen scene representation.

Q_H – Finite-Horizon Q Function: Target-conditioned finite-candidate action-value interface for the return after choosing one candidate first.

target-conditioned scorer – Target-Conditioned Scorer: VIN-style candidate scorer that receives scene state, a candidate view, and an encoding of the target of interest. The scorer predicts target-specific utility so view ranking can prioritize a selected entity or region instead of only optimizing scene-level RRI.

VIN – View Introspection Network: Learned candidate-view scorer that predicts RRI or an ordinal RRI-derived utility without capturing the candidate view. ARIA-NBV adapts VIN-NBV by placing a lightweight RRI prediction head on top of frozen EVL features and candidate-pose evidence from ASE snippets.

Planning

cost – Acquisition Cost: Budget consumed to acquire observations, measured by view count, path length, elapsed time, invalid-action rate, or a weighted combination. The thesis should first report cumulative root-normalized target

gain, diagnostic target RRI, and acquisition cost separately, then use scalarized objectives only when the tradeoff is explicit.

candidate – Candidate View: Proposed camera pose whose expected reconstruction utility is evaluated before selecting the next observation. ARIA-NBV samples candidate views around a reference pose or target, renders candidate depths from the mesh for oracle labels, and scores candidates with RRI or learned VIN predictions.

transition – Counterfactual Transition: Deterministic or replayable rollout update that renders or retrieves only the selected candidate observation, backprojects it to candidate geometry, and fuses it into the accumulated reconstruction proxy.

action set – Finite Candidate Action Set: Valid discrete action set for ARIA-NBV planning, obtained by masking a sampled candidate view table. The action is the candidate-table index, not a continuous pose; the selected pose is $q_t = q_{t,a_t}$. Invalid candidates are constraints, not low-quality actions.

return – Finite-Horizon Return: Symbolic H-step return used by ARIA-NBV rollout evaluation and Q_H targets. The formula keeps γ explicit while the thesis-core comparison reports cumulative root-normalized target gain under equal acquisition budget.

5DoF – Five Degrees of Freedom: Reduced camera-action parameterization commonly used when roll is fixed or otherwise constrained. ARIA-NBV uses 5-DoF language for candidate and planning abstractions where position and viewing direction matter while roll is not an independently optimized control dimension.

CF+ state – Geometry-Rich Counterfactual State: Implemented privileged S1 selected-surface control. Its CF-GT carrier stores only previously selected rendered depth, valid mask, calibration, selected-camera pose, source identity, and causal support; an identity-start point encoder consumes that carrier. It is neither the actor-visible scientific target nor evidence that selected geometry improves value prediction.

historic state – Historic Snippet State: Raw actor-visible state available along the logged Project Aria/ASE trajectory before counterfactual rollout synthesis. It includes captured camera streams, calibration, timestamps, trajectory/gravity, semidense points with support or uncertainty, EVL/EFM evidence, and detected or predicted OBBs; GT mesh and GT OBBs are excluded as actor inputs.

CF0 state – Minimal Counterfactual Actor State: Semantic actor-visible target for counterfactual planning. It must update causally from selected observations and preserve distinctions that can change target-specific return, including observed surface, observed free, unknown, finite support, uncertainty, source, and recency. It does not prescribe points, voxels, or rays as the carrier.

NBV – Next-Best-View: Problem of selecting the next sensor viewpoint to improve an active reconstruction or inspection objective under a limited acquisition budget. In ARIA-NBV, the preferred objective is reconstruction quality improvement rather than coverage alone, with candidate views ranked by oracle or learned RRI scores.

oracle state – Oracle Rollout State: Privileged state used only for ASE mesh-supervised label generation, upper-bound planning, and evaluation. It includes GT mesh and target crops, all-candidate renders, points, normals, mesh-face visibility, RRI and point-mesh error metrics, oracle scores, and GT OBBs.

offline state – Persisted Offline Sample State: Compact VIN offline-store state persisted for training and diagnostics. It contains the VinSnippetView payload, candidate table and cameras, candidate count, validity or count metadata, oracle metrics as labels, optional rendered depths, compact OBBs, trajectory metadata, and selected EVL numeric fields.

S0-pose – Pose-Only Counterfactual State: Implemented selected baseline $s_t^{S0\text{-pose}}$. It combines immutable root evidence, target and candidate geometry,

the strictly causal selected-pose prefix $\mathbf{H}_t^{\text{pose}}$, and remaining budget, but it does not update scene geometry from selected observations.

state – Rollout State: Family of rollout states used by ARIA-NBV. State source, dynamic carrier, implementation maturity, and evidential status are independent: an implemented privileged control is not an actor-visible scientific target, and a planned target state is not an empirical result. Oracle geometry and all-candidate labels remain supervision or evaluation only.

NBV MDP – Target-Conditioned NBV MDP: Finite-horizon Markov decision process used to define ARIA-NBV’s target-conditioned rollout and fitted Q_H training contract. The actor observes only the declared reconstruction state, sampled candidate views, action support, target descriptors, history, and budget state; the provenance of action support is part of the protocol. Oracle meshes, GT crops, and mesh-derived validity are privileged signals unless explicitly admitted as a named control.

mask – Validity Mask: Hard candidate-level feasibility mask used during candidate ranking, rollout generation, and Q_H training. The scalar mask $m_{t,i}$ gates selectable rows, while the reason code $\rho_{t,i}$ records why a row was rejected. Invalidity is represented separately from reconstruction quality and should not be encoded as the lowest RRI class.

Protocol

GT-EVAL – Ground-Truth Target Evaluation: Main thesis protocol component restricting ground-truth OBBs and target mesh crops to oracle labels, target-RRI computation, matching checks, and evaluation. GT target crops are not actor-visible inputs in the main result.

OBS-SEL – Observed Target Selection: Main thesis protocol component requiring target selection to use only actor-visible evidence such as predicted or tracked OBBs, class probabilities, confidence, projected area, and semidense or EVL support. Ground-truth target annotations are not visible to the selector in this protocol.

PRED-Q – Predicted-Target Q: Main thesis protocol component requiring the scorer or fitted Q_H model to condition on predicted or observed target descriptors rather than ground-truth target inputs. It covers target-conditioned one-step scoring and finite-candidate Q_H selection.

Reconstruction

MVS – Multi-view Stereo: Reconstruction family that estimates dense or semi-dense scene geometry from multiple calibrated views. MVS is part of the broader reconstruction context for NBV, although ARIA-NBV's current oracle labels are built from ASE meshes and Aria-style semi-dense point observations.

PC – Point Cloud: Set of 3D points representing observed scene geometry. ARIA-NBV compares current and candidate-fused point clouds against ground-truth meshes to compute oracle RRI and surface-distance diagnostics.

SLAM – Simultaneous Localization and Mapping: Method for estimating sensor motion while building a map of the surrounding scene. In ARIA-NBV, logged SLAM poses and semi-dense points form the current reconstruction state against which candidate-view RRI is evaluated.

track – Track: Temporal sequence of corresponding image-feature detections across frames, usually carrying per-frame image coordinates, timestamps, and camera IDs. Tracks can be triangulated or optimized into 3D points, making them a bridge between image evidence and the semi-dense point clouds used by ARIA-NBV.

VIO – Visual-Inertial Odometry: Pose-estimation method that combines visual measurements from cameras with inertial measurements from IMUs. Project Aria pose and mapping products build on visual-inertial estimation, and ARIA-NBV depends on those pose streams for snippet state and candidate frame contracts.

Representation

***Occupancy Grid* – Occupancy Grid:** Spatial grid whose cells encode whether space is occupied, free, unknown, or represented by a related occupancy probability. Occupancy-style voxel evidence appears in EVL outputs and VIN feature construction, where it helps summarize local 3D scene state for candidate scoring.

***3DGS* – 3D Gaussian Splatting:** Explicit radiance-field representation using optimized 3D Gaussian primitives for real-time novel-view synthesis. ARIA-NBV treats active 3DGS work as proposal, uncertainty, and simulator-bridge literature, not as a replacement for ASE mesh-supervised RRI.

Supervision

***oracle RRI* – Oracle RRI:** RRI label computed with privileged ground-truth geometry, used for supervised training and evaluation. The current oracle renders candidate depth maps from ASE ground-truth meshes, backprojects candidate point clouds, fuses them with the current semi-dense reconstruction, and scores the resulting surface-distance improvement.

List of Symbols

Symbol	Meaning	Key
Oracle and Reconstruction		
RRI	Relative Reconstruction Improvement score for a candidate view.	oracle.rri
$\mathcal{P}$	Actor-visible point set accumulated from observed or replayed views.	oracle.points
$\mathcal{P}_q$	Candidate-view point contribution used in oracle or counterfactual updates.	oracle.points_q
$\mathcal{Q}$	Finite candidate-view set considered by the planner.	oracle.candidates
$\mathcal{Q}_t$	Candidate-view set available at rollout step t.	oracle.candidates_t
$q_{t,i}$	Candidate pose i at rollout step t.	oracle.candidate_qti
$\mathbf{c}$	World-space camera-center vector; subscripts identify a rollout root or candidate pose.	oracle.center
D_q	Oracle-rendered or candidate-specific depth map for a proposed view.	oracle.depth_q
$D_{P \rightarrow M}$	Point-to-mesh directional error component.	oracle.acc
$D_{M \rightarrow P}$	Mesh-to-point directional error component.	oracle.comp
D	Aggregate point-mesh reconstruction error used by RRI definitions.	oracle.err
ASE Assets		
$\mathcal{M}^{\text{GT}}$	Ground-truth scene mesh used for oracle labels and evaluation.	ase.mesh
$\mathcal{M}_e^{\text{GT}}$	Ground-truth mesh crop for target e.	ase.mesh_target
$\mathcal{F}^{\text{GT}}$	Ground-truth mesh faces used in point-mesh distance calculations.	ase.faces
$\mathbf{T}_{\text{rig}}^{\text{w}}(t)$	Project Aria rig pose in the world frame at time t.	ase.traj
$\mathbf{T}_{\text{rig}}^{\text{w}}(T)$	Final rig pose used as an anchor for gravity-aligned local fields.	ase.traj_final
VIN Scorer		
r	Target or scene RRI value used as a model target.	vin.rri
$\hat{r}$	Predicted RRI value emitted by the learned scorer.	vin.rri_hat
$\mathcal{L}$	Training loss for scorer or value-model objectives.	vin.loss
m	Candidate validity indicator used by scorer diagnostics.	vin.cand_valid

Symbol	Meaning	Key
Entity and Target		
RRI_e	Target-specific Relative Reconstruction Improvement for entity e.	entity.rri_e
$RRI_{t,i}^e$	State-relative marginal target RRI for candidate i at rollout step t.	entity.target_rri_marginal
$C_t^{RRI,e}$	Running sum of selected state-relative target RRIs.	entity.target_rri_cumulative
J_t^e	Running cumulative target gain normalized by rollout-root error.	entity.target_root_gain_cumulative
Δ_t^e	Target-specific reconstruction error at rollout step t.	entity.target_error
φ_e	Actor-visible target/entity descriptor used to condition target-specific value prediction.	entity.target_desc
p_e^w	World-space center of the selected target entity e.	entity.center
Planning and Rollout		
s	Abstract planning state.	rl.s
a	Action index selecting a candidate row.	rl.a
r	Scalar reward or immediate gain.	rl.r
G	Finite-horizon return accumulated across rollout steps.	rl.G
j	Factual rollout-chain index used to preserve common trajectory heritage across selected steps.	rl.rollout_index
$\mathcal{M}_{NBV}$	Target-conditioned finite-candidate NBV decision process.	rl.mdp_nbv
$\mathcal{A}(s_t)$	Valid action set available in state s _t .	rl.action_set
T	State-transition operator for selected candidate actions.	rl.transition
s_t^{hist}	Logged historic snippet state at rollout step t.	rl.s_hist
s_t^{off}	Persisted offline sample state used by training readers.	rl.s_off
s_t^{cf0}	Minimal actor-visible counterfactual state.	rl.s_cf0
$s_t^{\text{S0-pose}}$	Implemented fixed-root pose-history state used by qh_cf0_v1.	rl.s_pose
$s_t^{\text{S1-surface}}$	Implemented privileged selected-surface control state; it consumes only previously selected rendered depth and is not deployable.	rl.s_surface
$s_t^{\text{S2-ray}}$	Planned actor-visible ray-aware state preserving observed surface, free, unknown, support, uncertainty, source, and recency.	rl.s_ray

Symbol	Meaning	Key
$s_t^{\text{cf}+}$	Geometry-enriched counterfactual successor state.	rl.s_cf_geom
$s_t^{\text{CF-GT-carrier}}$	Implemented privileged selected-depth carrier used by qh_cfplus_gt_depth_v1.	rl.s_cf_gt_carrier
s_t^{oracle}	Oracle-only state used for labels, upper bounds, and evaluation.	rl.s_oracle
H_t^{pose}	Strictly causal prefix of previously selected poses available at decision step t.	rl.selected_pose_prefix
r_t^e	Target-specific reward at rollout step t.	rl.reward_target
Entity and Target		
r_t^e	Canonical target-specific reward at rollout step t.	entity.target_reward
Planning and Rollout		
$G_t^{(H)}$	H-step finite-horizon return from rollout step t.	rl.return_h
Q_H	Finite-horizon candidate-value function.	rl.qh
$Q_{h,\theta,e,i}^{\text{cond}}$	Action-mask-independent conditional candidate value emitted by the scorer.	rl.conditional_q
$\ell_{t,i}^{\text{feas}}$	Physical or observed feasibility logit emitted by the scorer's feasibility head.	rl.feasibility_logits
e_k^Q	Fixed continuous fitted-Q boundary between adjacent CORAL classes.	rl.coral_q_edge
u_k^Q	Fixed continuous-Q representative used to decode one CORAL class.	rl.coral_q_value
c_n^Q	Ordinal class assigned to one continuous fitted-Q target.	rl.coral_q_label
γ	Return discount factor.	rl.gamma
H	Planning or rollout horizon length.	rl.H
H_{max}	Maximum residual horizon admitted by a scorer/data contract.	rl.H_max
h	Scalar residual horizon requested for one scorer query; it must satisfy $1 \leq h \leq h_{\text{max}}$.	rl.requested_horizon
$m_{t,i}$	Hard validity mask for candidate i at rollout step t.	rl.validity_mask
$m_{t,i}^{\text{cand}}$	Materialized candidate row versus structural padding.	rl.candidate_row_mask
$m_{t,i}^{\text{act}}$	Authoritative physically valid action support.	rl.action_mask
$m_{t,i}^{Q,h}$	Availability of a finite factual value target for candidate i at requested horizon h.	rl.q_label_mask

Symbol	Meaning	Key
$m_{t,i}^F$	Availability of a trustworthy feasibility label for candidate i.	rl.feasibility_label_mask
m_t^{succ}	Availability of a factual successor backup for transition t.	rl.successor_mask
$\ell_{t,i}^{\text{src}}$	Actor-oracle source-provenance role for candidate evidence.	rl.source_role
$\rho_{t,i}$	Invalid-action reason code for candidate i at rollout step t.	rl.invalid_reason
e_t	Selected target entity at rollout step t.	rl.target
b_t	Remaining acquisition budget at rollout step t.	rl.budget
$C(\tau)$	Acquisition cost of a selected trajectory.	rl.acquisition_cost
Shape and Size		
N_q	Number of candidate views in a candidate table.	shape.Nq
K	Number of ordinal bins or discrete levels when used in scorer losses.	shape.K
scene		
M_t^{ray}	Sparse actor-visible ray-aware memory separating occupied, free, and unknown support at step t.	scene.ray_memory_t
Observation and Actor State		
$F_t^{\text{DINO@pt}}$	Visibility-gated logged DINO feature bank attached to semidense or fused world points.	obs.dino_point_bank_t
X_t^{pt}	Point-token set combining geometry, support, uncertainty, history, and optional compressed logged descriptors.	obs.point_tokens_t
scene		
Φ_t^{scene}	Composite actor-visible scene-memory descriptor queried by the finite-candidate value model.	scene.scene_memory_t
$E_0^{\text{EVL-local}}$	Root local EVL evidence field used for target/candidate support reads.	scene.evl_local
$\omega_{t,i}^{\text{EVL}}$	Candidate or target-query fraction supported by the root EVL voxel extent.	scene.evl_support_frac
$g_{t,i}^{\text{EVL}}$	Pooled local EVL evidence token for a target, candidate frustum, or target-candidate intersection query.	scene.evl_support_token
g_e^{tgt}	Pooled target-support descriptor for the selected target hypothesis.	scene.target_support_pool
$g_{t,i}^{\text{fr}}$	Candidate-frustum support descriptor pooled from actor-visible scene memory.	scene.frustum_support_pool

Symbol	Meaning	Key
$\mathbf{g}_{t,e,i}^{\text{cap}}$	Pooled descriptor over the intersection between target support and candidate frustum support.	scene.target_frustum_pool
$\mathbf{g}_{t,i}^{\text{ray}}$	Candidate-conditioned ray query over occupied, free, unknown, hit, target-support, and uncertainty summaries.	scene.ray_query_ti
RenderQuery	Abstract query operator that summarizes scene memory in a candidate camera/frustum without introducing fresh counterfactual RGB features.	scene.render_query
spatial		
r_t	Reference pose for candidate-relative descriptor construction at decision step t.	spatial.ref_pose
$\mathbf{T}_{r_t,i}^{\text{rel}}$	Relative transform from the reference pose r_t to candidate camera i.	spatial.ref_candidate_transform
$\mathbf{h}_{t,i}^{\text{pose}}$	Relative/local candidate pose descriptor derived from a reference pose rather than raw world coordinates.	spatial.candidate_pose_feat
$\mathbf{h}_{t,e i}^{\text{rel}}$	Candidate-target relation descriptor in the candidate/query local frame.	spatial.candidate_target_rel_feat
$\mathbf{e}_{a i}^{\text{rel}}$	Query-local relative positional embedding for target, history, support, or candidate relations.	spatial.relation_rpe
r_e	Geometric-mean target-proxy scale derived from the target OBB semi-axis lengths.	spatial.target_obb_scale
$\hat{\delta}_{j,t}^e$	Unit factual selected-camera displacement expressed in target-object coordinates for rollout j and step t.	spatial.target_frame_motion_direction
$\hat{\mathbf{v}}_{j,t}^e$	Unit selected-camera forward optical axis expressed in target-object coordinates.	spatial.target_frame_view_direction
$\mathbf{M}^{\text{dir}}$	Second directional moment used to summarize selected target-relative approach directions.	spatial.dir_moment
$\mathcal{F}_{j,t}^e$	Front-facing target-proxy surface directions that project inside the calibrated selected-camera image.	spatial.target_frame_frustum
$\kappa_{j,t}^e$	Fraction of the target-centred proxy sphere geometrically supported by one selected calibrated frustum.	spatial.target_frame_frustum_fraction
model		
$\mathbf{h}_e^{\text{tgt}}$	Learned selected-target token built from root-relative target pose and metric extents.	model.target_token
spatial		

Symbol	Meaning	Key
$\Omega_{j,t}^{\text{FOV}}$	Intrinsic solid angle of one calibrated selected-camera pinhole field of view.	<code>spatial.frustum_solid_angle</code>
model		
$\mathbf{h}_{t,i}^{\text{phys}}$	Candidate-local physical token from root/current-relative candidate poses and compact root evidence.	<code>model.candidate_physical_token</code>
$\mathbf{x}_{t,i}$	Candidate-local conditional-value query formed by adding candidate-from-target pose encoding to the physical token.	<code>model.candidate_row</code>
$\mathbf{p}_{t,j}^{\text{hist}}$	Previously selected pose j encoded from the current camera at decision state t.	<code>model.history_pose_feature</code>
$a_{t,j}^{\text{hist}}$	Normalized relative age of selected pose j at decision state t; the immediate predecessor has age zero.	<code>model.history_relative_age</code>
$\mathbf{h}_t^{\text{hist}}$	Fixed-width causal selected-pose history token supplied to scorer state fusion.	<code>model.history_token</code>
Entity and Target		
$J_e^{(H)}$	Endpoint reconstruction gain for target e over horizon H.	<code>entity.endpoint_gain</code>
$J_{e,\log}^{(H)}$	Log-scale endpoint reconstruction gain for target e.	<code>entity.log_gain</code>
Δ_{look}	Gain available to an oracle lookahead policy beyond one-step selection.	<code>entity.lookahead_headroom</code>
η_Q	Fraction of the learned-myopic-to-oracle-lookahead endpoint gap closed by the Q policy.	<code>entity.q_recovery</code>
$G_t^{(H)}$	Target-conditioned finite-horizon return.	<code>entity.return_h</code>
Observation and Actor State		
$\mathbf{D}$	Depth observation sequence.	<code>obs.depth</code>
$\mathbf{n}$	Per-point or per-face normal observation.	<code>obs.face_normal</code>
$\mathbf{I}^{\text{rgb}}$	RGB image observation.	<code>obs.img_rgb</code>
$\mathcal{P}^{\text{cf}}$	Counterfactual point observation.	<code>obs.points_cf</code>
$\mathcal{P}^{\text{semi}}$	Semi-dense observed point set.	<code>obs.points_semi</code>
$\mathcal{P}_t$	Point set available at rollout step t.	<code>obs.points_t</code>
Planning and Rollout		
$\mathcal{A}_t$	Finite feasible action set at rollout step t.	<code>rl.action_set_t</code>
$\mathcal{Q}_t$	Finite candidate-view table at rollout step t.	<code>rl.candidate_table</code>
$y_t^{(2,\text{exact})}$	Exact two-step fitted-Q control using factual dense successor one-step rewards.	<code>rl.exact_q2_target</code>

Symbol	Meaning	Key
$\varepsilon_t^{(2)}$	Absolute learned-recursion target error against the exact two-step control.	rl.q2_recursion_error
VIN Scorer		
F_v	Learned value field evaluated at v.	vin.field_v

List of Figures

- Figure 1** Information roles in the offline experiment. The actor side contains calibrated logged observations, derived EFM3D evidence, an explicitly sourced target instruction, the selected-view prefix, remaining budget, candidates, and the hard action mask. The oracle side adds GT geometry, counterfactual renders, labels, and endpoint evaluation. Every stored field keeps its declared information role. . . . 20
- Figure 2** Oracle target-task construction. Geometry-valid GT OBB rows form the task population; seeded uniform sampling without replacement applies the manifest-defined cap. Candidate scoring later determines whether the selected target crop is evaluable. 25
- Figure 3** One stored candidate shell from ASE scene 81286 and sample ASE_81286_Atek_000035. The perspective panel relates the shell to the logged camera history, target OBB, and selected prefix. The bird’s-eye panel retains all 60 candidate centres: 25 are feasible, 35 fail clearance, and oracle-greedy selects row 47. The example’s evidential role is proposal support and hard feasibility. 27
- Figure 4** Conceptual topology of bounded oracle lookahead. Nodes are counterfactual states; edges are feasible actions labelled by immediate reward. The beam retains only a bounded set of prefixes. The symbolic inequalities show how a lower immediate reward can lead to a larger finite-horizon return. 33
- Figure 5** Geometry of the operational point–mesh error. Panel A isolates the point-to-triangle primitive; panel B separates

point-to-face accuracy from face-to-point completeness.
 Panel C keeps the surface and point set fixed but changes
 its tessellation, changing equal-face completeness from
 0.03640 to 0.02284 m². The synthetic fixture provides
 evidence for metric sensitivity. 35

Figure 6 Normalized replay lineage. A source defines target tasks;
 each task produces recipe-specific retained chains; each
 chain orders factual decision states; and each state owns
 one complete candidate shell. Model-ready tensors are
 derived by padding and batching these relations. 40

Figure 7 Evidence-gated claim graph. A failed stage retains its own
 measured outcome but blocks claims that depend on it;
 missing evidence remains unavailable rather than becoming
 zero. 67

Figure 8 Candidate-family survival for one deterministic scene from
 each of the ten persisted Phase-A audit strata. Each cell
 reports compact-valid-shell membership divided by
 attempted rows for one factual state and proposal family;
 the complete 100-state matrix remains in the hash-bound
 evidence bundle. The underlying artifact geometry uses
 correction revision **phase-a-target-aligned-z-up-v1**; the stored
 shells, support counts, verdict, and execution identity are
 unchanged. The artifact is bound to clean execution
 revision **2baf7cf6b276b81c50d01d45b152016d7cf68033** and
SHA-256
6d33e9e3d68737c8a6a5589ae5117c1e4d7fcaa89056fcfaec1d315e4509c83. 77

Figure 9 Critical development gates from experiment substrate to
 the 30 November 2026 full-draft freeze. 86

Figure 10 Attempted candidate centres from a two-scene, two-state
 real-data method audit. Each state is centred on its factual

expansion pose, yaw-aligned to the oracle-evaluation target with world Z-up, and divided by its current three-dimensional target distance. The cross marks the expansion/root pose, the stars mark the two oracle task target centres, colour denotes proposal family, and marker shape denotes actor-valid or selected status. The figure diagnoses proposal geometry under the `v0_gt_input` target protocol; it is neither a population estimate nor evidence of actor-visible target discovery or policy quality. 89

Figure 11 Target-frame $\mathcal{S}^2$ movement directions from the frozen pilot report. This view records how the camera moved; it is not a visibility measurement. 90

Figure 12 Target-frame $\mathcal{S}^2$ selected-camera optical-axis directions from the same frozen pilot report. This view records where the camera pointed; it is distinct from both camera motion and calibrated surface support. 91

Figure 13 Calibrated selected-frustum support on the target-centred proxy sphere. The surface heat map is the complete equal-solid-angle coverage field; overlaid incidence samples preserve rollout and time heritage. This is not measured target-mesh visibility. 92

Figure 14 Development-only directional-memory fixture. The sphere and Mollweide panels show the same logged ASE camera directions in a target-object frame; the orange query is another logged pose used to inspect one prospective second-moment novelty definition. No smooth field, spherical-harmonic representation, or learned benefit is claimed. 104

List of Tables

Table 1	Research-question-to-evidence map.	7
Table 2	Concept-centred comparison of the literature families that bound the thesis question. Rows are distinguished by utility, target scope, state evidence and representation role, candidate support and feasibility, and decision structure; they are not ranked by reported performance across incompatible settings.	16
Table 3	Geometric roles of the three families in the frozen <code>realistic_core_60</code> intervention. Counts and radii are specific to this profile.	26
Table 4	Candidate-generation diagnostics and their aggregation populations. Every state-level quantity is reduced before scene and cohort aggregation.	32
Table 5	Reconstruction measures answer different geometric questions. The equal-face point–mesh error is the selected outcome; the other families are diagnostics or alternative operationalizations.	36
Table 6	Scientific roles of the normalized replay relations. State, target, action support, supervision, and lineage are bound before model-specific tensorization.	40
Table 7	Interpretive status for Method. Scientific role, implementation maturity, and evidence maturity are orthogonal; no row is an empirical ranking.	43
Table 8	Current baseline and scientific target. The right column is a promotion contract: it states information that must be	

	represented and tested, not an architecture already implemented or a result already obtained.	46
Table 9	Representation design space ordered by scientific role, not presumed performance. A carrier is promoted only when it tests a diagnosed information loss under the same task, action, replay, and evaluation contracts.	48
Table 10	Logical data roles at the reader–model boundary. Storage readability and batch colocation do not imply actor visibility.	50
Table 11	Development-only architectural feature-integration ladder. Status records implementation maturity, while promotion remains conditional on a named scientific failure and a one-factor comparison.	59
Table 12	Value-model alternatives by scientific role and estimand. Executability, selection, and evidence are distinct; alternatives are retained without presenting them as co-equal validated methods.	64
Table 13	Learning-target admission. Dense one-step labels, exact two-step targets, and general recursive targets are different evidence populations.	71
Table 14	Matched policy roles. Oracle access defines a privileged reference, not a deployable upper bound; every row retains the same acquisition budget.	72
Table 15	Failure-attribution matrix. Interpretation stops at the earliest failed prerequisite for each claim; downstream architecture stories remain hypotheses.	73
Table 16	Evidence status by inferential stage. “Failed” retains a measured negative outcome, “blocked” marks a failed prerequisite, and “not available” preserves missingness. ..	75

Table 17 Evidence-gated schedule from the completed pilot
infrastructure to submission. Dependencies are checked
against the adjacent roadmap data. 86

Table 18 Loaded development-fixture provenance. The fixture verifies
the document interface only and is not scientific
evidence. 102

1 Introduction

Active perception begins from a geometric premise: what can be reconstructed depends on what a sensing action makes visible [AWB88, Baj88]. A new view changes both the available surface evidence and the state from which later views are chosen. Next-best-view (NBV) planning turns this coupling between action, observation, and reconstruction into the repeated choice of where to look next [Jia+25, SRR03].

With a limited acquisition budget, a view is “best” only relative to an objective. Coverage rewards newly observed surface, information criteria reward reduced uncertainty, and quality-driven methods evaluate the reconstruction itself [Fra+25, GML22, JLD24]. These objectives can disagree. A view may reveal more of a room without exposing the occluded surface of a requested object, while a small side-step may improve that object but add little scene-wide coverage. Target conditioning therefore specifies **whose geometry** should improve; it is not generic semantic awareness.

Existing methods supply the necessary ingredients under different assumptions. VIN-NBV ranks sampled candidates by direct reconstruction improvement, object-aware 3D Gaussian Splatting conditions an information criterion on a requested object, and GenNBV and Hestia learn sequential coverage policies [Che+24, Fra+25, Jeo+26, Lu+26]. What remains unresolved is their conjunction: whether target-specific reconstruction quality contains useful multi-step structure and whether that structure can be recovered from a bounded egocentric information state.

The actor-state question includes a representation-transfer hypothesis. Prior NBV pipelines already use generic 2D pretraining or foundation-model components for view prediction, depth estimation, or scene construction [AKC20, Gué+23, Str+24a]. These uses do not establish whether an egocentric 3D representation that combines a frozen 2D foundation feature extractor with learned upsampling and 3D processing improves target-conditioned endpoint-value learning. ARIA-NBV therefore treats EFM3D/EVL evidence as a hypothesis to be

tested against matched actor-visible geometric controls, not as an assumed source of gain [Str+24b].

Project Aria supplies calibrated egocentric observations, EFM3D lifts them into local 3D evidence, and ARIA Synthetic Environments (ASE) provides the privileged geometry needed for controlled target tasks and counterfactual evaluation [Eng+23, Met25, Str+24b]. This combination makes the information boundary testable. Task construction, action-set construction, hard admission, supervision, and evaluation may use privileged geometry, while the scorer consumes only the fields admitted by its declared actor-facing protocol. Scientific interpretation must account for both paths. Chapter 3 therefore treats actor visibility as a property of the complete decision process, not merely of the scorer tensor.

1.1 Aim and bounded problem

The evaluation proceeds in two stages. First, it measures whether bounded oracle lookahead yields better fixed-budget endpoint reconstruction of a requested object than one-step oracle-greedy selection. Conditional on such headroom, it then measures how much of the endpoint gap from an actor-visible learned-myopic control to oracle lookahead an offline finite-horizon model closes. This order separates the existence of a planning opportunity from the ability of a learned model to exploit it; the gap-closure ratio and oracle headroom have different baselines.

At every step, each policy selects from the same finite generated candidates; hard-invalid actions are excluded, and the target, horizon, candidate support, and endpoint evaluation are matched. The current oracle tasks use geometry-valid ground-truth boxes, while actor-visible target discovery and matching remain separate research requirements. The objective adapts VIN-NBV rather than copying it: VIN-NBV defines RRI from point-cloud Chamfer distance [Fra+25], whereas ARIA-NBV evaluates a target-cropped point-mesh error after an equal acquisition budget. The resulting claims concern this bounded finite-candidate setting, not a deployable target-selection system or an objective invariant to other reconstruction protocols.

1.2 Contributions and current evidence

The thesis operationalizes this evaluation through four scientific functions:

1. It separates one-step RRI, finite-horizon return, and fixed-budget endpoint gain within target-specific reconstruction quality.
2. It defines an end-to-end information boundary across task and action construction, state updates, model inputs, supervision, and evaluation.
3. It defines candidate value relationally over target, causal state, admissible action, budget, horizon, and continuation rule.
4. It evaluates measurement, support, oracle headroom, immediate-value learning, recursion, and learned-control endpoint gap closure in that order.

Before model capacity becomes the main question, the evidence needed to define the learning problem must itself be identifiable. The outcome must be repeatable, the study population and candidate support must contain the relevant opportunity, bounded lookahead must exhibit headroom, and factual replay must identify the required targets. Only after those conditions hold can a failure to recover non-myopic value be attributed primarily to representation or learning. The order makes a negative result informative: it distinguishes a measurement or support failure from a failure of immediate-value learning or finite-horizon recursion.

The present evidence supports a methodological contribution rather than a policy result. The implementation executes target-specific oracle scoring and selected-action replay, while rendering memory limits scale. Actor-visible target matching, metric repeatability, a validated held-out population, oracle-lookahead headroom, and paired policy outcomes remain unestablished. Accordingly, the thesis does not yet claim policy superiority or deployment readiness.

1.3 Research Questions

1.3.1 Research aim and evaluation logic

The study evaluates a sequence of linked questions. RQ1 defines the endpoint outcome. RQ2 then asks whether bounded oracle lookahead improves that outcome relative to one-step oracle-greedy selection and, conditionally, how much of

the separate actor-visible myopic-to-oracle-lookahead gap an offline model closes. RQ3 defines the admissible information, and RQ4 establishes the target, action, and replay support over which both comparisons are interpretable. RQ5 and RQ6 extend the study only if the offline finite-candidate evidence warrants online or continuous-control claims.

1.3.2 RQ1 — Target-specific objective and endpoint outcome

Can target-conditioned finite-candidate NBV be evaluated through a stable, target-specific reconstruction objective under a fixed acquisition budget?

For a target e , the oracle measures reconstruction error on a frozen target crop. The primary estimand is the quality change from the root state to the endpoint after H acquisitions:

$$J_e^{(H)} = \frac{\Delta_0^e - \Delta_H^e}{\Delta_0^e + \varepsilon}$$

State-relative RRI is a one-step diagnostic, root-normalized gain supplies the training reward, and endpoint gain remains the policy outcome. A positive answer requires a frozen, repeatable metric and equal acquisition horizons; runtime, invalid actions, and oracle calls are reported separately.

1.3.3 RQ2 — Bounded lookahead and learned gap closure

Does bounded oracle lookahead improve fixed-budget endpoint target gain over one-step oracle-greedy selection and, if so, how much of the actor-visible myopic-to-oracle-lookahead endpoint gap does an offline finite-horizon value model close?

The paired endpoint difference under the same scene, target, candidates, validity rules, horizon, and budget defines oracle-lookahead headroom:

$$\Delta_{\text{look}} = J_e^{(H)}(\pi_{\text{oracle-look}}) - J_e^{(H)}(\pi_{\text{oracle-1}})$$

The learned comparison proceeds only if this headroom passes a predeclared meaningful-effect and uncertainty rule. Its ratio uses a different baseline:

$$\eta_Q = \frac{J_e^{(H)}(\pi_Q) - J_e^{(H)}(\pi_{\text{learned-1}})}{J_e^{(H)}(\pi_{\text{oracle-look}}) - J_e^{(H)}(\pi_{\text{learned-1}}) + \varepsilon}$$

Here the numerator is the gain of the finite-horizon learned policy over the actor-visible learned-myopic control, and the denominator is the full gap from that control to oracle lookahead. The ratio is therefore conditional learned gap closure, not a fraction of the oracle-lookahead headroom above oracle-greedy. It is reported only when the paired denominator $J_e^H(\pi_{\text{oracle-look}}) - J_e^H(\pi_{\text{learned-1}})$ exceeds a predeclared positive tolerance $\tau_{\text{gap}} > 0$; otherwise the analysis reports the raw paired policy effects and classifies gap closure as not estimable. Because the headroom rule and required gap-closure fraction are not yet frozen, RQ2 remains prospective. If headroom is absent, the evaluated support does not expose a non-myopic advantage; this does not establish its universal absence.

1.3.4 Conditions for interpretation

1.3.4.1 RQ3 — Actor-visible target and information state

Which end-to-end target, action-support, and history protocol supports one-step and finite-horizon candidate scoring without privileged target geometry, oracle labels, or unselected candidate renders at decision time?

Ground-truth target tasks may define supervision, evaluation, and a separately reported privileged **V0/GT** task-construction sanity control; they do not satisfy the core RQ3 gate. The core learned comparison requires an observation-derived target instruction and matching path (**V1/observed**) in addition to actor-visible candidate support, hard validity, selected-state updates, and scorer inputs. The **v1_observed** admission and writer–reader route exists, but no frozen scene-disjoint observed-target corpus, causal observation-updated state, actor-visible action-support protocol, or end-to-end evidence bundle yet supports the core comparison; it therefore remains unavailable. The scorer then receives only the declared target descriptor, causal egocentric evidence, remaining budget, and

requested horizon; evaluation separately audits matching failures, ranking and calibration, and actor/oracle leakage.

1.3.4.2 RQ4 — Candidate, replay, and population support

Do the candidate generator, validity rules, rollout recipes, and replay population provide adequate and diverse support for the RQ1–RQ3 estimands?

The candidate table combines exploration and target-bearing views, records generator provenance, and excludes geometric infeasibility through a hard mask. RQ4 measures candidate-family survival, invalidity, scene and target coverage, replay and horizon support, and resource cost. Scene-disjoint splits and scene-level aggregation prevent dense sampling from masquerading as population coverage.

1.3.5 Scope extensions

1.3.5.1 RQ5 — Online discrete bridge

If offline headroom, replay support, and actor-visible scoring are established, does online interaction over the unchanged discrete candidate set and validity rules improve endpoint target gain or calibration over the offline policy?

RQ5 keeps the target, information, candidate, validity, and endpoint assumptions of RQ1–RQ4 unchanged. It is not required for the offline finite-candidate evaluation.

1.3.5.2 RQ6 — Continuous or simulator-backed control

If the finite-candidate evidence is stable, does a continuous or hierarchical target-then-pose policy provide measurable headroom over the best discrete policy under the same target-specific objective and a comparable acquisition or motion budget?

RQ6 changes the action space and feasibility assumptions and therefore requires separate evidence for support, safety, simulator realism, and comparable cost. It remains deferred; the current implementation implies no continuous, simulator, or real-device result.

1.3.6 Shared evidence constraints

Action validity, actor support, target validity, oracle-label validity, and training eligibility remain distinct; missing evidence is not encoded as low utility. Confirmatory comparisons are paired under the same scene, target, candidates, validity regime, and budget. The scene is the experimental unit, and the analysis plan freezes exclusions, aggregation, uncertainty, and decision rules before policy outcomes are inspected.

1.3.7 Research-to-evidence map

RQ	Question role	Primary evidence	Interpretation gate
RQ1	outcome measurement	target reconstruction endpoint gain	frozen repeatable metric; fixed horizon and budget
RQ2	lookahead and learned gap closure	paired greedy, lookahead, myopic- control, and learned-policy outcomes	meaningful oracle headroom before learned gap closure
RQ3	information boundary	matching, ranking, calibration, and leakage audits	end-to-end actor-visible protocol
RQ4	population and support	candidate, replay, validity, and coverage diagnostics	scene-disjoint aggregation
RQ5	conditional extension	matched online discrete-policy evaluation	offline gates satisfied first
RQ6	deferred extension	continuous or simulator-backed evaluation	separate action space and cost comparison

Table 1: Research-question-to-evidence map.

1.4 Thesis structure

The remaining chapters follow that dependency chain. Chapter 2 derives the objective, support, temporal, and representation concepts needed to state the gap precisely. Chapter 3 constructs the experimental world by defining information access, admissible actions, state transitions, and measured outcomes. Chapter 4 then specifies the one finite-horizon value interface evaluated within that world. Chapter 5 states which evidence admits or blocks each claim. Chapter 6 reports the admitted answers in gate order; Chapter 7 attributes their implications and alternatives; and Chapter 8 synthesizes what those bounded answers establish.

2 Foundations and Related Work

The Introduction located this thesis within active perception: sensing actions change the evidence on which later inference depends [AWB88, Baj88]. Three-dimensional view planning turns that principle into a repeated choice of where to observe next [SRR03]. The unresolved dependency is not another planner implementation but a precise account of what a view is valuable **for**, which actions are available, and what information a sequential score may condition on. By the end of the chapter, the thesis question is reduced to three coupled requirements—utility alignment, temporal dependence, and support/state adequacy—before the experimental world is constructed.

The argument follows the dependencies of the decision problem. It first separates the next-best-view mechanism from the objective used to rank views, then distinguishes target-conditioned reconstruction quality from coverage and uncertainty proxies. It next separates candidate-view support from endpoint, transition, and human-motion feasibility before explaining why partial observability and delayed consequences require a finite-horizon information state. The chapter then derives the geometric properties a candidate scorer should preserve and positions the thesis question against the resulting literature dimensions.

2.1 Active Perception and the Next-Best-View Problem

In active perception, an observation is also an intervention on the future information state. A next-best-view (NBV) method operationalizes this idea by generating possible viewpoints, estimating the benefit of observing from each one, and selecting an action. PB-NBV makes this generate–score–select decomposition explicit for a finite candidate set [Jia+25]. The decomposition identifies three different scientific objects: the proposal mechanism determines which views can be considered, the utility determines how they are ordered, and the selection rule turns that ordering into an action.

How those objects are represented varies across NBV systems. Finite-candidate methods evaluate a bounded set of poses, whereas learned continuous policies

predict camera position and orientation directly. GenNBV formulates the latter as a five-degree-of-freedom reinforcement-learning problem, while Hestia factorizes the continuous action into a look-at point followed by a camera position [Che+24, Lu+26]. A finite set makes the evaluated choices explicit and comparable; a continuous or hierarchical policy can search a broader space but couples view proposal and control more tightly.

These action formulations do not determine what makes a view useful. PB-NBV, for example, accelerates candidate evaluation through projected frontier and occupied evidence, so its proposal and scoring mechanism is tied to a coverage objective [Jia+25]. The same finite candidate table could instead be ranked by uncertainty reduction or by reconstruction error. Consequently, changing the utility changes the scientific task even when candidate generation and selection remain unchanged.

The NBV mechanism therefore answers only **how** a sensing choice is made. To interpret the choice, the utility must state **which reconstruction consequence** is valued and **for which spatial support**. This distinction leads from the general active-perception loop to the objective families compared next; only after fixing that objective can candidate support and feasibility be separated from view value [Che+24, Jia+25].

2.2 View Utility and Target-Conditioned Reconstruction

Coverage-based utilities value evidence about previously unobserved surface. SCONE estimates surface-coverage gain from occupancy and visibility fields, MACARONS anticipates coverage gain while reconstructing online, and Hestia rewards newly visible voxel faces [GML22, Gué+23, Lu+26]. These objectives connect a view to geometric completeness, but they weight every counted surface element through the coverage definition; they do not directly ask whether the reconstructed geometry became more accurate.

Information-based utilities instead value a view through its expected reduction of model uncertainty. FisherRF derives an information-gain criterion from

the Fisher information of a radiance field [JLD24]. Such a criterion can prefer a view that constrains uncertain model parameters even when its immediate surface coverage is small. The benefit is model-aware exploration; the limitation is that lower parameter uncertainty and lower geometric reconstruction error are related only through assumptions about the model and the evaluated scene.

Direct quality utilities close this proxy gap by evaluating the reconstruction itself. VIN-NBV defines Relative Reconstruction Improvement (RRI) as the reduction in Chamfer distance between reconstructed and ground-truth point clouds after adding a candidate observation, and learns to rank sampled query cameras by this quantity [Fra+25]. The present work follows this direct reconstruction-improvement principle but not VIN-NBV’s metric unchanged: it evaluates a target-cropped point-mesh error instead of scene-level point-cloud Chamfer distance. This project-specific definition makes the requested target, rather than all reconstructed geometry, determine which errors count. It does not remove the myopia of one-step scoring, because each candidate is still valued by its immediate improvement rather than by the sequence it may enable.

Target conditioning changes which reconstruction errors contribute to that quality. Object-centric 3DGS planning associates features and confidence with individual Gaussian primitives and gates its information score by the requested object [Jeo+26]. This mechanism demonstrates that a view can be informative for one object without being equally useful for the scene as a whole. It also exposes a separate assumption: supplying an object identifier or region for conditioning is not the same problem as discovering that object and maintaining its identity from observations.

The thesis therefore treats coverage and uncertainty as non-equivalent diagnostics and adopts target-specific reconstruction quality as the endpoint of interest [Fra+25, Jeo+26]. These works motivate direct reconstruction improvement and object conditioning separately; their combination in a target-cropped endpoint is the thesis-specific objective. This choice determines what the decision should improve, but not which viewpoints may be proposed or reached. Candidate sup-

port and motion feasibility therefore remain separate from utility before delayed consequences are introduced through the finite-horizon decision problem [Jia+25].

2.3 Candidate-View Support and Motion Feasibility

A finite-candidate policy selects only from the pose rows in its proposed candidate table $\mathcal{Q}_t = \{q_{t,i}\}_{i=1}^{N_t}$. PB-NBV uses a reachability- and camera-conditioned partial hemisphere of target-facing candidates [Jia+25]. Proposal geometry therefore bounds what even a perfect scorer can select.

Target-relative coordinates provide a useful proposal prior without defining the objective. PB-NBV points candidates toward the object centre, while Hestia factorizes a viewpoint into a predicted look-at point and a camera position [Jia+25, Lu+26]. Such parameterizations concentrate support on a requested target, but an orbit or target-facing arc remains a coverage geometry, not evidence that a wearer normally moves that way. Project Aria explicitly contrasts carefully curated scanning motion with natural, unconstrained egocentric activity, so wearable plausibility must be checked against trajectory evidence rather than inferred from a target-relative construction [Eng+23].

Geometric admissibility has at least two levels. Hestia shifts a proposed camera position to its nearest collision-free endpoint, whereas Next Best Sense obtains a list of feasible candidate views but still falls back to the next-ranked view when inverse kinematics or trajectory planning fails [Lu+26, Str+24a]. A collision-free endpoint therefore does not by itself establish that the transition from the current state is executable.

Physical feasibility and wearable plausibility are likewise different requirements. Robot reachability constrains PB-NBV’s hemisphere, while Project Aria supplies calibrated device trajectories recorded during natural activities [Eng+23, Jia+25]. Robot-specific reachability can define a hard admissibility test for that platform; it cannot establish a distribution of realistic human head and body motion. For egocentric data generation, that distribution is an empirical prior to be measured, reported, and validated.

These distinctions map the proposed poses $q_{t,i} \in \mathcal{Q}_t$ to the canonical admissible row-index set $\mathcal{A}(s_t) = \{i : m_{t,i} = 1\}$, where $m_{t,i}$ is the hard-valid indicator after endpoint and transition checks. The utility or value function ranks only those admitted rows; a rejected candidate is unavailable rather than merely low-value [Jia+25, Str+24a]. Separating proposal, admission, and ranking makes the state dependence of available actions explicit [SB18, Secs. 3.1 and 3.5, pp. 47–53 and 58–62] and prepares the sequential decision model introduced next.

2.4 Sequential Decision-Making under Partial Observability

An egocentric reconstruction system never observes the complete physical scene. Its decision can depend on the sequence of earlier actions and observations, not only on the latest frame. In a partially observable Markov decision process, a belief state is a sufficient statistic of that action–observation history for policy choice [LHP23]. ARIA-NBV does not attempt general belief inference; the relevant consequence is narrower: any finite representation used for view selection is a hypothesis about which parts of the history are sufficient for predicting future target improvement.

The preceding section defined the admissible choices as the state-dependent row-index set $\mathcal{A}(s_t)$ after proposal and feasibility checks. That set can change as the target-conditioned reconstruction and current pose change, while a hierarchical controller can also condition camera position on an earlier look-at decision [Jia+25, Lu+26]. Sequential value must therefore be conditioned on the currently available action rather than treating every camera pose as permanently selectable.

Sequentiality matters when the best immediate view is not the best first step of a sequence. For target e , candidate i , and requested horizon h , the finite return and corresponding conditional value are

$$G_{t,e}^{(h)} = \sum_{k=0}^{h-1} \gamma^k r_{t+k}^e, \quad 1 \leq h \leq b_t \leq H_{\max}$$

$$Q_{h,e}^*(s_t, i) = \sup_{\pi \in \Pi^{\text{act}}} \mathbb{E}_{\pi} \left[G_{t,e}^{(h)} \mid s_t, a_t = i \right], \quad i \in \mathcal{A}_t, \quad 1 \leq h \leq b_t \leq H_{\max}, \quad Q_{0,e}^*(s, i) = 0$$

where continuation policies choose only from the admissible candidates at each step. These equations specialize the standard discounted-return and action-value constructions to a requested target and finite horizon [SB18, Secs. 3.3–3.6, pp. 54–67]. Fixed-Horizon TD defines horizon-indexed predictions with the boundary $Q_0 = 0$ and relates horizon h to the shorter horizon $h - 1$ [De +20]. If an episode ends before all h rewards are collected, later rewards are zero. The value thus states the consequence of taking a candidate now and then following a named continuation rule; it is not an intrinsic property of the camera pose alone.

The horizon is distinct from two related bounds. The requested horizon h states how many future acquisition rewards the query represents; the remaining budget b_t limits how many acquisitions are still possible; and $H_{\max}$ marks the largest horizon supported by the chosen model and evidence. Universal Value Function Approximators establish the general possibility of conditioning one approximator on an additional task variable, while Fixed-Horizon TD shares representations across horizon-indexed predictions [De +20, Sch+15]. Neither result licenses extrapolation beyond the horizons actually supported and evaluated.

Learning such values from a fixed data collection introduces an epistemic constraint in addition to the mathematical horizon. Batch-constrained and conservative offline reinforcement-learning methods are motivated by the error that arises when a learned policy assigns high value to actions outside the data distribution [FMP19, Kum+20]. These methods differ in how they control that error, but they share the warning relevant here: a longer horizon does not create evidence for transitions that the data never resolves.

Together, partial observability, state-dependent choices, and finite-horizon return make candidate value explicitly relational [De +20, LHP23]. A camera pose has no context-free value. Its value depends on the causal information state, requested target, currently admissible support, horizon, state-update rule, and continuation policy. The next dependency is therefore representational. The actor state must preserve the distinctions on which those future consequences depend.

2.5 Egocentric and Geometric Representations

The preceding section established that future view value depends on a causal information state rather than on a camera pose alone. Project Aria provides calibrated, time-aligned egocentric streams, and EFM3D lifts posed image features and semidense geometry into a local gravity-aligned representation [Eng+23, Str+24b]. The representation question is therefore not how faithfully to reproduce the entire world, but which distinctions must survive so that target-specific counterfactual return remains predictable.

Pretraining provenance and decision-time role must be distinguished. 3D-NVS uses ImageNet initialization in its NBV classifier, whereas MACARONS reuses pretrained image features within an online mapping pipeline [AKC20, Gué+23]. Next Best Sense instead uses SAM2 and monocular-depth priors to improve 3DGS construction while retaining an analytic Fisher-information selector [Str+24a]. EVL freezes its DINOv2.5 foundation feature extractor, then learns the feature upsampler, 3D U-Net, and task heads that form a local 3D field from posed egocentric streams and semidense geometry [Str+24b]. This structure may preserve semantic and geometric distinctions relevant to target visibility, but whether those distinctions improve candidate value is an empirical property of the downstream task rather than a consequence of pretraining alone.

Four requirements follow. First, **causal sufficiency** requires the actor state to retain the observation history relevant to future return without importing future or unselected evidence. EFM3D supplies strong local evidence but has a finite voxel extent, while Hestia demonstrates that accumulated directional history can alter later choices [Lu+26, Str+24b]. A spatial state is therefore not sufficient merely because it is geometric; sufficiency is a hypothesis to be tested against the task.

Second, **spatial relationality** requires the target, observer, candidates, and observed support to be comparable in a common local geometry. Query-centric models express scene elements relative to the active query, reducing dependence on an arbitrary global origin [Zho+23]. For target-conditioned NBV, the relevant

quantity is therefore not candidate position in isolation but candidate geometry relative to the current observer, target, and accumulated evidence. Coordinate choice also determines which changes should be treated as nuisances. Translating the world origin should not change a score, whereas gravity, metric scale, camera direction, target orientation, occlusion, and temporal order can remain informative [Bro+21]. The useful prior is thus relative geometry that removes arbitrary coordinates without erasing task-relevant physical structure.

Third, **candidate-order equivariance** follows from the action set itself. A candidate table denotes physical viewpoints rather than a ranked sequence, so permuting its rows must induce the same permutation of their scores; an invariant output would instead discard which score belongs to which action. Deep Sets provides invariant aggregation for shared set context, while Set Transformer provides permutation-equivariant candidate interaction [Lee+19, Zah+17]. This is a behavioral contract rather than an architecture prescription: independent row scoring, pooled context, and attention can all satisfy it.

Fourth, **physical observability** requires the representation to preserve what the causal sensing process actually distinguishes. EFM3D derives a surface mask from observed semi-dense points and a free-space mask from camera-to-surface rays [Str+24b]. Voxels supported by neither mask remain unobserved; they must not be treated as observed free space or negative surface evidence. This surface/free/unobserved distinction is the observation contract. The choice of carrier is a separate architecture tradeoff: point models retain irregular samples and local relations, sparse-voxel models regularize space at a chosen resolution and extent, and equivariant message passing commits to a chosen symmetry family [CGS19, SHW21, Zha+21]. These families trade spatial fidelity, computational structure, and strength of geometric prior, but none by itself guarantees the observation contract or improves the endpoint utility.

These four requirements—causal sufficiency, spatial relationality, candidate-order equivariance, and physical observability—complete the conceptual dependency chain. The literature synthesis can now compare methods by the scientific

distinctions they preserve, while the Method chapter remains responsible for one concrete realization and its tests [Bro+21, Str+24b].

2.6 Related-Work Synthesis and Research Gap

The reviewed approaches differ along six coupled dimensions: utility, target scope, represented evidence, representation provenance and decision-time role, candidate support and admission, and decision horizon [Fra+25, GML22, Jia+25, JLD24]. Comparing only one dimension can obscure the scientific task: a sequential coverage controller does not answer the same question as a greedy reconstruction-quality scorer, and neither comparison is meaningful if the relevant view lies outside candidate support.

Family	Utility and target scope	State evidence and representation role	Support, feasibility, and decision
PB-NBV [Jia+25]	Projected frontier and occupied coverage; reconstruction object	Analytic classified voxel clusters and compact projection proxies	Reachability- and camera-conditioned partial hemisphere; greedy score
3D-NVS [AKC20]	Pairwise voxel reconstruction; object	ImageNet-pretrained VGG16 directly scores fixed next-view classes	Eleven fixed view classes; one-step classifier
SCONE / MACARONS [GML22, Gué+23]	Expected surface coverage; scene-wide	Task-trained occupancy and visibility; pretrained ResNet features support MACARONS depth	Iterative candidate selection
GenNBV / Hestia [Che+24, Lu+26]	Coverage gain; reconstruction object	Task-trained geometric, history, and shallow-image encoders; Hestia uses no external pretraining	Continuous or hierarchical policy; collision handling
FisherRF [JLD24]	Information gain; radiance-field scene model	Per-scene model uncertainty enters an analytic score	Candidate views; greedy score
VIN-NBV [Fra+25]	Direct reconstruction improvement; reconstructed object	Task-trained CNN over candidate-conditioned projections of current reconstruction	Sampled candidates; greedy score
Next Best Sense [Str+24a]	Color-depth information gain; radiance-field scene model	SAM2 and depth priors support 3DGS; Fisher selection remains analytic	Feasible candidates; fallback after kinematic or planning failure
Object-centric 3DGS [Jeo+26]	Object-conditioned information gain; requested object	Foundation-model masks and CLIP support object state and selection; target score remains analytic	Candidate views; greedy score

Table 2: Concept-centred comparison of the literature families that bound the thesis question. Rows are distinguished by utility, target scope, state evidence and representation role, candidate support and feasibility, and decision structure; they are not ranked by reported performance across incompatible settings.

Pretraining alone does not determine the scientific role of a representation. 3D-NVS uses ImageNet initialization directly in a fixed-view classifier, while MACARONS uses pretrained image features inside an online mapping pipeline [AKC20, Gué+23]. Next Best Sense places foundation-assisted masks and depth upstream of an analytic Fisher selector [Str+24a]. By contrast, VIN-NBV, GenNBV, and Hestia learn their decision state from task-specific reconstruction cues or shallow encoders [Che+24, Fra+25, Lu+26]. These approaches establish different uses for transferred visual features, but they do not isolate the value of an egocentric 3D representation that combines frozen 2D foundation features with learned 3D processing for target-conditioned endpoint prediction.

The comparison yields two requirements for the value definition. First, **utility alignment** requires the score to reflect the endpoint of interest: coverage and information criteria reward proxies for reconstruction, whereas VIN-NBV evaluates realized reconstruction error [Fra+25, GML22, JLD24]. Second, **temporal dependence** requires the score to represent delayed consequences. Sequential policies retain history and optimize delayed reward, but the reviewed examples optimize coverage rather than the reconstruction quality of a requested target [Che+24, Lu+26]. Existing work therefore establishes direct quality, target conditioning, and sequentiality separately under different assumptions.

A third requirement is **support/state adequacy**: valuable actions must be available, and the causal state must retain the consequences needed to value later actions. PB-NBV constrains proposal support; Hestia carries directional history while separately repairing infeasible endpoints; and Next Best Sense distinguishes feasible candidates from execution fallback [Jia+25, Lu+26, Str+24a]. Even a perfect scorer cannot exploit a view outside its support, and a longer horizon cannot recover a distinction erased from the state.

Within the actor-state question, representation transfer is a downstream diagnostic. A pretrained carrier can help only if it preserves distinctions aligned with target-specific return; semantic strength in 2D does not by itself provide causal scene memory, metric support, or candidate-relative geometry. EFM3D/

EVL supplies a foundation-derived local field in which a frozen 2D feature extractor feeds learned upsampling, 3D processing, and task heads alongside observed surface and free-space evidence [Str+24b]. Whether that field improves target-conditioned candidate value must therefore be measured against matched actor-visible geometric controls after the core evidence gates, rather than inferred from pretraining.

This thesis studies the core requirements jointly within a bounded setting: whether a finite-candidate policy can improve endpoint reconstruction quality for a specified target by valuing more than the next observation. Actor-visible foundation-derived egocentric evidence is then a conditional representation diagnostic for the corresponding value model [De +20, Fra+25, Str+24b]. Among the reviewed reconstruction-NBV methods, none establishes this conjunction: direct-quality prediction, target conditioning, bounded lookahead, and a matched test of a foundation-derived egocentric representation with frozen 2D features and learned 3D processing. The literature therefore cannot determine whether non-myopic headroom exists or whether EFM3D/EVL-derived evidence improves its recovery; the latter remains a conditional diagnostic rather than another principal question.

The resulting principle is relational rather than architectural: view value is defined by a target, a causal state, admissible support, a horizon, a state update, and a continuation rule. The thesis does not generalize this bounded relation to exhaustive novelty priority, continuous control, actor-visible target discovery, closed-loop motion execution, or natural human movement [Che+24, Eng+23, Lu+26, Str+24a]. Chapter 3 makes the bounded relation operational by separating observable state from the privileged geometry used to construct tasks and evaluate outcomes.

3 Oracle and Data Generation

This chapter defines the experimental objects from which ARIA-NBV constructs training and evaluation data. A logged snippet first induces an actor-side information state s_t^{hist} . A target task then turns that state into a finite candidate table $\mathcal{Q}_t$ the hard action mask $m_{t,i}^{\text{act}}$ selects its feasible subset. Privileged rendering assigns target-specific reconstruction outcomes, while a selected action creates the only factual successor that may extend the causal history. Replay storage preserves these relations. Final padded tensors are model-specific projections.

State determines what may be conditioned on. The target and proposal mechanism determine which actions can be compared. The oracle defines their supervision, and lineage determines which multi-step returns can be reconstructed from factual transitions. These dependencies connect the chapter’s information, action, measurement, and storage contracts to the learned scorer in the Method chapter.

3.1 State as an Information Contract

An NBV state contains exactly the information on which an action value may depend. ARIA-NBV therefore distinguishes three nested roles: logged evidence from the observed snippet, causal evidence produced by actions that were actually selected, and privileged quantities used only to construct or evaluate counterfactuals. ASE provides both sensor-like streams and ground truth, whereas EFM3D transforms only logged image evidence into a local 3D representation [Met25, Str+24b]. Provenance and acquisition time determine information role independently of file location and tensor shape.




Figure 1: Information roles in the offline experiment. The actor side contains calibrated logged observations, derived EFM3D evidence, an explicitly sourced target instruction, the selected-view prefix, remaining budget, candidates, and the hard action mask. The oracle side adds GT geometry, counterfactual renders, labels, and endpoint evaluation. Every stored field keeps its declared information role.

The boundary is also temporal. A render from an unselected candidate is a counterfactual oracle query. After that candidate is selected, an observation from its pose may enter the next state, but only under a declared acquisition model: a real observation, an admitted sensor simulation, or a privileged mesh render. Actor visibility requires both factual selection and an admitted source role.

3.1.1 Logged Information State

The logged historic state s_t^{hist} groups the evidence available before any counterfactual acquisition:

$$s_t^{\text{hist}} = (\mathbf{I}^{\text{rgb}}, \mathbf{X}, \mathcal{P}^{\text{semi}}, \mathbf{F}_v, e_t, b_t) \in \mathcal{S}^{\text{hist}}$$

Calibrated images and trajectory poses establish where and when evidence was acquired. Semi-dense points provide sparse surface observations, and the EFM3D field $\mathbf{F}_v$ supplies a learned local representation. EFM3D lifts posed image features into a finite gravity-aligned local volume [Str+24b]. Absence from that volume denotes missing representation support. Empty-space evidence requires measured

rays, and pose feasibility requires the geometric and motion checks defined later in this chapter.

Semi-dense points $\mathcal{P}^{\text{semi}}$ retain uncertainty, timestamps, and observing-camera lineage. EFM3D uses that lineage to distinguish surface samples from free-space samples along measured rays [Met25, Str+24b]. A point set supplies surface samples; camera rays additionally supply free-space evidence. Missing surface samples carry no free-space meaning without that lineage.

The target e_t and remaining budget b_t are task variables. They make the value conditional, conceptually $Q(s_t, e_t, a_t)$, and distinguish two otherwise identical observations posed with different requested objects or acquisition budgets. In the present oracle task, e_t is derived from a selected GT box. A deployable task derives the target instruction from an observation-side proposal and association path.

3.1.2 Causal Successor State

After an action is selected, the state may grow only with evidence generated by that factual choice. The conceptual counterfactual actor state is

$$s_t^{\text{cf0}} = (\Phi_t^{\text{scene}}, \varphi_e, \mathcal{Q}_t, m_{t,i}, \mathbf{H}_t^{\text{pose}}, b_t) \in \mathcal{S}^{\text{cf0}}$$

and is read together with the ordered selected-view history:

$$\mathbf{h}_{t,e,i} = \text{Read}(\Phi_t^{\text{scene}}, \mathbf{h}_e^{\text{tgt}}, q_{t,i}, \mathbf{H}_t^{\text{pose}}, t, H)$$

Here the root field is immutable, while the accumulated evidence may grow only through selected observations. Candidate rows $q_{t,i}$ encode possible poses and proposal provenance. Observations from those poses become available only after selection. The hard mask $m_{t,i}^{\text{act}}$ defines the action set. Invalid-reason codes explain exclusions as diagnostic fields outside the value-model input. This separation prevents the proposal table from leaking the consequences it is supposed to predict.

The implemented **S0-pose** baseline materializes the weaker state

$$s_t^{\text{S0-pose}} = (\mathcal{S}_{\text{root}}^{\text{VIN}}, \mathbf{F}_v^{\text{root}}, (\mathbf{T}_{r,e}, \mathbf{l}_e), \mathcal{Q}_t, m_{t,i}, \mathbf{H}_t^{\text{pose}}, b_t)$$

It carries root evidence, target context, the current candidate table and mask, the selected-pose prefix, and the remaining budget. This state represents the

actor’s motion history. A geometry-enriched successor additionally retains what the selected observations revealed:

$$s_t^{\text{cf}+} = \left(s_t^{\text{cf}0}, (\mathbf{D}, \mathbf{V}, \mathcal{P}^{\text{cf}}, \mathbf{n})_{1:t} \right)$$

The implemented richer control is the intermediate carrier

$$s_t^{\text{CF-GT-carrier}} = \left(s_t^{\text{S0-pose}}, (\mathbf{D}, \mathbf{V}, \mathcal{K}, \mathbf{T}_{\text{root} < \text{cam}})_{1:t}^{\text{sel}} \right)$$

It retains calibrated depth only for selected actions. This makes the carrier causal with respect to the selected path, while its mesh-rendered source keeps the control privileged. Fusion into surface, free-space, uncertainty, and EFM3D evidence defines the fuller geometry state. The implemented carrier isolates the value of selected depth before that fusion.

More generally, the transition from t to $t + 1$ may consume only the observation produced by a_t . RGB and EFM3D updates require the corresponding calibrated image stream; depth updates the declared geometry channels. $s_t^{\text{cf}0}$ is an **information state** for value prediction. Its sufficiency for the physical process is an empirical question: it is adequate only if it preserves the distinctions needed to predict future target-specific return.

3.1.3 Oracle Augmentation

The oracle augments the causal state with quantities needed to answer counterfactual questions:

$$s_t^{\text{oracle}} = \left(s_t^{\text{cf}+}, \mathcal{M}^{\text{GT}}, \mathcal{M}_e^{\text{GT}}, \mathbf{D}_q, \mathcal{P}_q, \text{RRI} \right) \in \mathcal{S}^{\text{oracle}}$$

ASE provides per-frame metric GT depth aligned with RGB, per-pixel instance identifiers, class mappings, and the scene trajectory. The EFM3D release adds ASE OBB metadata and validation meshes for object-detection and surface-reconstruction supervision [Met25, Str+24b]. EFM3D uses depth to supervise occupancy at sampled free, surface, and behind-surface points, while visible OBBs supervise centerness, class, and box geometry [Str+24b]. These uses establish label provenance; the information source determines actor observability.

For ARIA-NBV, the scene mesh $\mathcal{M}^{\text{GT}}$ supports candidate rendering, and the target crop $\mathcal{M}_e^{\text{GT}}$ fixes the surface on which error is measured. All-candidate depth $\mathbf{D}_q$ and backprojected points $\mathcal{P}_q$ instantiate the hypothetical observation associated with each candidate. The resulting RRI labels the consequence of an action and becomes available only to the oracle.

A GT OBB defines oracle identity, pose, extent, crop, and target bearing. Actor detection and association require an observation-derived proposal. In a deployable protocol, that proposal defines the target hypothesis, while ASE identity and geometry serve as matching and evaluation references.

The three states form a controlled augmentation chain: s_t^{hist} fixes logged evidence, s_t^{cf0} adds only factual selected history, and s_t^{oracle} adds counterfactual supervision. This ordering is the conceptual data contract inherited by candidate generation and replay.

3.1.4 Visibility, Feasibility, and Utility

Three predicates apply to a candidate and answer different questions. **Visibility** asks whether a view supplies evidence about a spatial element. **Feasibility** asks whether the pose belongs to the admissible action set. **Utility** asks how the target reconstruction changes after an admissible action. A candidate can be physically feasible yet outside the local EFM3D volume, or it can frame the target yet yield little reconstruction gain. These cases coexist because the predicates refer to different structures.

Accordingly, $m_{t,i}^{\text{act}}$ owns infeasibility, and reward is defined only on the feasible action set. A feasible row may have negative gain. Oracle evaluation can fail for a separate reason—for example, an unusable target crop or render—in which case the label mask $m_{t,i}^{Q,h}$ is false while the geometric reason code is unchanged. Keeping the masks distinct preserves the difference between actions outside the feasible set and feasible actions lacking an oracle training label.

Actor visibility is a property of the complete pipeline. It includes the scorer tensor, the source of the target, the construction and masking of candidate support, the selected observation used for the successor, and the fields passed to the

model. A scorer with actor-only tensors can still depend indirectly on privileged support construction. Such a protocol supports an oracle-assisted experiment. Claims about independently actor-constructed actions require an actor-side proposal and feasibility path.

3.2 Task and Action Construction

An action has value only relative to a target and to the alternatives available in the same state. Data generation therefore begins by constructing a **target task** and a finite **candidate shell**. The target task fixes the requested entity e , its evaluation support $\mathcal{M}_e^{\text{GT}}$, and the instruction used by target-bearing proposal families. Its reusable descriptor is

$$\varphi_e = \text{Enc}_{\text{tgt}}(\hat{\mathbf{B}}_e, \hat{\mathbf{y}}_e, \hat{\pi}_e, A_e^{\text{proj}}, n_e^{\text{semi}}, n_e^{\text{EVL}}, \omega_e^{\text{EVL}}, \ell_e^{\text{src}}, \mathbf{T}_{r_t, e}, \mathbf{T}_{c_t, e})$$

where geometry, class evidence, observation support, source role, and target-relative transforms stay explicit. In the present oracle protocol these fields originate from a selected GT OBB. The descriptor therefore defines a controlled target-conditioned task. Actor-visible target discovery is evaluated through the separate proposal-association path below.

At decision step t , the proposal mechanism produces $\mathcal{Q}_t$. Each row $q_{t,i}$ represents one possible camera pose together with its proposal provenance. The hard mask $m_{t,i}^{\text{act}}$ then induces the feasible action set

$$\mathcal{Q}_t = \{q_{t,i}\}_{i=1}^{N_q}, \quad \mathcal{A}_t = \{i \in \{1, \dots, N_q\} : m_{t,i}^{\text{act}} = 1\}$$

Candidate generation and feasibility are deliberately separate operations. The proposal distribution determines which parts of the local action space are represented; the mask determines which proposed rows may be selected. Invalid rows are retained for action-space coverage and failure analysis. Selection, stochastic normalization, loss targets, and bootstrap maximization use the feasible subset. A feasible row with poor target gain is a valid low-utility example.

3.2.1 Target Selection

The oracle task sampler enumerates non-padding GT OBBs with finite positive geometry and samples them uniformly without replacement up to the manifest-defined per-snippet cap. This choice makes task selection independent of candidate reward: near-solved, difficult, and negative-gain targets are eligible. The stored **matched** status records a geometry-valid oracle task. Successful actor proposal-to-identity matching is a separate observed-target event.




Figure 2: Oracle target-task construction. Geometry-valid GT OBB rows form the task population; seeded uniform sampling without replacement applies the manifest-defined cap. Candidate scoring later determines whether the selected target crop is evaluable.

This sampler controls computational cost and keeps the task population independent of oracle headroom. Actor-visible target selection uses observation-derived proposals and association as a distinct stage. Oracle task admission supplies the controlled reference population for that evaluation.

3.2.2 Observed-Target Admission

The observed-target matcher implements that distinct stage. Given an observed or predicted proposal $\hat{e}$, it compares only same-class reference OBBs using oriented three-dimensional IoU,

$$\mu_{\text{IoU}}(\hat{e}, e) = \text{IoU}_{3D}(\hat{\mathbf{B}}_{\hat{e}}, \mathbf{B}_e)$$

$$\tau_{\text{IoU}} = 0.20$$

A reference row qualifies only when its IoU is strictly greater than 0.20. The proposal is admitted exactly when one reference row qualifies:

$$n_{\text{qual}}(\hat{e}) = \text{card}(\{e \in \mathcal{E} : \kappa(\hat{y}_{\hat{e}}, y_e) = 1 \wedge \mu_{\text{IoU}}(\hat{e}, e) > \tau_{\text{IoU}}\})$$

$$a_{\text{id}}(\hat{e}) = 1 \text{ iff } n_{\text{qual}}(\hat{e}) = 1$$

Equality at the threshold, no qualifying row, and multiple qualifying rows leave the proposal unresolved. The rule is intentionally simple: it turns an observation proposal into an unambiguous evaluation identity. Association confidence is a separate diagnostic, and the current oracle task sampler operates directly on GT rows.

3.2.3 Candidate View Generation

Candidate generation is a proposal distribution over a bounded local camera motion space. It determines the alternatives over which the NBV scorer ranks actions. At step t , it constructs the shell


$$\mathcal{Q}_t = \{q_{t,i}\}_{i=1}^{N_q}, \quad N_q = 60$$

and records a family label $k(i)$ for every row. The frozen `realistic_core_60` profile allocates 24 forward-local, 24 target-bearing-local, and 12 lateral-target-bypass rows. The three families encode complementary geometric hypotheses:

Family	Geometric role	Profile support
Forward-local	Preserve egocentric motion continuity while perturbing the current forward direction.	24 rows; radius [0.25, 1.1] m
Target-bearing-local	Move approximately along the current target bearing and orient the camera towards the target.	24 rows; radius [0.4, 1.1] m
Lateral-target-bypass	Introduce side-step parallax while retaining a component towards the target.	12 rows; radius [0.25, 1.1] m

Table 3: Geometric roles of the three families in the frozen `realistic_core_60` intervention. Counts and radii are specific to this profile.

Optional refinement, revisit, and target-orbit families activate only in challenger profiles. Because a profile changes the represented action support, its family counts, radii, angular caps, and pruning parameters belong to the resolved run manifest. A comparison fixes that manifest across methods.



Scene 81286, sample ASE_81286_Atek_000035, rollout row 73, step row 121 24 forward-local + 24 target-bearing
 1. 12 lateral-bypass candidates. Dense mesh layers are
 z-buffered raster; OBBs, paths, centres, and wire frusta remain vector geometry.

Figure 3: One stored candidate shell from ASE scene 81286 and sample ASE_81286_Atek_000035. The perspective panel relates the shell to the logged camera history, target OBB, and selected prefix. The bird's-eye panel retains all 60 candidate centres: 25 are feasible, 35 fail clearance, and oracle-greedy selects row 47. The example's evidential role is proposal support and hard feasibility.

Each row begins with a random unit direction in the reference rig frame. The frozen profile uses a forward-biased Power Spherical distribution with $\kappa = 8$ [DAT20]:

$$\mathbf{u}_i \sim \text{PS}(\mathbf{e}_z, \kappa), \quad p(\mathbf{u}) = c_\kappa (1 + \mathbf{e}_z^T \mathbf{u})^\kappa, \quad \kappa = 8$$

The sampled direction is then mapped into the configured azimuth and elevation caps. With $\psi = \text{atan2}(u_x, u_z)$ and $u_y = \sin \theta$, the transform is

$$\psi' = \psi \frac{\Delta\psi}{2\pi}, \quad y' = \sin \theta_{\min} + \frac{u_y + 1}{2} \cdot (\sin \theta_{\max} - \sin \theta_{\min})$$

$$\mathbf{d}_i^0 = \left(\sqrt{1 - y'^2} \sin \psi', y', \sqrt{1 - y'^2} \cos \psi' \right), \quad \|\mathbf{d}_i^0\|_2 = 1$$

The cap controls angular extent without changing the row count. Family-specific maps then reinterpret the same bounded random variable relative to rig forward, the target bearing, or the lateral direction. This factorization separates **diversity within a family** from the **geometric meaning of the family**.

Let $\mathbf{f} = \mathbf{e}_z$ denote rig forward, $\mathbf{b}_e$ the supplied target bearing, $\mathbf{l}_e$ its normalized horizontal lateral direction, and $\mathbf{e}_y$ world up in the sampling frame. The three core maps are

$$\mathbf{g}_i^{\text{forward}} = \mathbf{f} + \alpha_f(\mathbf{d}_i^0 - (\mathbf{d}_i^0 \cdot \mathbf{f})\mathbf{f}), \quad \mathbf{d}_i^{\text{forward}} = \frac{\mathbf{g}_i^{\text{forward}}}{\|\mathbf{g}_i^{\text{forward}}\|_2}, \quad \alpha_f = 0.45$$

$$\mathbf{g}_i^{\text{target}} = \mathbf{b}_e + \alpha_t(\mathbf{d}_i^0 - (\mathbf{d}_i^0 \cdot \mathbf{b}_e)\mathbf{b}_e), \quad \mathbf{d}_i^{\text{target}} = \frac{\mathbf{g}_i^{\text{target}}}{\|\mathbf{g}_i^{\text{target}}\|_2}, \quad \alpha_t = 0.4$$

$$\mathbf{g}_i^{\text{bypass}} = 0.55\mathbf{b}_e + 0.85 \text{sign}(d_{i,x}^0)\mathbf{l}_e + \text{clip}(d_{i,y}^0, -0.35, 0.35)\mathbf{e}_y, \quad \mathbf{d}_i^{\text{bypass}} = \frac{\mathbf{g}_i^{\text{bypass}}}{\|\mathbf{g}_i^{\text{bypass}}\|_2}$$

A family-conditioned radius turns direction into translation; the reference pose then maps that offset into world coordinates:

$$r_i \sim \mathcal{U}(r_{\min}^{k(i)}, r_{\max}^{k(i)}), \quad \mathbf{c}_i^w = \mathbf{T}_r^w(r_i \mathbf{d}_i^{k(i)})$$

Forward-local rows retain the reference orientation. Target-looking rows use the supplied target centre $\mathbf{p}_e$ to construct an orthonormal look-at frame:

$$\mathbf{z}_i^w = \frac{\mathbf{p}_e - \mathbf{c}_i^w}{\|\mathbf{p}_e - \mathbf{c}_i^w\|_2}, \quad \mathbf{g}_i^{\text{up}} = \mathbf{e}_y - (\mathbf{e}_y^T \mathbf{z}_i^w) \mathbf{z}_i^w, \quad \mathbf{y}_i^w = \frac{\mathbf{g}_i^{\text{up}}}{\|\mathbf{g}_i^{\text{up}}\|_2}, \quad \mathbf{x}_i^w = \mathbf{y}_i^w \times \mathbf{z}_i^w$$

The target centre in these equations is privileged under the present oracle protocol. A deployable candidate generator must replace it with an observation-derived target estimate. These equations specify the sampling prior; oracle opportunity and downstream policy evidence evaluate the utility of the actions that survive feasibility pruning.

The optional `target_orbit` challenger makes this distinction explicit. It proposes paired signed angular offsets at the current horizontal target standoff, thereby testing bilateral partial-orbit coverage. Pruning can remove one side, so proposal symmetry and feasible symmetry are measured separately. The family serves as a geometric ablation and contains no learned human-motion model.

Pruning evaluates each proposal independently of its oracle reward. A row is feasible only if it lies in the snippet support, stays clear of the GT mesh, admits a collision-free straight path, and satisfies the local motion bounds

$$\|\mathbf{o}_i\|_2 \leq 1.0 \text{ m}, \quad |\Delta h_i| \leq 0.25 \text{ m}, \quad \max(0, -o_{i,z}) \leq 0.25 \text{ m}, \quad \Delta\psi_i \leq 70 \text{ deg}$$

The full shell retains proposal probability, family, rule masks, diagnostics, and invalid-reason bitsets. The hard mask then produces $\mathcal{A}(s_t)$ while retaining rejected rows for analysis. A separate root-support gate rejects a task if pruning leaves too few alternatives:

$$N_{\text{valid}} \geq \max(12, \text{ceil}(0.25N_q))$$

This gate protects the statistical support of the decision problem: one or two surviving actions provide insufficient variation for a candidate-ranking test. Oracle opportunity measures the utility available among the surviving candidates. Rejected roots and family-specific failure reasons enter the dataset population report.

3.2.4 Candidate-Support and Geometric-Coverage Diagnostics

Learning can rank only the alternatives supplied by candidate generation. The proposal profile is therefore part of the experimental design, and its support must be characterized before reward or policy results are interpreted. Let $\mathcal{J}_s$ contain all attempted non-padding rows, let $\mathcal{V}_s \subset \mathcal{J}_s$ contain the feasible rows selected by $m_{t,i}^{\text{act}}$, and let $\mathcal{L}_s \subset \mathcal{V}_s$ contain feasible rows with finite oracle labels. These nested populations separate three questions: what was proposed, what could be selected, and what could supervise learning.

The actor-valid fraction

$$f_{\text{actor}(s)} = \frac{\sum_{i \in \mathcal{J}_s} m_{s,i}^{\text{act}}}{|\mathcal{J}_s|}$$

measures pruning severity, whereas

$$n_{\text{valid}(s)} = |\{i \in \mathcal{J}_s : m_i^{\text{act}} = 1\}|$$

counts the number of alternatives that survive pruning. The distinction matters because two profiles can retain the same fraction while presenting very different numbers of choices. Family collapse is measured separately by

$$z_{\text{family}(s)} = \left(\frac{1}{|\mathcal{F}_s|} \right) \sum_{f \in \mathcal{F}_s} \mathbb{1}[n_{\text{valid}(s,f)} = 0]$$

which is one exactly when every configured family has zero feasible rows and zero only when every family survives. It detects a failure that a pooled valid fraction can hide: a large forward family may dominate the row count even when all target-conditioned families disappear.

Because many states can originate from one scene, row-level pooling would let large scenes dominate the study. Every state statistic $q(s)$ is therefore averaged within scene and then macro-averaged across scenes:

$$\mathcal{S}_{c,q} = \{s \in \mathcal{S}_c : q(s) \text{ is defined}\}, \quad |(q)_c| = \left(\frac{1}{|\mathcal{S}_{c,q}|} \right) \sum_{s \in \mathcal{S}_{c,q}} q(s), \quad \mathcal{C}_q = \{c \in \mathcal{C} : |\mathcal{S}_{c,q}| > 0\}, \quad |(q)| = \left(\frac{1}{|\mathcal{C}_q|} \right) \sum_{c \in \mathcal{C}_q} |(q)_c|$$

The metric-specific set $\mathcal{S}_{c,q}$ makes undefinedness explicit. Failed generation roots contribute zero support and complete family collapse, while a directional statistic is undefined if no relevant direction was attempted. Undefined states and scenes are reported separately, and paired profile contrasts use the shared eligible-scene intersection. This preserves the difference between absence of support and absence of a defined measurement.

Support geometry is compared in the target-relative normalized frame $d_{t,e}^{\text{current}} = \|\mathbf{p}_e^w - \mathbf{c}_{r_t}^w\|_2$, $\tilde{\mathbf{c}}_{t,i}^{\text{support}} = (\mathbf{B}_{r,t}^{\text{Z-up}})^\top \frac{\mathbf{c}_{t,i}^w - \mathbf{c}_{r_t}^w}{d_{t,e}^{\text{current}}}$, $\tilde{\mathbf{p}}_{t,e}^{\text{support}} = (\mathbf{B}_{r,t}^{\text{Z-up}})^\top \frac{\mathbf{p}_e^w - \mathbf{c}_{r_t}^w}{d_{t,e}^{\text{current}}}$, $\|\tilde{\mathbf{p}}_{t,e}^{\text{support}}\|_2 = 1$. The factual expansion pose is the origin, the target lies at unit distance, and world up fixes the vertical axis. The normalization removes absolute standoff while preserving whether candidates move towards, around, above, or behind the target. This coordinate system describes candidate motion; calibrated projection and visibility are measured in their respective image and observation domains.

For attempted target-conditioned rows, side balance

$$^t > \varepsilon], \quad n_-(s) = \sum_{i \in \mathcal{J}_s^{\text{target}}} \mathbb{1}[y_{s,i}^{\text{target}} < -\varepsilon], \quad n_0(s) = \sum_{i \in \mathcal{J}_s^{\text{target}}} \mathbb{1}[|y_{s,i}^{\text{target}}| \leq \varepsilon], \quad b_{\text{side}(s)} = 1 - |n_+(s) - n_-(s)| \frac{1}{n_+(s) + n_-(s)}$$

compares positive and negative target-relative sides after excluding the neutral band $|y_{s,i}^{\text{target}}| \leq 10^{-9}$. Circular orbit span

$$\alpha_{s,1} \leq \dots \leq \alpha_{s,n_s}, \quad \alpha_{s,n_s+1} = \alpha_{s,1} + 2\pi, \quad s_{\text{orbit}(s)} = 2\pi - \max_{1 \leq j \leq n_s} (\alpha_{s,j+1} - \alpha_{s,j})$$

measures the smallest angular arc covering the attempted directions. These statistics characterize proposal diversity before pruning. Family survival and oracle opportunity quantify the later feasibility and utility stages.

For evaluated candidates, the projection fraction

$$f_{\text{proj}(s)} = \frac{\sum_{i \in \mathcal{J}_s^{\text{evaluated}}} \mathbb{1}[i \in \mathcal{J}_s^{\text{projected}}]}{|\mathcal{J}_s^{\text{evaluated}}|}, \quad |\mathcal{J}_s^{\text{evaluated}}| > 0$$

records whether the target centre falls inside the calibrated image domain. It measures geometric framing. Visibility additionally requires object extent and occlusion evidence. The finite-support oracle opportunity

$$o_{\text{oracle}(s)} = \max_{i \in \mathcal{L}_s} g_{s,i}^{\text{target-root}}, \quad |\mathcal{L}_s| > 0$$

is the best target-root gain present in $\mathcal{L}_s$. It measures whether the finite shell contains a useful action, independent of which policy selects it. Low opportunity identifies an action-support limitation; high opportunity with poor policy return instead leaves room for a ranking limitation.

Finally, jitter compliance checks whether within-family perturbations obey their declared angular caps:

$$f_{\text{jitter}(s)} = \frac{\sum_{i \in \mathcal{J}_s^{\text{jitter}}} \mathbb{1}[b_i = 1] \cdot \mathbb{1}[|\Delta\psi_i| \leq |(\psi)_i] \cdot \mathbb{1}[|\Delta\theta_i| \leq |(\theta)_i]}{\sum_{i \in \mathcal{J}_s^{\text{jitter}}} \mathbb{1}[b_i = 1]}, \quad \sum_{i \in \mathcal{J}_s^{\text{jitter}}} \mathbb{1}[b_i = 1] > 0$$

Uncapped spherical profiles form a different support class and are reported separately. Jitter compliance provides generator quality control; visibility and policy quality use their own measurements.

Quantity	Population	Unit	Interpretation boundary
Actor-valid fraction	Attempted rows per state	fraction	Hard feasibility before label and training filters.
Valid support	Actor-valid rows per state	rows	Report the distribution, lower tail, and failed roots.
Configured-family zero rate	Configured families per state	fraction	Retains absent and zero-support families; state values are scene-macro reduced.
Side balance / circular span	Non-neutral attempted target-conditioned rows	fraction / angle	Neutral rows and undefined states are reported separately; proposal coverage precedes pruning.
Target-centre projection	Evaluated candidate views per state	fraction	Calibrated framing; zero-evaluated states are reported as undefined.
Oracle opportunity	Finite actor-valid rewards per eligible state	gain	Available one-step headroom; undefined when no finite actor-valid label exists.
Jitter compliance	Bounded-jitter rows per state	fraction	State-scene-cohort macro; report undefined, nonzero, and uncapped support separately.

Table 4: Candidate-generation diagnostics and their aggregation populations. Every state-level quantity is reduced before scene and cohort aggregation.

3.2.5 Rollout Branch Sampling and Dataset Impact

Rollout recipes convert feasible candidate shells into factual training chains. Uniform sampling explores support without a utility preference; one-step oracle greedy selects the largest immediate target gain; bounded oracle lookahead ranks finite action sequences; and temperature-softmax samples according to oracle score while retaining stochastic diversity. These recipes define the behavior distribution from which replay is collected. They serve as data-generation policies and reference strategies. Chapter 5 evaluates the learned policy on held-out tasks.

Lookahead can select a different first action from one-step greedy because an action changes both the causal state and the next candidate shell (Figure 4). It ranks the return of retained chains, so future rewards can outweigh the first reward. This is decision-time planning: computation expands possible consequences to improve the current decision [SB18, Sec. 8.8, pp. 180–181]. The store preserves selected and beam-retained chains with their per-state shells. Search expands the other counterfactual branches transiently.




Figure 4: Conceptual topology of bounded oracle lookahead. Nodes are counterfactual states; edges are feasible actions labelled by immediate reward. The beam retains only a bounded set of prefixes. The symbolic inequalities show how a lower immediate reward can lead to a larger finite-horizon return.

For stochastic collection, the temperature-softmax recipe converts each finite oracle score s_i into the robust logit

$$\ell_i = \frac{s_i - \text{median}(s)}{\text{IQR}(s)\tau}, \quad P(i \mid m_i = 1) = \frac{\exp(\ell_i)}{\sum_{j:m_j=1} \exp(\ell_j)}$$

Median and interquartile-range normalization reduce sensitivity to the absolute score scale, while temperature controls concentration over the feasible rows. The resulting replay distribution is determined jointly by target sampling, candidate support, pruning, and rollout recipe. Reporting these factors isolates the population and behavior distribution on which learned ranking is assessed. Train-only pilots provide generation-cost and support-diversity evidence; held-out evaluation supplies policy evidence.

3.3 Measurement and Outcomes

The candidate shell defines possible interventions; the measurement operator assigns their target-specific consequences. For each feasible row, the oracle renders depth from $\mathcal{M}^{\text{GT}}$, backprojects $\mathcal{P}_q$, fuses the candidate evidence with the root evaluation cloud, crops both points and mesh to target e , and computes reconstruction error [Fra+25]. Applying the same operator before and after the hypothetical acquisition yields Δ_t^e and $\Delta_{t|i}^e$. The candidate label depends jointly on camera pose, current state, target crop, and update rule.

Validation TODO: Validate repeatability, numerical-failure handling, construct interpretation, root normalization, and matched endpoint re-evaluation before treating target gain as a reliable study outcome.
 Source: docs/literature/tex-src/arXiv-VIN-NBV; metric-validation artifacts
 Gate: frozen metric protocol, repeatability table, and claim-level source locators

3.3.1 Operational Reconstruction Error

Let C_e crop geometry to the selected target region. The evaluation point set $P_t^e = C_e(\mathcal{P}_t)$ is reconstructed from the observed prefix of ASE GT depth. Candidate i contributes renderer-derived depth, which is backprojected and fused with the same root cloud to obtain $P_{t|i}^e$. MPS semi-dense points serve as actor-side representation input and as a legacy diagnostic source. ASE depth defines the current oracle evaluation cloud.

For target mesh $M_e = \mathcal{M}_e^{\text{GT}}$, the point-to-mesh term is the mean squared distance from each reconstructed point to its closest reference triangle,

$$D_{P \rightarrow M}(\mathcal{P}, \mathcal{M}^{\text{GT}}) = \frac{1}{\|\mathcal{P}\|} \sum_{\mathbf{p} \in \mathcal{P}} \min_{\mathbf{f} \in \mathcal{F}^{\text{GT}}} d(\mathbf{p}, \mathbf{f})^2$$


It measures how closely the reconstructed evidence lies on the target surface and is therefore sensitive to extra, noisy, or misregistered points. The reverse mesh-to-point term is

$$D_{M \rightarrow P}(\mathcal{P}, \mathcal{M}^{\text{GT}}) = \frac{1}{\|\mathcal{F}^{\text{GT}}\|} \sum_{\mathbf{f} \in \mathcal{F}^{\text{GT}}} \min_{\mathbf{p} \in \mathcal{P}} d(\mathbf{p}, \mathbf{f})^2$$

and measures which target faces lack nearby reconstructed evidence. Their sum is the target error

$$\Delta_t^e = d(C_e(\mathcal{P}_t), \mathcal{M}_e^{\text{GT}}) = D_{P \rightarrow M, t}^e + D_{M \rightarrow P, t}^e$$

with units of square metres. Root and candidate-augmented clouds pass through the same deterministic fusion and point-cap operator, so the difference isolates the candidate evidence under a fixed update rule. The resulting estimand is defined on the capped point set, fixed target mesh, and declared fusion operator.



Controlled validity fixture computed by the repository’s PyTorch3D metric and exact Trimesh closest-point witnesses; it is not an ASE performance result.

Figure 5: Geometry of the operational point–mesh error. Panel A isolates the point-to-triangle primitive; panel B separates point-to-face accuracy from face-to-point completeness. Panel C keeps the surface and point set fixed but changes its tessellation, changing equal-face completeness from 0.03640 to 0.02284 m^2 . The synthetic fixture provides evidence for metric sensitivity.

Because completeness weights faces equally, its value depends on mesh tessellation and differs from an area-weighted surface integral. Subdividing one region changes its contribution even when the geometric surface is unchanged (Figure 5). Comparisons are internally consistent when they share one fixed target mesh and scoring protocol, but the numeric error is specific to that mesh representation.

3.3.2 Reliability under the Frozen Protocol

For fixed source state, target, candidate, and manifest, repeated oracle evaluation should reproduce both $\Delta_{t|i}^e$ and the induced candidate ordering. This requirement covers the complete measurement operator—rendering, backprojection, target

cropping, fusion, capping, and numerical failure handling—because instability in any component changes the label. Deterministic code and a fixed mesh make exact replay possible, but RQ1 still requires repeatability measurements over the intended scene and target population.

Metric family	Question answered	Dependency and thesis role
Equal-face bidirectional point-mesh error	Are reconstructed points close to GT faces, and does every GT face have nearby point evidence?	Depends on point fusion and mesh tessellation; selected operational RRI error.
Point-sampled Chamfer / accuracy-completeness	Are two sampled surface point sets mutually close?	Depends on sampling density, randomness, norm, and squaring convention; alternative operationalization.
Thresholded precision, recall, and F-score	What fraction of predicted and reference samples fall within tolerance τ , and how do the two fractions balance?	Interpretable at a physical tolerance but threshold-dependent; secondary diagnostic alongside scalar RRI.
TSDF, occupancy, coverage, or information gain	How much volumetric surface, free space, or previously unknown space is reconstructed or observed?	Depends on grid, truncation, and visibility definitions; answers a different question from target surface quality.

Table 5: Reconstruction measures answer different geometric questions. The equal-face point-mesh error is the selected outcome; the other families are diagnostics or alternative operationalizations.

Point-sampled Chamfer distance draws a finite reference set Q_e from the mesh and averages nearest-neighbour distances in both directions:

$$D_{\text{PS-Chamfer}(\mathcal{P}, \mathcal{Q})} = \frac{1}{\|\mathcal{P}\|} \sum_{\mathbf{p} \in \mathcal{P}} \min_{\mathbf{q} \in \mathcal{Q}} \|\mathbf{p} - \mathbf{q}\|_2^2 + \frac{1}{\|\mathcal{Q}\|} \sum_{\mathbf{q} \in \mathcal{Q}} \min_{\mathbf{p} \in \mathcal{P}} \|\mathbf{q} - \mathbf{p}\|_2^2$$

Sampling density, randomness, norm, and squaring convention determine its finite value despite the symmetric formula. Thresholded precision, recall, and F-score add a physically interpretable tolerance τ , but discard error magnitude on either side of that threshold. Volumetric coverage and information gain measure observed space or uncertainty reduction. These are useful complementary quantities. Each operationalization maps a different geometric construct.

3.3.3 Construct Validity

The operational construct is **target-surface evidence after a declared acquisition and update**. ASE depth, the selected target crop, equal-face mesh reduction, and the fusion operator jointly define it. Perceived object quality and whole-scene completeness require different constructs. Two systems that share the final point-mesh formula but apply different state updates estimate different intervention effects.

The implementation includes a mesh face in the target crop when any vertex lies inside the target Oriented Bounding Box (OBB). A geometry-invalid candidate is excluded by $m_{t,i}^{\text{act}}$. An empty crop, insufficient root support, or unusable render makes the oracle outcome undefined and clears $m_{t,i}^{Q,h}$. This preserves whether failure occurred in the action domain or in the measurement operator.

Comparative policy evidence therefore requires the same target mesh and crop, ASE-depth source, renderer, backprojection, fusion, point cap, and metric configuration. Replacing ASE depth with MPS semi-dense evidence changes both the observation process and the error population and therefore defines a separate experimental outcome.

3.3.4 RRI, Training Reward, and Endpoint Gain

The learning and evaluation quantities are derived from the same error sequence $\Delta_0^e, \Delta_1^e, \dots, \Delta_H^e$, but answer different questions. State-relative RRI asks how much one action improves the **current** state; target-root reward assigns additive credit along a chain; endpoint gain asks how much the target improved after a fixed budget. Their denominators determine whether they can be accumulated or compared across states.

The VIN-compatible marginal diagnostic is

$$\text{RRI}_{t,i}^e = \frac{\Delta_t^e - \Delta_{t|i}^e}{\max(\Delta_t^e, \varepsilon)}$$

Its selected-chain running sum is

$$C_t^{\text{RRI},e} = \sum_{k=0}^{t-1} \text{RRI}_{k,a_k}^e$$

Because each term is normalized by its own Δ_t^e , the sum describes accumulated local relative improvements. Fixed-budget policy outcomes use the root-to-end-point definition below.

The training reward normalizes every step by the rollout-root target error [Fra+25]:

$$r_t^e = \frac{\Delta_t^e - \Delta_{t+1}^e}{\max(\Delta_0^e, \varepsilon)}$$

With unit discount, summing these rewards over a factual chain telescopes:

$$J_t^e = \sum_{k=0}^{t-1} r_k^e = \frac{\Delta_0^e - \Delta_t^e}{\max(\Delta_0^e, \varepsilon)}$$

The horizon-conditioned return used to supervise Q_H is

$$G_{t,e}^{(h)} = \sum_{k=0}^{h-1} \gamma^k r_{t+k}^e, \quad 1 \leq h \leq b_t \leq H_{\max}$$

This return evaluates a candidate first action under the continuation rule that generated the remaining factual chain. Its scope is offline finite-candidate control.

The primary policy outcome is endpoint gain

$$J_e^{(H)} = \frac{\Delta_0^e - \Delta_H^e}{\Delta_0^e + \varepsilon}$$

because it compares target error after the same acquisition budget H . The log-gain variant

$$J_{e,\log}^{(H)} = \log(\Delta_0^e + \varepsilon) - \log(\Delta_H^e + \varepsilon)$$

is retained as a scale-sensitivity diagnostic. Thus $\text{RRI}_{t,i}^e$ describes local candidate improvement, $G_t^{(H)}$ supplies the learning target, and $J_e^{(H)}$ supplies the fixed-budget evaluation outcome. The endpoint denominator $\Delta_0^e + \varepsilon$ differs from the clamped root denominator $\max(\Delta_0^e, \varepsilon)$ used by the additive reward. Near zero root error, each quantity must be interpreted under its own denominator. The resolved

manifest must therefore freeze the discount, clipping, crop, point update, and all evaluation-geometry parameters for each reported experiment.

3.4 Dataset and Storage Semantics

The replay dataset is organized around decision-state relations. Each state binds an information state s_t and target e to a candidate table $\mathcal{Q}_t$, row-specific feasibility and label masks, oracle outcomes where defined, and the selected successor. Its normalized lineage connects a source to target tasks, retained chains, decision states, and candidate rows.

One logged source can support several target tasks. Each task can be explored by several rollout recipes. A retained chain contains ordered selected steps, and each step owns the full finite candidate shell. This layout preserves two forms of supervision. Candidate rows record the counterfactual outcomes available at one state; the selected row a_t links that state to the factual successor s_{t+1} . The first supports dense one-step ranking, while the second supports finite-horizon returns along a causal trajectory.

chains, per-state candidate shells, selected-action successors, and recipe provenance. Changing a rollout recipe creates a new replay dataset over the same source facts. Observations and one-step oracle products keep stable identities across those experiments, which enables paired comparison.

For each selected action, replay stores one calibrated depth raster. The ordered rasters reconstruct the selected-observation prefix without duplicating every candidate render inside each chain. Source role stays attached to each raster: mesh-rendered depth defines the privileged causal control, and sensor-like or observed depth defines an actor-side successor. Candidate renders for unselected rows stay with the one-step oracle products used to construct labels.

The stored data-generation transition is

$$(x_{t+1}, \mathbf{H}_{t+1}^{\text{pose}}, b_{t+1}, \mathcal{Q}_{t+1}) = \text{Step}(x_t, \mathbf{H}_t^{\text{pose}}, b_t, q_{t,a_t}, \xi_t)$$

It updates reference pose, selected history, budget, and the next candidate table under generation context ξ_t . A visual successor adds a separately specified observation operator and state update for RGB, depth, EFM3D features, or surface memory. The rollout recipe retains the chains it selects or keeps in its beam. Matched oracle endpoint re-evaluation then supplies a common fixed-budget outcome for policy comparison.

3.4.2 Missingness and Derived Views

Every candidate row carries separate predicates for materialization, actor-feasible action $m_{t,i}^{\text{act}}$, finite value label $m_{t,i}^{Q,h}$, and factual successor m_t^{succ} . These predicates describe structure, action support, supervision support, and trajectory heritage. A training objective may select their intersection; the stored fields preserve the reason each row enters or leaves that population.

Model-ready arrays are deterministic projections of the normalized store. They pad candidate shells, expand a chain into state queries, and cache features while preserving source identity, target identity, state order, masks, requested horizon, and label provenance. Several states from one chain may share a batch, but each query receives only its own causal prefix.

Reproducibility requires content identity and transformation identity. The source manifest fixes scene, snippet, target, geometry, and oracle products; the rollout manifest fixes proposal profile, pruning, recipe, seed, and retention; the training view fixes tensorization and supported horizons. Coverage, failed roots, family survival, undefined metrics, source-role counts, and exclusions define the population to which a learned value claim applies. Compression, chunking, hashes, and execution commands complete the reproducibility record. Development bandwidth pilots contribute cost estimates; held-out policy evidence comes from the evaluation protocol in Chapter 5.

4 Method

Chapter 3 fixed the experimental world: a target-specific task, a action set, a causal transition, and a root-normalized reconstruction-gain outcome. This chapter fixes the learner that operates inside that world. Three classifications are kept independent throughout:

- **scientific role** asks whether an element is the selected realization, a matched control, a requirement of the scientific target, a contingent alternative, or an exploratory idea;
- **implementation maturity** asks whether its data path and computation are implemented, partial, planned, or only conceptual; and
- **evidence maturity** asks whether the corresponding scientific claim is validated, pending, conflicted, or inapplicable.

These axes prevent two recurrent category errors. Implemented does not mean selected or empirically supported, and scientific target does not mean one particular architecture. The target specifies the actor-visible information and evidence needed to answer the research questions; candidate realizations may then compete to satisfy it.

Scientific role	Example	Meaning in this chapter	Present maturity
Selected realization	$s_t^{S0\text{-pose}}$, H0, A1, direct regression	Executable reference method against which one-factor changes are interpreted.	implemented; scientific evidence pending
Matched controls	A0; $s_t^{S1\text{-surface}}$ H1; CORAL	Executable contrasts that preserve a named comparison but are not part of the selected method.	implemented; comparisons pending
Scientific target	actor-visible target source and causal observation-updated state	Information and evidence contract required for the core claim; architecture-neutral.	partly implemented; not validated
Contingent alternatives	$s_t^{S2\text{-ray}}$ candidate-set interaction; separate horizon heads	Promoted only after a measured failure identifies the missing distinction.	planned or executable; unselected
Exploratory ideas	3DGS memory; SceneScript context; quantile value	Conceptually relevant extensions whose estimand, comparison, or implementation is not yet frozen.	ideas only

Table 7: Interpretive status for Method. Scientific role, implementation maturity, and evidence maturity are orthogonal; no row is an empirical ranking.

All admitted state realizations share a scalar-horizon, target-conditioned candidate-value interface. The scorer reads admitted scene and target evidence, relative candidate geometry, causal history, and remaining budget; it predicts a conditional value for every materialized candidate before the authoritative hard mask is applied. What differs is whether the admitted actor state retains the target-specific information on which future return can depend.

The selected executable configuration uses the S0-pose ($s_t^{\text{S0-pose}}$), H0 mean-pooled $\mathbf{H}_t^{\text{pose}}$, A1 candidate-to-state cross-attention, and direct continuous Huber regression. A0 is the matched interaction control. The privileged selected-surface state $s_t^{\text{S1-surface}}$, ordered H1 history, and CORAL decoder are also executable controls, but none has a frozen comparative result. The planned ray-aware state $s_t^{\text{S2-ray}}$ is a candidate realization of the actor-state target, not a claim that rays are already known to be the best carrier. Candidate-set interaction and distributional value prediction remain contingent hypotheses rather than silently discarded possibilities.

The value query distinguishes factual remaining budget b_t from requested residual horizon h . Current supervision supports dense one-step queries and recursive queries on the factual budget diagonal; wider executable inputs do not establish wider learned support. Exact horizon two is therefore the first epistemic test of learned lookahead: it can compare learned recursion with an exact finite-support endpoint without trusting a learned longer-horizon continuation. Passing that test is necessary but not sufficient for a policy claim, which additionally requires positive oracle headroom and held-out endpoint recovery.

The chapter proceeds from the state design space and selected encoding, through finite action and replay semantics, to interaction alternatives, acceptance properties, and the finite-horizon objective. This order keeps failures interpretable: lost state information, malformed action or replay semantics, insufficient relational structure, and value-learning error remain distinct explanations rather than an undifferentiated model defect.

4.1 Current and Target Actor State

Implementation: partial **Evidence:** pending

The selected $s_t^{S0\text{-pose}}$ baseline and the independent `v1_observed` descriptor path are implemented and tested. The current selected experiment still uses privileged `v0_gt_input` and `oracle_action_mask_v1`; no frozen actor-visible target, action-support, or causal observation-updated state evidence yet supports the scientific target.

Literature: [LHP23, Str+24b]

Promotion gate: freeze and evaluate an actor-visible target corpus and action-support protocol; implement the causal state-update contract; pass source-dropout, no-future-observation, leakage, and held-out task-sufficiency tests

Development source: `aria_nbv/aria_nbv/targets/protocol.py`; `aria_nbv/aria_nbv/targets/selection.py`; `aria_nbv/aria_nbv/oracle/pipelines/campaign.py`; `aria_nbv/aria_nbv/oracle/pipelines/rollout_dataset.py`; `aria_nbv/aria_nbv/data_handling/qh_data/views.py`; `aria_nbv/aria_nbv/vin/modules/qh_scene_encoders.py`; `aria_nbv/aria_nbv/rollouts/qh_reader.py`; `aria_nbv/aria_nbv/vin/models/target_finite_horizon.py`

The state is the information available before selecting action a , not every quantity co-located in a replay row. Oracle target gains, mesh crops, ground-truth associations, current all-candidate renders, and labels are excluded from the scorer graph. Three independent protocol axes matter. The target-source axis distinguishes non-deployable `v0_gt_input` geometry from a `v1_observed` descriptor constructed from actor-visible observations with bound source and construction provenance. The dynamic-state axis distinguishes the pose-only $s_t^{S0\text{-pose}}$ carrier from a state updated with selected observations. The action-support axis distinguishes the current mesh-derived `oracle_action_mask_v1` from actor-observed support or a separately calibrated learned-feasibility route. Sharing tensor shapes or field names across these axes does not make their scientific claims interchangeable.

Both the current and target states retain the same decision interface,

$$\mathbf{h}_{t,e,i} = \text{Read}(\Phi_t^{\text{scene}}, \mathbf{h}_e^{\text{tgt}}, q_{t,i}, \mathbf{H}_t^{\text{pose}}, t, H)$$

but differ in what the state makes available to that read. The current root component contains semidense evidence and lossy global moments of locally supported EFM3D features. Its dynamic component is only the strictly causal selected-pose prefix. Candidate regeneration updates the reference pose, prefix, remaining

budget, and action table; it does not imply a new RGB observation, a refreshed EFM3D field, or fused selected depth.

Component	Current realization	Scientific target
Root scene	immutable semidense points and global moments of locally supported EFM3D features	immutable actor-visible evidence with explicit finite support and missingness; absent evidence must not mean observed free space
Target	selected experiment: privileged $\mathbf{v0_gt_input}$; an independent $\mathbf{v1_observed}$ admission and I/O path is implemented but not frozen or evaluated	$\mathbf{v1_observed}$: actor-visible identity and geometry with support, source, construction provenance, and explicit matching failures
Dynamic state	$\mathbf{H}_t^{\text{pose}}$ only	a strictly causal update from the selected observation that preserves observed surface, observed free, unknown, support, uncertainty, source, and recency
Action support	$\mathbf{oracle_action_mask_v1}$: privileged mesh-derived physical-validity support	$\mathbf{actor_observed_action_mask_v1}$, or a learned-feasibility route with frozen calibration, threshold, abstention, and false-valid evidence
Decision context	$\mathcal{Q}_t$, b_t , and h under the named oracle-support protocol	the same finite-candidate interface under actor-visible target, state, and action-support protocols

Table 8: Current baseline and scientific target. The right column is a promotion contract: it states information that must be represented and tested, not an architecture already implemented or a result already obtained.

The CF0 state is intentionally a semantic contract rather than a premature choice among points, voxels, rays, or another carrier. A proposed carrier is admissible only if its deterministic fusion rule exposes source and availability, uses the selected actor-visible observation and no future or unselected render, and preserves the distinctions in Table 8. Source-dropout, no-future-observation, and leakage tests establish the information boundary; held-out task-sufficiency evidence must then show that the retained distinctions are useful for the target-specific value task.

4.1.1 Sufficiency and the local-evidence limit

A decision state is task-sufficient for this study only if it preserves the distinctions needed to predict target-specific future return. Under partial observability, this is an empirical representation requirement rather than a claim that the study solves a general belief-state POMDP [LHP23]. $s_t^{\text{S0-pose}}$ cannot satisfy that condition by construction when two histories share poses but reveal different surfaces, occlu-

sions, or free-space evidence. It is therefore a deliberately compact baseline and an interface test, not a claim that pose history is a complete reconstruction state.

EFM3D reinforces this limit. Its EVL field has an explicit pose and finite extent [Str+24b]. A target or candidate outside that extent may remain physically valid while lacking local lifted evidence. Global moments further remove spatial correspondence. The selected method can test whether this compact context contains usable signal, but a failure cannot be attributed to value learning alone, and a success cannot establish spatial state sufficiency.

4.1.2 Candidate realizations of the state contract

The scientific target determines which distinctions must survive; it does not determine the carrier. This separates a representational question—what information is retained—from an encoder question—how that information is compressed and queried. The relevant alternatives therefore remain visible, but their status is explicit.

Carrier or context	Scientific role	Distinction retained	Limitation and promotion condition
Root moments + H_t^{pose}	selected $s_t^{\text{S0-pose}}$	Cheap root context and causal pose history.	No selected-observation update or spatial correspondence; interface baseline only.
Selected-surface point residual	implemented privileged control $s_t^{\text{S1-surface}}$	Causal selected surfaces in the current-camera frame.	Conflates observed free with unknown and is density weighted; requires a frozen matched receipt.
Sparse ray-aware memory M_t^{ray}	planned candidate $s_t^{\text{S2-ray}}$	Observed surface, free, unknown, support, uncertainty, source, and recency [Str+24b].	Requires deterministic actor-visible fusion, leakage tests, and candidate-relative query evidence.
Target-centred EVL re-lifting	contingent representation ablation	Logged EFM3D evidence when target support lies outside the root field [Str+24b].	May shift the learned 3D-neck distribution; compare first with simpler logged-feature pooling.
Appearance attached to observed points	contingent representation ablation	Logged semantics or texture beyond the local EVL extent [Str+24b].	Requires visibility, source lineage, compression, and explicit missingness.
Sparse TSDF/SDF or point encoder	contingent encoder ablation	Metric spatial structure with fixed-width learned tokens [Str+24b].	Must preserve observation weight and unknown-space semantics rather than only occupied surfaces.
Object-aware 3DGS	exploratory idea	Renderable memory, soft target membership, and primitive confidence [Jeo+26].	Per-scene optimization and mask supervision redefine cost and state; justify only after a renderability failure.
SceneScript context	exploratory idea	Broad layout and object hypotheses [Ave+24].	Global semantics do not replace causal local free/unknown evidence.

Table 9: Representation design space ordered by scientific role, not presumed performance. A carrier is promoted only when it tests a diagnosed information loss under the same task, action, replay, and evaluation contracts.

The first comparison is $s_t^{\text{S0-pose}}$ against the matched privileged $s_t^{\text{S1-surface}}$ control because it asks whether any selected-surface signal survives the current compression seam. A positive result would motivate an actor-visible version; a null result would remain ambiguous between absent signal and a lossy point summary. $s_t^{\text{S2-ray}}$ is therefore the primary planned realization of the complete state contract, but it is promoted only after actor-visible observations and deterministic fusion exist. The remaining carriers are conditional probes of specific failures, not parallel thesis methods and not discarded ideas.

4.2 Target, Candidate, and History Encoding

Implementation: implemented **Evidence:** pending

The selected encoder materializes one target-conditioned query per candidate from root evidence, local relative geometry, H0 selected-pose history, remaining budget, and requested horizon. Its tensor contract is tested; scientific usefulness remains pending.

Literature: [Bro+21, Sch+15]

Promotion gate: preserve local-frame, mask, source, and strictly causal prefix tests; evaluate held-out ranking

Development source: `aria_nbv/aria_nbv/data_handling/qh_data/views.py`; `aria_nbv/aria_nbv/vin/mod-`
`els/target_finite_horizon.py`; `aria_nbv/aria_nbv/vin/modules/qh_history_encoders.py`; `aria_nbv/tests/vin/`
`test_target_finite_horizon.py`; `aria_nbv/tests/vin/test_qh_history_encoders.py`

For target e and candidate $q_{t,i}$, the scorer input is the structured relation

$$\mathcal{J}_{t,e} = \left(\Phi_t^{\text{scene}}, \mathbf{T}_{r \leftarrow e}, \mathbf{a}_e, \left\{ \mathbf{T}_{r \leftarrow c_i}, \mathbf{T}_{c_t \leftarrow c_i}, \mathbf{T}_{c_i \leftarrow e}, m_{t,i}^{\text{cand}} \right\}_{i=1}^{N_q}, \mathbf{H}_t^{\text{pose}}, b_t, h \right)$$

rather than a flat replay row. Source and availability masks distinguish absent evidence from measured zeros; padding is separate from physical invalidity; supervision and audit lineage never enter the learned query.

4.2.1 Logical data roles

The normalized reader may collate actor inputs, supervision, transition linkage, and audit lineage in one batch, but tensor colocation does not grant the model access. The Method therefore describes the logical contract rather than enumerating storage columns:

Logical role	Scientific content	Model access
Scene and target state	root actor evidence, source-bound target descriptor, $\mathbf{H}_t^{\text{pose}}, b_t$	candidate-value input
Candidate table	stable row identity, local pose relations, materialization, physical validity, reason and proposal provenance	geometry and materialization only; $m_{t,i}^{\text{act}}$ is external
Value query	scalar h with $1 \leq h \leq b_t \leq H_{\max}$	input only when bundle-supported
Candidate supervision	root-normalized gain, fitted target, $m_{t,i}^{Q,h}$	loss admission after prediction
Selected transition	factual selected row, reward, terminal flag, successor identity	Bellman linkage only
Audit lineage	scene, target, store, generator, policy, seed, source, and configuration identity	never a learned feature

Table 10: Logical data roles at the reader–model boundary. Storage readability and batch colocation do not imply actor visibility.

This separation matters for causal interpretation. Target-source provenance is experiment identity rather than a shortcut feature; generator family is proposal-support evidence rather than a semantic action label; and oracle returns supervise $Q_{h,\theta,e,i}^{\text{cond}}$ only after that value has been emitted. The same reader can therefore support privileged controls and actor-visible experiments without pretending that they estimate the same value function.

4.2.2 Target-conditioned query

The selected target protocol supplies a root-relative target pose and metric OBB extents. The target token is therefore

$$\mathbf{h}_e^{\text{tgt}} = \text{TargetProj}(\text{concat}(\text{PoseEnc}(\mathbf{T}_{r \leftarrow e}), \mathbf{a}_e))$$

The source protocol that produced this geometry is frozen experiment identity, not another model input. Current oracle tasks use ground-truth-derived target geometry; this privileged source limits deployability even though gains, meshes, associations, and audit lineage remain outside the prediction graph.

4.2.3 Relative candidate geometry

Canonical world poses remain reproducibility facts, but the selected scorer encodes each candidate both relative to the rollout root and relative to the factual current camera. Its physical trunk is

$$\mathbf{h}_{t,i}^{\text{phys}} = \text{PhysicalProj}(\text{concat}(\text{PoseEnc}(\mathbf{T}_{r \leftarrow c_i}), \text{PoseEnc}(\mathbf{T}_{c_t \leftarrow c_i}), \Phi_t^{\text{scene}}))$$

The conditional-value query then adds the target pose expressed from that candidate:

$$\mathbf{x}_{t,i} = \text{ValueQueryProj}\left(\text{concat}\left(\mathbf{h}_{t,i}^{\text{phys}}, \text{PoseEnc}\left(\mathbf{T}_{c_i \leftarrow e}\right)\right)\right)$$

This factorization removes arbitrary world origin while retaining the complete relative rotations and translations represented by the shared PoseTW encoder. It does not append handcrafted range, bearing, height, frustum, sampler-family, or generator-provenance features. The physical token feeds feasibility without target or horizon information; only the conditional-value query receives the candidate-from-target transform.

4.2.4 Strictly causal pose history

For every realized state, the reader supplies the complete $\mathbf{H}_t^{\text{pose}}$ for $j < t$, expressed from the factual current camera. The selected H0 encoder applies the shared pose map and a padding-aware arithmetic mean:

$$\mathbf{p}_{t,j}^{\text{hist}} = \text{PoseEnc}\left(T_{c_t \leftarrow c_j}\right), \quad j < t$$

$$a_{t,j}^{\text{hist}} = \frac{t - 1 - j}{H_{\text{max}}}$$

$$\mathbf{h}_t^{\text{histH0}} = \text{HistProj}\left(\frac{1}{\max(1, t)} \sum_{j < t} \mathbf{p}_{t,j}^{\text{hist}}\right)$$

$$\mathbf{h}_t^{\text{histH1}} = \text{HistProj}\left(\text{LastValid}\left(\text{CausalTransformer}\left(\left[\mathbf{e}_{\text{empty}}, \left\{\mathbf{p}_{t,j}^{\text{hist}} + g(a_{t,j}^{\text{hist}})\right\}_{j < t}\right]\right)\right)\right)$$

Only the H0 branch of this equation belongs to the selected method. The H1 branch is implemented and replaces the order-invariant mean with a causal Transformer and relative age; it remains an unselected pose-only control. The root state uses a learned empty token; padded entries cannot affect the history summary. Prefix cardinality and chronology are not substitutes for remaining budget b_t or requested horizon h , which remain separate tokens. Because H0 observes poses rather than selected surfaces or free-space evidence, it inherits the $s_t^{\text{S0-pose}}$ sufficiency limitation identified above.

Pose order and observation content answer different questions. H1 tests whether chronology within the same pose-only state matters. A signed first view-direction moment together with $\mathbf{M}^{\text{dir}}$ would instead test whether approach-direction coverage is lost by generic history pooling. Such a directional descriptor is scientifically meaningful, but it remains a contingent idea until held-out errors are shown to depend on target-relative approach direction; neither H1 nor directional moments supply missing surface or free-space observations.

The current encoder implements the $s_t^{\text{S0-pose}}$ actor-state carrier described above. The implemented `v1_observed` protocol can supply its target token without adding selected-observation state; target provenance and dynamic state are separate axes. Reaching the scientific target therefore does not merely append a larger scene token: the encoder must preserve source and availability while letting each candidate query both the actor-visible target and the causal dynamic state in relative geometry. Whether points, voxels, rays, or another carrier best meets that requirement remains an empirical comparison.

The maximum horizon H_{max} binds data, model, and checkpoint. Remaining budget b_t records the factual state and h selects one return from $1 \leq h \leq b_t \leq H_{\text{max}}$. Current bundles record trained horizons rather than budget–horizon pairs, so deployed inference requests the factual diagonal $h = b_t$. A manifest-bound pair gate is required before any off-diagonal $h > 1$ query is promoted beyond the syntactic scorer interface. The public interface scores one scalar horizon at a time and preserves candidate order in its $[B, S, N_q]$ output.

4.3 Finite Action and Causal Replay

Implementation: partial **Evidence:** pending

Finite candidate tables, mesh-derived `oracle_action_mask_v1` support, and the $s_t^{\text{SO-pose}}$ replay transition are implemented and tested. This is a privileged oracle-support control. The scientific target additionally requires actor-visible action support and a causal observation update, which remain pending.

Promotion gate: retain deterministic row identity, source roles, generator and action-support identities, hard-mask isolation, and selected-transition validation; add actor-visible action support, selected-observation fusion, and no-future-observation tests

Development source: `aria_nbv/aria_nbv/targets/protocol.py`; `aria_nbv/aria_nbv/rollouts/replay/engine.py`; `aria_nbv/aria_nbv/rollouts/zarr_store.py`; `aria_nbv/aria_nbv/rollouts/qh_reader.py`; `aria_nbv/tests/rollouts/test_qh_reader.py`

At step t , the named generator protocol returns the full candidate table $\mathcal{Q}_t$, the hard mask $m_{t,i}^{\text{act}}$, and versioned failure reasons. The admissible action set is

$$\mathcal{Q}_t = \{q_{t,i}\}_{i=1}^{N_q}, \quad \mathcal{A}_t = \{i \in \{1, \dots, N_q\} : m_{t,i}^{\text{act}} = 1\}$$

Invalid rows remain auditable but cannot enter selection, Q supervision, or a bootstrap maximum. A feasible row may have negative utility; infeasibility is therefore never encoded as a small value. Padding, materialization, `valid_action_mask`, `q_train_mask`, and any predicted feasibility remain distinct. Under the dense-valid supervision profile, the writer proves and the reader revalidates equality of Q-label support and hard action support on every realized state. Unknown or legacy label-density profiles fail closed for this claim.

Candidate orientations factor a component-specific base gaze from bounded yaw–pitch perturbations. Paired proposal components may reuse a camera centre across two base-gaze families, but both remain separate actions. This makes translation and viewing direction distinguishable without collapsing their values. Generator family, resolved configuration, randomness derivation, sampled support, hard rejection, mask semantics, and stable shell identity remain attached to every row so proposal coverage can be audited before policy quality is interpreted.

4.3.1 Supervision and audit contracts

The persisted supervision profile makes label density an auditable contract rather than an inference from array shape. Historical stores declare `legacy_unspecified` together with `subset_of_action_v1`: their Q-label support may be any subset of hard-valid materialized rows and cannot establish dense one-step supervision. A fitted-Q store may instead declare `dense_valid` together with `equals_action_on_realized_steps_v1`. The writer then proves, and the reader revalidates, exact equality between `q_train_mask` and `valid_action_mask` on every realized state. Padding is excluded from both masks. Hard-invalid materialized rows remain useful binary examples for the physical feasibility head, but receive neither a fabricated Q target nor bootstrap support. Unknown or crossed profile pairs fail closed.

For benchmark reporting, **proposal support** is the persisted set of candidate rows that a declared generator attempted, with applicability and validity masks retained per family. **Proposal regret** is a descriptive comparison between the selected row and the best admissible row under the same persisted oracle contract; it is unavailable when no compatible oracle labels exist. Scene-paired aggregation first computes state-level quantities and then gives each scene equal weight, preventing scenes with more states from dominating a macro estimate. An immutable evidence bundle is the hash-bound JSON manifest and Parquet projection of these facts. Readers reject incomplete, stale, schema-mismatched, fixture, or hash-mismatched bundles; presentation layers consume the reader and do not recompute scientific quantities.

The candidate-family preflight separates two estimands. Per factual state, the root-support floor is $\max(12, \lceil 0.25N_q \rceil)$; thus 14 hard-valid rows fail and 15 pass for $N_q = 60$. Independently, the versioned diagnostic family floor requires at least one final-shell row from every applicable family in each audited state and at least three final-shell rows in total from applicable non-forward target-aware families. Forward-local selections cannot satisfy the second requirement. Inapplicable family/state cells remain explicit and do not fail the gate, whereas

missing legacy applicability is an unknown provenance state and fails closed for deployment. These rules diagnose proposal support; they do not assign low utility to invalid or absent candidates. Here selection means admission into the compact valid action shell, not the single action later chosen by the rollout policy.

Reward variation is label-conditional. When finite target-root-gain labels exist under one manifest-bound oracle contract, the preflight compares their observed range with a versioned tolerance and reports the exact label-support denominator. A Phase-A audit has no such reward labels, so `flat_gain` is unavailable with both its label and eligible-state denominators reported as zero rather than being inferred from geometric dispersion. It is specifically a no-render, no-reward-label proposal-support audit with privileged ground-truth target instruction and mesh validity, not an oracle-free audit. A passing Phase-A family gate is necessary but not sufficient for broad rollout generation; the final hash-bound pre-scale decision remains a later issue-120 gate.

All persisted Phase-A geometry uses the proposal-support normalization

$$d_{t,e}^{\text{current}} = \|\mathbf{p}_e^w - \mathbf{c}_{r_t}^w\|_2, \quad \tilde{\mathbf{c}}_{t,i}^{\text{support}} = (\mathbf{B}_{r,t}^{\text{Z-up}})^\top \frac{\mathbf{c}_{t,i}^w - \mathbf{c}_{r_t}^w}{d_{t,e}^{\text{current}}}, \quad \tilde{\mathbf{p}}_{t,e}^{\text{support}} = (\mathbf{B}_{r,t}^{\text{Z-up}})^\top \frac{\mathbf{p}_e^w - \mathbf{c}_{r_t}^w}{d_{t,e}^{\text{current}}}, \quad \|\tilde{\mathbf{p}}_{t,e}^{\text{support}}\|_2 = 1$$

. The columns of $\mathbf{B}_{r,t}^{\text{Z-up}}$ are the horizontal expansion-to-target direction, its world-Z-up left direction, and world up. Candidate centres and target-relative vectors are divided by the current three-dimensional target distance, while camera-forward directions are rotated by the same basis but remain unit vectors. Hence target-forward and target-lateral plots are invariant to factual rig yaw; rig-frame components cannot carry those axis labels directly.

The authenticated Phase-A outcome belongs to Results rather than to this definition of the audit. Method fixes the estimand, denominators, coordinate normalization, and failure semantics; Chapter 6 reports what the frozen evidence actually showed.

4.3.2 Current replay transition

After selecting an admitted row, the replay engine applies

$$(x_{t+1}, \mathbf{H}_{t+1}^{\text{pose}}, b_{t+1}, \mathcal{Q}_{t+1}) = \text{Step}(x_t, \mathbf{H}_t^{\text{pose}}, b_t, q_{t,a_t}, \xi_t)$$

where x_t is the current reference pose, $\mathbf{H}_t^{\text{pose}}$ the selected-pose history, b_t the remaining budget, and ξ_t the deterministic generation context. The next candidate table is regenerated around the selected pose under the same target task and versioned generator. Proposal and action-selection randomness use separate streams keyed by the factual selected-action history, so reordering retained trajectories cannot change a previously defined state.

This is the complete transition of the selected $s_t^{\text{S0-pose}}$ method under the current `oracle_action_mask_v1` protocol. It changes reference pose, selected-pose history, remaining budget, and finite action support. It does not claim that the actor received RGB, fused depth, or updated a spatial reconstruction, nor that the mesh-derived action mask is available to a deployable actor. Selected mesh-rendered depth may be persisted as privileged counterfactual evidence for a separate control, but unselected candidate renders never enter a successor state.

4.3.3 Scientific target transition

The CF0 state retains the same factual action, budget, and finite- candidate interface, but it replaces privileged action support with an actor-observed mask or a separately calibrated learned-feasibility route. Its dynamic state must also update from the observation acquired at the selected pose. That update is strictly causal: it may use the previous actor-visible state, the selected action, and the selected observation, but neither a future observation nor any unselected candidate render. Its carrier must preserve observed surface, observed free, unknown, finite support, uncertainty, source, and recency so that two pose-identical histories with different observations need not collapse to one state.

This target transition is a scientific requirement, not the behavior of the current replay engine. Promotion therefore requires deterministic fusion, source-dropout and no-future-observation tests, target-source leakage checks, and held-out evidence that the added state improves the target-specific value task. Persisted

mesh depth remains a privileged control until an actor-visible observation path satisfies those gates.

The normalized replay view expands one retained chain into its realized decision states. Each state sees only the prefix available before its action. Several states from one chain may share a batch, but that tensor colocation does not create temporal communication. The factual selected action supplies the only successor link used by finite-horizon learning; no counterfactual transition is fabricated for unselected rows.

Training, validation, and test populations used for one claim must share the target, actor-state, generator-family and resolved-configuration, randomness-derivation, action-mask, reward, horizon, and label-support protocols. Different random seeds may realize new candidate tables under that frozen stochastic generator; changing the generator protocol or admissibility semantics instead changes the supported decision problem. Such a comparison is reported as an explicit support-transfer experiment, not absorbed into an ordinary split. Behavior-policy changes likewise remain identified because they alter replay coverage even when the greedy target continuation is unchanged. Acquisition remains scene-disjoint and independent of oracle outcomes or model error.

4.4 Selected Interaction and Acceptance Conditions

Implementation: implemented **Evidence:** pending

A1 candidate-to-state cross-attention is the selected interaction; A0 is its feature-matched independent-row control. Both pass the same executable acceptance tests. Their scientific comparison remains pending.

Literature: [Bro+21, Zah+17]

Promotion gate: retain row-equivariance, invalid-row isolation, frame, source, mask-independence, and horizon tests; measure the A0/A1 control

Development source: `aria_nbv/aria_nbv/vin/models/target_finite_horizon.py`; `aria_nbv/aria_nbv/vin/modules/qh_state_fusion.py`; `aria_nbv/tests/vin/test_qh_state_fusion.py`; `aria_nbv/tests/vin/test_target_finite_horizon.py`

Each materialized candidate is an independent query over shared scene, target, history, budget, and horizon tokens. A1 uses the candidate as query and those five state tokens as keys and values; candidates never attend to other candidates. A0

flattens the same ordered tokens and applies a row-shared MLP. Both expose the same-width context to the same value-decoder seam.

The A0/A1 choice is orthogonal to the state distinction in Table 8. A1 is the current interaction choice, not the scientific target itself: the same row-equivariant interface can read a richer causal state without introducing candidate-to-candidate communication. Architecture should escalate only if a frozen comparison shows that an admitted state is informative but A0/A1 cannot recover its value.

4.4.1 Architectural feature-integration ladder

The ladder below is a development plan, not a family of co-equal thesis methods. It orders feature integration by the smallest additional dependency introduced after the current A0/A1 controls. “Implemented” means that the tensor path and acceptance properties exist; it does not mean that the feature has demonstrated scientific value. “Planned” names the next frozen one-factor comparison. “Possible extension” preserves a conditional test whose promotion still depends on a diagnosed failure.

Stage and status	Integrated feature	Scientific purpose	Principal risk and promotion gate
F0 — implemented	candidate-relative geometric query	Complete candidate pose relative to root and current camera, plus target pose relative to the candidate.	Already present; retain transform-direction and local-frame tests. It supplies geometry, not selected-observation state.
F1 — implemented	A0/A1 state fusion	Compare identical-input independent-row fusion with candidate-to-state cross-attention.	The controls are not parameter matched; report parameters and runtime and require a frozen held-out comparison.
F2 — planned	candidate-relative relation embeddings	Embed each candidate’s relative transform to the target and pose-bearing causal history or state elements before A1 reads them.	May duplicate information already contained in F0. Freeze one relation map, causal and padding masks, and a one-factor A1 comparison before promotion.
F3 — possible extension	permutation-invariant candidate-set summary [Zah+17]	Test whether the sampled admissible set supplies context beyond the physical state and query-local relations.	Changes conditioning from one physical action to the sampled support; require duplicate-row and absolute-value tests.
F4 — possible extension	masked candidate self-attention [Lee+19]	Test whether pairwise candidate relations explain residual error beyond an invariant set summary.	Adds quadratic candidate interaction and stronger mask sensitivity; promote only after F3 fails under matched support.
F5 — possible extension	iterative or recurrent state reading	Test whether one candidate query needs repeated access to an already adequate causal state.	May hide missing state information behind capacity; promote only after the state, F2 relations, and simpler fusion controls pass.

Table 11: Development-only architectural feature-integration ladder. Status records implementation maturity, while promotion remains conditional on a named scientific failure and a one-factor comparison.

Candidate-relative relation embeddings are the planned next architectural addition because they preserve the current per-candidate value interface while making query-to-context geometry explicit. Query-centric models use local coordinate systems and relative positional embeddings so that the active query can read other elements through their relation to it [Zho+23]. In ARIA-NBV, the same construction can attach a shared embedding of candidate–target and candidate–history transforms to the A1 key/value tokens. Applied independently to every candidate, it preserves candidate-row equivariance and does not make an action’s value depend on unrelated candidate rows.

The case against immediate promotion is equally important. F0 already exposes complete relative poses, so F2 adds an inductive bias rather than new raw information. It can redundantly parameterize the same transform, overfit the geometry of one proposal generator, increase cost with candidate–history pairs, or encode an inappropriate symmetry if gravity, metric scale, camera direction, or target orientation is suppressed. Geometric priors can reduce the hypothesis space, but only when the chosen symmetry matches the task [Bro+21]. F2 is therefore justified by a matched failure analysis showing that A1 cannot recover relevant relations from F0, not by the existence of relation-embedding architectures in another domain.

Full SE(3)-Transformer and geometric-algebra Transformer variants are out of scope for the core study. The current local frames already remove arbitrary origin dependence while retaining gravity-aligned and metric quantities that may affect sensing. No measured symmetry failure presently warrants the added representation, implementation, and compute commitments of an exact equivariant backbone. Candidate-set interaction remains a later and separate hypothesis because it changes the conditioning context of each action rather than only how its physical relations are encoded.

Candidate order has no semantic meaning. Jointly permuting aligned rows by Π must permute predictions identically:

$$f_{\theta}(\Pi \mathbf{X}_t, \Pi \mathbf{m}_t^{\text{cand}}) = \Pi f_{\theta}(\mathbf{X}_t, \mathbf{m}_t^{\text{cand}})$$

Changing only hard-invalid row contents must not alter valid-row predictions:

$$\text{Mask}(\mathbf{X}_t, \mathbf{m}_t^{\text{cand}}) = \text{Mask}(\mathbf{X}'_t, \mathbf{m}_t^{\text{cand}}) \Rightarrow \text{Mask}(f_{\theta}(\mathbf{X}_t, \mathbf{m}_t^{\text{cand}}), \mathbf{m}_t^{\text{cand}}) = \text{Mask}(f_{\theta}(\mathbf{X}'_t, \mathbf{m}_t^{\text{cand}}), \mathbf{m}_t^{\text{cand}})$$

The scorer reads $m_{t,i}^{\text{cand}}$ only to sanitize padding. It never reads $m_{t,i}^{\text{act}}$ training owns label and bootstrap admission, and online inference owns the final selectable set. Thus changing the hard mask cannot change raw conditional values or feasibility logits. Duplicate-row and valid-count tests further prevent accidental set normalization from redefining the value of an unchanged physical candidate.

Local-frame encoding removes arbitrary global origin conventions without claiming exact SE(3) equivariance. Complete root/current-relative candidate poses and the candidate-from-target pose retain translation and rotation relationships represented by the shared PoseTW encoder [Bro+21]. Provenance tests additionally exclude target gains, meshes, associations, and current candidate renders from the actor graph. Previously selected privileged depth belongs only to an explicit non-deployable state protocol.

Finally, $h = 1$ has no bootstrap, $h > b_t$ is rejected rather than clamped, and an $h > 1$ target may depend only on a factual successor queried at $h - 1$. No monotonicity across horizons is assumed because admissible actions may have negative immediate gain.

4.5 Target-Conditioned Finite-Horizon Value

Implementation: partial **Evidence:** pending

The selected method is A1-H0- $s_t^{\text{S0-pose}}$ with a scalar requested horizon, direct continuous Huber regression, fitted-Q recursion, and hard-masked online selection. The implementation path is tested, but learned exact-Q2 accuracy and policy evidence remain absent.

Literature: [De +20, CMR19, EGW05, HGS15, Sch+15, SB18]

Promotion gate: populate a frozen held-out exact-Q2 receipt; establish oracle headroom; evaluate endpoint recovery on untouched scenes

Development source: `aria_nbv/aria_nbv/vin/models/target_finite_horizon.py`; `aria_nbv/aria_nbv/lightning/`
`qh_module.py`; `aria_nbv/aria_nbv/lightning/qh_q2_certification.py`; `aria_nbv/aria_nbv/oracle/pipelines/online_qh.py`;
`aria_nbv/tests/lightning/test_qh_q2_certification.py`

4.5.1 Conditional scorer

The actor input contains root semidense and supported EFM3D summaries, target geometry admitted by the declared source protocol, relative candidate geometry, the factual selected-pose prefix, remaining budget, and materialization support. Oracle gains, target crops, meshes, associations, and audit lineage are excluded. The current `root_moments_v1` scene summary and H0 pose history make this the executable S0-pose, denoted $s_t^{\text{S0-pose}}$, not a task-sufficient reconstruction state.

Before the hard action mask is applied, the scorer emits

$$\left(Q_{h,\theta,e,i}^{\text{cond}}, \ell_{t,i}^{\text{feas}}\right) = f_{\theta}(s_t, e, q_{t,i}, h), \quad 1 \leq h \leq b_t \leq H_{\max}, \quad h = b_t \text{ when omitted}$$

for every materialized row. `conditional_q` alone carries return units and enters regression, Bellman backup, and online ranking. Feasibility logits are a separate auxiliary prediction; multiplying them by value would mix meanings. The training adapter owns hard-valid label and bootstrap admission, while online inference owns the final masked selection.

The selected A1 interaction lets each $q_{t,i}$ query shared scene, target, history, budget, and horizon tokens. A0 supplies the matched independent-row control. Both use the same geometric query and value decoder. Their comparison therefore asks whether query-dependent state reading helps under a fixed state and objective; it does not isolate attention independently of parameter count or runtime.

4.5.2 Scalar requested horizon

Using the finite-horizon return and conditional value defined in Section 2.4, the return for target e over exactly the requested residual horizon h is

$$G_{t,e}^{(h)} = \sum_{k=0}^{h-1} \gamma^k r_{t+k}^e, \quad 1 \leq h \leq b_t \leq H_{\max}$$

and the learned value is

$$Q_{h,e}^*(s_t, i) = \sup_{\pi \in \Pi^{\text{act}}} \mathbb{E}_{\pi} \left[G_{t,e}^{(h)} \mid s_t, a_t = i \right], \quad i \in \mathcal{A}_t, \quad 1 \leq h \leq b_t \leq H_{\max}, \quad Q_{0,e}^*(s, i) = 0$$

One model call follows the frozen interface

$$f_{\theta} \left(s_t^{\text{S0-pose}}, \mathbf{r}_{r \leftarrow e}, \mathbf{a}_e, \{q_{t,i}\}_{i=1}^{N_q}, h \right) \rightarrow \left(\{Q_{h,\theta,e,i}^{\text{cond}}\}_{i=1}^{N_q}, \{\ell_{t,i}^{\text{feas}}\}_{i=1}^{N_q} \right)$$

The notation Q_H Q_H names the bounded family rather than a collection of separately configured models. Remaining budget b_t describes the factual state; h selects one return from the triangular domain $\{(b_t, h) : 1 \leq h \leq b_t \leq H_{\max}\}$. Omitting h requests the full factual residual budget. The mathematical boundary $Q_0 = 0$ closes recursion, whereas zero in executable tensors denotes padding and is never a learned query.

Dense labels train $h = 1$ for every hard-valid candidate. For $h > 1$, only the factual selected action has a successor observation, so recursive supervision is necessarily narrower. The current recursive construction and exact- Q_2 surface populate the factual diagonal $h = b_t$, with exact Q_2 executable at $(b_t, h) = (2, 2)$ but its held-out receipt still pending. Dense Q_1 additionally supports $(b_t, 1)$ across realized budgets. Current bundles record trained horizons rather than realized pairs, while deployed online inference requests the diagonal and applies that horizon gate. Before any off-diagonal $h > 1$ query is exposed, the bundle and runtime must add pair-bound training and evaluation evidence and reject unsupported pairs. A wide syntactic interface is not evidence of wide learned capability.

4.5.3 Direct continuous objective

The decoder maps each candidate representation directly to the continuous, root-normalized finite-horizon return. This regression target preserves metric order and additive return units. The implemented Huber penalty is quadratic for small residuals and linear beyond its fixed threshold. Losses are first averaged within realized state and then within horizon. Consequently, a state with more valid candidates does not receive more weight merely because it contributes more rows, and abundant one-step states do not silently dominate a sparser recursive horizon. Reported calibration and support remain stratified by horizon.

At $h = 1$, evaluation can use dense candidate labels to report continuous error, within-state pairwise ranking, and greedy regret. Factual $h > 1$ transitions do not supply counterfactual values for unselected rows, so longer-horizon ranking and regret remain undefined until exact or independently oracle-rescored candidate tables exist. Recording zero support is more faithful than inventing a metric from the selected transition alone.

4.5.4 Decoder and estimator design space

Direct regression is selected because the Bellman target, endpoint gain, and reported calibration share continuous return units. That choice does not make the other formulations conceptually irrelevant. Each changes a distinct part of the inference problem and therefore remains visible with an explicit role:

Alternative	Scientific role	Quantity estimated	Condition for a meaningful comparison
Direct Huber regression	selected	continuous finite-horizon return in root-gain units	frozen per-horizon support, calibration, ranking, and endpoint evaluation
CORAL decoder	implemented matched decoder control	the same scalar fitted-Q target after ordinal discretization and fixed representative decoding	training-only support provenance and saturation diagnostics; hard-invalid rows remain excluded
Separate or fixed-horizon heads	contingent parameter-sharing control	the same horizon-indexed values without the shared scalar-horizon interface	compare only after each horizon has sufficient matched support
Double-Q backup	implemented selected estimator	greedy finite-support continuation with separated selection and delayed evaluation	report support and max-bias diagnostics; it cannot repair state aliasing
Behavior-return regression	distinct-estimand control	return under the chain-generating behavior policy	must not be called greedy value unless behavior and target continuation coincide
One-step base + residual	contingent decomposition	immediate gain plus learned continuation residual	use additive units and no candidate-table mean centering
Quantile value	exploratory idea	a return distribution rather than one scalar expectation	justify with measured stochasticity or aliasing and freeze projection, decoding, and calibration

Table 12: Value-model alternatives by scientific role and estimand. Executability, selection, and evidence are distinct; alternatives are retained without presenting them as co-equal validated methods.

Q-CORAL is an executable decoder over the same continuous fitted targets, not a different action-value definition. It preserves ordinal order but requires an additional support artifact and representative-value rule to return to continuous units [CMR19]. Behavior-return regression changes the estimand more fundamentally: without off-policy correction it follows the continuation policy that generated the retained chain rather than the greedy continuation used by fitted Q [SB18, Secs. 5.2 and 5.5, pp. 96–97 and 103–109]. The one-step residual and quantile formulations remain hypotheses because neither the need for the decomposition nor a return distribution has yet been demonstrated.

4.5.5 Fitted-Q recursion

Batch fitted Q iteration turns the fixed replay collection into successive supervised problems [EGW05]. The immediate reward is

$$r_t^e = \frac{\Delta_t^e - \Delta_{t+1}^e}{\max(\Delta_0^e, \varepsilon)}$$

and for $h > 1$ the selected method uses the lower-horizon target

$$y_t^{(h,e)} = r_t^e + \gamma B_t^{(h,e)}$$

where the continuation value is detached, frozen, or supplied by a delayed copy. The structural rule is $Q_h \leftarrow Q_{h-1}$: a horizon never bootstraps from itself. The successor maximum intersects the factual horizon support with the hard action set,

$$\mathcal{A}_{t+1}^{(Q,h-1)} = \{i : m_{t+1,i}^{\text{act}} = 1 \wedge m_{t+1,i}^{Q,h-1} = 1\}$$

and Double Q separates row selection from delayed evaluation [HGS15]. This estimator may limit max bias, but it cannot repair unsupported actions, aliased state, or missing selected observations.

The recursion estimates greedy continuation over the generated, hard-valid finite support under the named target-source, actor-state, generator, action-mask, reward, discount, and horizon protocols. Unlike an immediate per-row target, the $h > 1$ continuation value changes when the successor-support protocol changes because that protocol changes the bootstrap maximum. It is not the Monte Carlo return of the behavior policy that produced the retained chain, unless that behavior policy selects the same continuation; nor is it an optimum over ungenerated continuous camera poses. This distinction fixes the meaning of a successful fit before any policy comparison is attempted.

4.5.6 Why exact horizon two is decisive

When the selected action has a successor table with dense one-step labels, its finite-support two-step target is computable exactly:

$$y_t^{(2,\text{exact})} = r_t^e + \gamma_t \max_{j: m_{t+1,j}^{\text{train}}=1} r_{t+1,j}^e$$

This makes horizon two the first non-myopic falsification test. A unit test can inject exact Q_1 values and prove that the tensor path implements the equation. A frozen population test instead measures learned recursion error,

$$\varepsilon_t^{(2)} = \left| y_t^{(2,\text{recursive})} - y_t^{(2,\text{exact})} \right|, \quad \varepsilon_t^{(2)} \leq \tau_{\text{abs}} + \tau_{\text{rel}} \left| y_t^{(2,\text{exact})} \right|$$

on held-out supported rows. The latter tests the learned Q_1 approximation, state encoding, masks, successor linkage, and recursive target construction together. It is therefore model evidence rather than another implementation test.

The held-out evaluation begins with the complete eligible chain population rather than only rows on which an exact target is available. Chains that terminate or lose support before $h = 2$ remain in the denominator with zero exact rows. Scene- and target-stratified coverage, error, uncertainty, numeric tolerance, candidate width, generator, recipe, and behavior policy are frozen before inspection. Low error on a small supported subset cannot justify a longer horizon when minimum support or coverage fails.

Exact Q_2 is necessary but not sufficient: policy claims additionally require positive equal-budget oracle headroom and held-out endpoint recovery.

The resulting evidence boundary is explicit. The current implementation supports the $s_t^{\text{S0-pose}}$ scorer, direct objective, `oracle_action_mask_v1` recursion, and exact- Q_2 evaluation path. The scientific target additionally requires an evaluated actor-visible `v1_observed` corpus, a causal observation-updated state, actor-visible or calibrated action support, a frozen held-out exact- Q_2 receipt, positive oracle headroom, and endpoint recovery. The `v1_observed` admission and writer–reader path is implemented, but that intermediate has not passed these evidence gates. Until they pass, this chapter establishes an executable method and its falsification tests, not a successful non-myopic actor. The alternatives in Table 12 remain scientifically available, but promotion requires a diagnosed failure and a one-factor comparison rather than a larger uninterpretable model bundle.

5 Experimental Design

The experiment is a graph of admissibility gates rather than one aggregate model score. Measurement and population/action support are prerequisites for policy comparison. After support is established, the actor-protocol audit is reported independently of oracle headroom; actor-visible one-step value and exact two-step recursion follow the protocol branch, while meaningful headroom follows the oracle branch. Both branches are required before fixed-budget endpoint gap closure is admitted. Each stage has its own population, estimate, uncertainty, and stopping rule.




Figure 7: Evidence-gated claim graph. A failed stage retains its own measured outcome but blocks claims that depend on it; missing evidence remains unavailable rather than becoming zero.

5.1 Population, Estimands, and Gate Order

Validation TODO: Preregister the eligible population, exclusions, scene aggregation, independent-run structure, uncertainty interval, meaningful headroom, learned-control gap-closure fraction, and comparison family before inspecting confirmatory outcomes.

Source: experiment manifest and analysis specification

Gate: immutable analysis plan plus matched held-out policy table

5.1.1 Candidate-Support Estimands and Aggregation

RQ4 treats the scene as the independent unit. Candidate-support quality-control observations are computed per state, averaged across states within scene, and finally macro-averaged across the scene cohort; no candidate row is treated as an independent scene replicate. These within-profile diagnostics remain descriptive. Any comparative claim instead uses a paired per-scene difference over eligible scene/target/root tasks observed under both profiles, equal candidate budgets, and the frozen candidate-generation manifest. Its report identity is `paired_scene_mean_difference`, with the paired scene count as its denominator and an interval computed over the paired scene differences.

Missing or failed states are retained in the audit tables. Support counts and fractions use the preregistered failure values stated in Table 4; an undefined target-relative direction prevents that paired scene from entering a geometric contrast rather than creating an angle by imputation. The excluded-scene count and reason remain reported. Thus the descriptive state-scene macro summaries and the paired comparative estimand have distinct populations and cannot be substituted for one another.

Minimum operational support is reported as a worst-case diagnostic; the fifth percentile of per-state valid support is the more stable lower-tail estimand, accompanied by actor-valid fraction, configured-family zero rate, target-side balance, and circular target-relative orbit span. Failed roots and zero-valid configured family/state pairs remain part of the denominator. The projection fraction, oracle opportunity, and jitter compliance are secondary diagnostics: projection is calibrated framing rather than visibility, oracle opportunity is finite-support

headroom rather than policy performance, and jitter compliance must preserve the nonzero production seminar-jitter invariant.

The source population comprises ASE/ATEK snippet windows whose configured GT mesh and object-box table resolve. A frozen manifest assigns entire scenes to train, validation, or test before model selection; snippets from one scene cannot cross those boundaries. Every report records the manifest hash, scenes, snippets, admitted target tasks, rollout chains, realized transitions, candidate rows, exclusions, and failure strata. A capped training pilot may test throughput or support, but it cannot estimate held-out policy performance.

The current task generator samples geometry-valid ground-truth boxes. It can therefore establish oracle-task coverage, not actor-visible target discovery. Ground-truth target geometry may define task construction, labels, and bounded oracle references, but an actor-facing claim additionally requires an observation-derived descriptor and an audit of its matching and failure population. The same boundary excludes unselected candidate renders and oracle-derived labels from decision-time input.

The evidence graph in Figure 7 answers the research questions through seven stages. After population and action support are established, actor-protocol validity and oracle headroom form independently reportable branches; they meet again only when endpoint recovery is interpreted:

1. **Measurement validity (RQ1):** freeze crop, render, fusion, and point-mesh metric identity; show repeatability within a declared tolerance.
2. **Population and action support (part of RQ4):** establish scene-disjoint target-task coverage, candidate-family survival, hard validity, and acquisition feasibility with exact denominators.
3. **Actor-visible protocol (RQ3):** bind target-matching failures to the attempted matching population and audit the complete actor-input identity and actor-oracle leakage over the held-out scene population, without using policy headroom.

4. **Oracle headroom (first half of RQ2):** compare bounded lookahead with one-step oracle greedy under the same acquisition budget,

$$\Delta_{\text{look}} = J_e^{(H)}(\pi_{\text{oracle-look}}) - J_e^{(H)}(\pi_{\text{oracle-1}})$$

5. **Actor-visible Q_1 (RQ3 and RQ4):** under the admitted protocol, evaluate target-conditioned one-step ranking and calibration with scene-clustered uncertainty, plus dense-label replay coverage.
6. **Exact Q_2 (RQ2 and RQ4):** measure held-out two-step error, factual-successor coverage, and horizon support against the finite-support target under one receipt-bound independent-unit contract.
7. **Endpoint gap closure (second half of RQ2):** only after meaningful headroom and admitted learned-value prerequisites, estimate the prespecified learned-myopic-to-oracle-lookahead gap-closure fraction

$$\eta_Q = \frac{J_e^{(H)}(\pi_Q) - J_e^{(H)}(\pi_{\text{learned-1}})}{J_e^{(H)}(\pi_{\text{oracle-look}}) - J_e^{(H)}(\pi_{\text{learned-1}}) + \varepsilon}$$

from matched endpoint oracle evaluation. Its denominator is distinct from oracle-lookahead headroom above oracle greedy.

RQ5 and RQ6 are evaluated only if the offline finite-candidate evidence justifies extending the action or interaction setting. This dependency graph prevents an attractive downstream policy estimate from compensating for an unstable metric, an unsupported action set, privileged actor input, or failed recursion.

5.2 Replay Admission for the Learning Gates

Implementation: partial **Evidence:** pending

Dense one-step supervision, selected-transition fitted-Q recursion, scalar horizon support, and exact-Q2 certification are executable. No qualifying held-out learning or policy evidence is currently available.

Literature: [De +20, EGW05, HGS15]

Promotion gate: held-out dense-Q1 checkpoint; qualifying exact-Q2 receipt; matched endpoint oracle evaluation

Development source: `aria_nbv/aria_nbv/rollouts/gh_reader.py`; `aria_nbv/aria_nbv/lightning/gh_module.py`; `aria_nbv/aria_nbv/lightning/gh_q2_certification.py`

Replay eligibility is defined by factual evidence, not padded tensor shape. The hard action mask identifies selectable rows. The stricter one-step training mask also requires a finite target-root-gain label. A recursive row additionally requires the factual selected action, its reward and terminal role, and—when nonterminal—a reproducible successor state with lower-horizon support. Modality presence, source role, target validity, action validity, Q-label support, transition support, and horizon support remain separate masks.

Learning gate	Estimand	Minimum factual evidence
actor-visible Q_1	immediate root-normalized target gain	hard-valid candidate, finite one-step label, actor-visible target protocol, and held-out scene role
exact Q_2	first recursive finite-support target	selected reward plus terminal outcome or a successor table with complete one-step labels over every hard-valid row
recursive Q_h	greedy value within represented support	factual selected transition, terminal and discount; nonterminal rows also require successor state, hard mask, and admitted Q_{h-1} support

Table 13: Learning-target admission. Dense one-step labels, exact two-step targets, and general recursive targets are different evidence populations.

The exact- Q_2 completeness rule is strict. One finite successor label is insufficient because an unlabelled hard-valid action could hold the maximum. The certification denominator therefore begins with every eligible held-out chain, retains chains that terminate or lose support, and reports how many reach a complete successor table and a factual exact row. Missing support fails the gate; it is not zero recursion error.

One retained rollout contributes several causal decision states. State t receives only the prefix observed before a_t , even if later states share the same batch. Dense Q_1 loss uses all admitted candidate rows; $h > 1$ loss uses only factual selected transitions. Candidate losses are averaged within state, then states within horizon, then non-empty horizon means uniformly. This keeps large action tables and abundant one-step rows from hiding recursive failure.

The learner implements fitted-Q recursion from Q_h to a detached or delayed Q_{h-1} target and uses the hard-valid successor set for every maximum [De +20, EGW05]. Double-Q selection and evaluation are the implemented estimator; they

address max bias but do not create transition support [HGS15]. A single-estimator frozen lower-horizon maximum remains a planned matched control until it has an executable configuration and receipt. The experiment therefore reports horizon-specific population, calibration, terminal fraction, bootstrap support, and exact- Q_2 coverage before any endpoint comparison.

The promotion boundary is deliberately narrow. Executable training, a valid checkpoint bundle, or low error on a few rows cannot establish finite-horizon value. Actor-visible Q_1 must first pass on held-out scenes; exact Q_2 must then pass its frozen independent-unit, coverage, and tolerance rules; endpoint recovery is interpreted only if oracle headroom is already meaningful.

5.3 Matched Policies and Failure Attribution

Implementation: partial **Evidence:** pending

Scene-paired aggregation and reporting that preserves missingness are implemented; the confirmatory policy population is absent.

Promotion gate: frozen held-out manifest, completed paired endpoints, and validated confirmatory report bundle

Development source: `aria_nbv/aria_nbv/rollouts/reporting.py`; `docs/typst/thesis/experiment_data.typ`

The scene is the independent experimental unit. Comparisons are paired by scene, target task, root state, candidate distribution, hard-validity regime, and acquisition horizon. Repeated snippets, targets, and seeds are aggregated within scene before inference so dense sampling cannot masquerade as independent evidence. Policies receive the same number of acquisitions even when their planner depth differs.

Policy	Decision information	Budget	Scientific role
π_{rand}	actor-visible	H	valid-action lower reference
$\pi_{\text{learned-1}}$	actor-visible	H	learned myopic control
$\pi_{\text{oracle-1}}$	oracle immediate gain	H	one-step oracle-greedy comparator
$\pi_{\text{oracle-look}}$	bounded oracle lookahead	H	privileged bounded-lookahead reference
π_Q	actor-visible	H	learned finite-horizon gap-closure policy

Table 14: Matched policy roles. Oracle access defines a privileged reference, not a deployable upper bound; every row retains the same acquisition budget.

The primary estimand is the paired per-scene difference in fixed-budget endpoint target-quality gain. Crop, mesh, rendering, backprojection, fusion, point cap, and point-mesh metric identity are shared across policies. The analysis manifest freezes exclusions, within-scene aggregation, interval procedure, comparison family, meaningful headroom, and learned-control gap-closure threshold before learned-policy outcomes are inspected. The gap-closure ratio is reported only after oracle headroom passes the meaningful-effect gate, but its denominator remains the actor-visible learned-myopic-to-oracle-lookahead endpoint gap rather than oracle headroom above oracle greedy. That denominator must exceed the predeclared positive tolerance τ_{gap} ; otherwise the analysis reports the raw paired policy effects and marks learned gap closure as not estimable.

First failed gate	Admitted interpretation	Not implied	Next discriminating evidence
measurement	the oracle outcome is not stable enough for comparison	anything about planning, learning, or support	repair and repeat the frozen metric protocol
population / action support	the requested estimand lacks an adequate study or action population	zero utility or policy failure	report exclusions, family survival, horizon coverage, and resource failures
actor protocol	target matching, actor-input identity, or the leakage boundary failed	absence of task signal or model inadequacy	repair and repeat the failed protocol audit
oracle headroom	no meaningful non-myopic structure was detected in the frozen setup	universal myopia or model inadequacy	change support or horizon only in a separately declared study
actor-visible Q_1	the evaluated learner does not recover immediate target value under the admitted protocol	absence of signal or a specifically long-horizon failure	separate model class, optimization, capacity, sample size, calibration, and state support
exact Q_2	the first learned recursion is unsupported or inaccurate	endpoint planning value or a need for more architecture	separate coverage, Q_1 error, successor linkage, and bootstrap error
endpoint gap closure	the admitted learned policy does not close the prespecified endpoint gap	which mechanism failed	stratify by support, target observability, replay coverage, and state aliasing

Table 15: Failure-attribution matrix. Interpretation stops at the earliest failed prerequisite for each claim; downstream architecture stories remain hypotheses.

Validation TODO: Populate all seven evidence stages. Missing or failed prerequisites block downstream admission while their own measured outcomes remain reportable.

Source: confirmatory report bundle and exact-Q2 receipt

Gate: artifact-backed Results chapter

Open decision: Freeze aggregation, interval level, comparison family, meaningful headroom, and learned-control gap-closure threshold in the resolved analysis manifest.

Source: confirmatory analysis plan

Gate: analysis freeze before outcome inspection

Train-only pilots and failed generation attempts remain feasibility evidence. Renderer out-of-memory events can motivate batching and resource measurement, but they cannot support or refute candidate utility, oracle headroom, or policy quality. Likewise, a policy that cannot continue remains in the paired analysis with no further gain rather than disappearing from the denominator.

6 Results

Validation TODO: Populate one result row per inferential stage from a confirmatory bundle: population, estimate, uncertainty, admitted conclusion, and blocking condition. Every value must resolve to its raw and derived artifact provenance.

Source: confirmatory report bundle, exact-Q2 receipt, and analysis manifest

Gate: all seven evidence stages resolve without fixture or pilot substitution

The loaded report declares evidence class *pilot*. Schema validity proves that provenance and missingness are readable; it cannot promote a development fixture, training pilot, or incomplete store. Evidence presence, the gate’s own decision, and satisfaction of upstream prerequisites remain separate. A failed gate retains its estimate as a bounded negative result while preventing downstream admission.

Gate / RQ	Population	Estimate	Uncertainty	Admitted conclusion	Blocking condition
measurement / RQ1	frozen repeated oracle evaluations	repeatability: not available	declared numeric tolerance	no conclusion: required evidence is absent	mismatched identity, absent repeats, or tolerance failure
population / action / RQ4	held-out scenes, targets, and full candidate tables	support: not available	exact denominators and scene strata	no conclusion: required evidence is absent	missing population, exclusions, or valid-action support
actor protocol / RQ3	target-match attempts and held-out protocol scenes	audit: not available	match failures plus exact attempt and scene denominators	no conclusion: required evidence is absent	support failure, privileged input, or target mismatch
headroom / RQ2	paired lookahead and one-step oracle scenes	endpoint effect: not available	paired scene interval	no conclusion: required evidence is absent	measurement/support failure or non-meaningful effect
actor Q ₁ / RQ3	held-out actor-visible candidate states	ranking and calibration: not available	scene-clustered ranking and calibration intervals	no conclusion: required evidence is absent	protocol failure or inadequate ranking/calibration
exact Q ₂ / RQ2	eligible held-out factual successors	error and coverage: not available	positive independent-unit denominator and receipt-bound rule	no conclusion: required evidence is absent	failed Q ₁ , incomplete support, or recursion error
endpoint recovery / RQ2	paired held-out learned and reference policies	gap closure: not available	positive paired-scene denominator and interval	no conclusion: required evidence is absent	headroom, recursion, or recovery-rule failure

Table 16: Evidence status by inferential stage. “Failed” retains a measured negative outcome, “blocked” marks a failed prerequisite, and “not available” preserves missingness.

6.1 Measurement Validity

The loaded evidence does not establish repeatability of the frozen target-specific endpoint metric. All downstream policy quantities therefore remain unavailable.

6.2 Population and Action Support

No validated held-out bundle currently defines the study population and complete candidate-support denominators. The available Phase-A audit is a bounded proposal-support result, not a substitute for that confirmatory population.

The frozen 100-scene Phase-A proposal audit attempted 6,000 candidates and admitted 3,146 into the compact valid shells. It nevertheless failed its support gate: 44 applicable state-family cells had no selected row, 24 states missed the aggregate non-forward target-aware-family floor, and 8 states missed the root-support threshold. All 100 reviewed source rows, scenes, and target states were represented without exclusions. Because this no-render audit contains no reward labels, flat-gain status is unavailable with label and eligible-state denominators both zero. The result diagnoses proposal support; it provides no evidence about RRI, candidate quality, rollout throughput, or policy performance and does not admit broad rollout generation.

Separately, training-source rollout attempts reached mesh rendering, target-specific oracle scoring, and selected-action replay. An out-of-memory failure during unbatched candidate rendering and subsequent memory-bounded attempts identify rendering as a scale gate; they do not establish population throughput or storage feasibility.

State $\times$ family applicability and selected survival




Figure 8: Candidate-family survival for one deterministic scene from each of the ten persisted Phase-A audit strata. Each cell reports compact-valid-shell membership divided by attempted rows for one factual state and proposal family; the complete 100-state matrix remains in the hash-bound evidence bundle. The underlying artifact geometry uses correction revision `phase-a-target-aligned-z-up-v1`; the stored shells, support counts, verdict, and execution identity are unchanged. The artifact is bound to clean execution revision `2baf7cf6b276b81c50d01d45b152016d7cf68033` and SHA-256 `6d33e9e3d68737c8a6a5589ae5117c1e4d7fcaa89056fcfcaec1d315e4509c83`.

6.3 Actor-Visible Protocol

The loaded evidence does not establish the target-matching, actor-input, and leakage boundary required by RQ3.

6.4 Oracle Headroom

Meaningful equal-budget lookahead headroom over one-step oracle greedy is not estimable from the loaded evidence.

6.5 Actor-Visible One-Step Value

No validated held-out result currently supplies the complete target-matching, leakage, aggregation, uncertainty, ranking, and calibration contract required for actor-visible one-step value.

6.6 Exact Two-Step Recursion

No qualifying held-out exact-Q2 receipt is available; recursive finite-horizon accuracy is therefore unestablished.

6.7 Endpoint Recovery

The thesis cannot estimate learned endpoint recovery because the prerequisite gate chain is incomplete.

6.8 Resource Feasibility

Renderer memory failures motivate bounded rendering and retained failure provenance, but no validated completed-store evidence supports throughput or dataset-volume extrapolation.

7 Discussion

The available evidence supports an implementation conclusion, not a policy conclusion. ARIA-NBV represents target-specific finite actions, separates actor-visible inputs from oracle supervision, applies invalidity as a hard constraint, records factual selected transitions, and reports observed values without replacing missing evidence. These properties make the proposed study auditable. They do not establish metric repeatability, held-out population support, non-myopic headroom, learned value accuracy, or endpoint improvement.

The current evidence supports a narrow implementation conclusion. ARIA-NBV has an executable finite-candidate path that keeps oracle gains, meshes, associations, and candidate renders outside the scorer graph, applies invalidity as a hard decision constraint, evaluates target-specific reconstruction change offline, and exposes the resulting store through one typed reporting seam. The selected experiment still conditions on privileged `v0_gt_input` target geometry; the actor-visible `v1_observed` path is implemented but not frozen or evaluated. This establishes that the proposed experiment can be represented and audited; it does not establish deployability or that one candidate family, rollout policy, representation, or learned model performs better than another.

The evidence graph in Figure 7 makes this distinction constructive. Each claim is interpreted at its earliest failed prerequisite or prerequisites. Here, repeatability and the held-out population are not established, so headroom, actor-visible Q_1 , exact Q_2 , and endpoint recovery remain unavailable—not negative. Treating them as zeros would collapse missing evidence into measured failure and could make a pipeline limitation appear to answer the scientific question.

7.1 Interpreting future outcomes

The failure-attribution matrix in Table 15 separates mechanisms that an endpoint score alone cannot identify. Unstable measurement invalidates every downstream comparison. Adequate measurement but insufficient target or action support locates the problem in the study population rather than the policy. Stable support

with negligible oracle headroom means that the frozen candidate generator, horizon, and metric expose little exploitable non-myopic structure; it does not imply universal myopia.

The actor-protocol audit is interpretable independently of oracle headroom: it tests target matching, the complete actor-input path, and actor-oracle leakage. Conditional on that audit, actor-visible Q_1 evaluates whether the selected learner recovers immediate target value from the declared information state. Failure does not by itself show that the state lacks signal; model misspecification, optimization, capacity, and finite-sample error remain viable explanations alongside target association, representation support, and calibration. Passing Q_1 but failing exact Q_2 narrows the diagnostic to successor coverage, state aliasing, recursive targets, or bootstrap estimation. Only after those prerequisites pass does failed endpoint recovery become evidence about planning under the admitted model and replay support. Even then, the endpoint estimate alone does not identify the failed mechanism.

This logic also disciplines positive outcomes. Exact Q_2 would validate the first recursive prediction on its admitted support, not the complete policy. Endpoint recovery would remain conditional on the target protocol, finite candidate generator, hard-validity regime, horizon, oracle metric, and ASE population. The privileged bounded-lookahead policy is a reference within that finite experimental world, not a representation-independent or deployable upper bound.

7.2 Systems and external-validity boundaries

Counterfactual scoring repeatedly renders a large mesh, so memory and latency constrain branch factor and rollout volume before statistical efficiency can be assessed. An out-of-memory event motivates batching and resource measurement; it says nothing about candidate utility or planning value. Completed validated stores are required before throughput, storage, failure rates, or population coverage can be generalized.

The offline ASE setting additionally supplies geometry and target identity that a deployed egocentric actor would have to infer. The thesis therefore separates oracle task construction from actor-visible scoring end to end: target instruction, candidate proposal, hard mask, selected observation, and scorer input. Real-device or continuous-control claims would change observation, action, safety, and cost assumptions and remain outside the present evidence.

7.3 Evidence-conditioned architecture bridges

Implementation: exploratory Evidence: pending

Alternative representations and policies remain part of the thesis design registry, but they are promoted only as responses to diagnosed bottlenecks.

Promotion gate: a measured limitation that the proposed bridge directly tests

Development source: Table 8; Table 11

If the local EFM3D field is the limiting variable, broad semidense memory, sparse occupied/free/unknown state, target-centred re-lifting, or logged appearance features provide progressively stronger representation tests. If the current query contains the necessary geometry but A1 fails to recover candidate–target or candidate–history relations, the planned candidate-relative relation embedding is the smallest architectural test. If errors instead remain correlated with the sampled admissible set, DeepSets or masked candidate interaction becomes relevant. If selected-view history fails specifically by approach direction, signed first- plus second-moment memory, followed only if needed by spherical harmonics, becomes the targeted ablation. A renderable 3DGS state is appropriate only when candidate rendering or soft instance membership is itself the measured missing capability.

Continuous and hierarchical policies change the action contract rather than merely increasing model capacity. A target-then-pose policy could first select or mask a target token and then condition pose queries on target, horizon, and step; multiple targets can share scene encodings and be evaluated in parallel. This remains post-core work because continuous pose generation requires separate feasibility, safety, support, and simulator-realism evidence. The finite-candidate

model is retained as the calibrated control even if a continuous policy is later introduced.

Sequence or recurrent models are similarly conditional. They become scientifically motivated when errors persist after explicit history, time, horizon, and scene-memory conditioning, not simply because the dataset contains trajectories. Looping or recurrent refinement must preserve causal masks and candidate-row equivariance and must be compared against the same fixed-budget endpoint oracle evaluation.

Scientific-core TODO C2: After the primary evidence chain is complete, promote only bridges tied to a measured failure: EVL support, target observability, directional history, valid-set interaction, long-horizon credit, or action-support mismatch.

Rank C2 (diagnostic discriminator) — domain architecture — readiness contingent

Affected claim: which representation or interaction limitation explains residual held-out value error

Source: method design-space registry

Gate: artifact-backed failure analysis

The next evidential step remains procedural: complete validated stores under resolved manifests, freeze the eligible population and exclusions, establish oracle stability and headroom, train the matched controls, and only then interpret architectural differences. The registry preserves concrete follow-up hypotheses without presenting them as validated conclusions.

The conceptual contribution is consequently a diagnostic admissibility structure for target-specific view planning. View value is relational, invalidity is distinct from utility, privileged support is distinct from actor visibility, and each scientific claim is attached to the evidence and prerequisites required to admit it. This structure narrows alternative explanations without claiming causal identification and remains useful whether the eventual policy result is positive, negative, or blocked.

8 Conclusion

This thesis formulates target-specific egocentric next-best-view planning as a relational finite-action problem: value depends jointly on target, causal information state, admissible candidate support, horizon, update rule, and continuation policy. It contributes a target-cropped reconstruction outcome, an explicit actor-oracle information boundary, a strict separation of feasibility from utility, causal selected-transition replay, a scalar-horizon fitted-value method, and a sequential evaluation from measurement validity to endpoint recovery.

The current evidence answers the four core research questions only at the level of implemented study design. For RQ1, the target-cropped endpoint objective is specified, but confirmatory repeatability has not established it as an admissible comparison metric. For RQ2, neither meaningful equal-budget oracle headroom nor learned endpoint-gap closure is established. For RQ3, the declared actor-visible protocol and scorer are executable, but held-out target matching, input-identity and leakage audits, scene-clustered uncertainty, ranking, and calibration remain unestablished. For RQ4, candidate, replay, and validity diagnostics exist, but no validated held-out scene-target population passes the frozen support rule. These are unresolved answers, not negative empirical results.

Training-source rollouts establish that the pipeline reaches mesh rendering, target-specific oracle scoring, selected-action replay, and the fitted-value interface; renderer memory limits the evaluated scale. These observations justify an auditable method and study design, not policy superiority, a population effect, deployment readiness, or a scale estimate.

The eventual conclusion follows the prerequisite graph rather than a single undifferentiated score. Unstable measurement blocks policy comparison; inadequate support blocks population claims; a failed actor protocol blocks actor-visible interpretation; negligible headroom is a setup-specific negative result; and failed Q_1 , exact Q_2 , or endpoint recovery progressively narrows—but does not uniquely identify—the learned-control limitation. RQ5 and RQ6 remain conditional extensions: online interaction and continuous control change the evidence, action, safety,

and cost assumptions and are not missing parts of the present offline finite-candidate evaluation.

Development roadmap

This page is the compact strategic view of thesis development. It is omitted from submission output. Scientific claims remain owned by the active thesis; executable behavior and measurements remain owned by source, tests, configuration, and immutable evidence artifacts.

Implementation: partial

Evidence: pending

Reviewed 31 Aug 2026 refresh due 14 Sep 2026. Target: reproducible full-draft freeze by 30 November 2026, followed by a safety buffer until submission on 7 January 2027.

Promotion gate: Close P1 before confirmatory policy claims.

Development source: docs/typst/thesis/development/roadmap.toml; active thesis, package, test, configuration, and evidence owners

TL;DR

Done

Pilot infrastructure

Candidate-benchmark smoke evidence, a bounded two-scene orbit MVP, and a pilot rollout report are committed. These prove the pipeline can produce reviewable artifacts, not that population-scale support or policy gains are established.

Now

Evidence substrate

Freeze the multi-scene store, scene-level split, camera/depth and target-RRI fidelity, candidate-support admission, and throughput receipts while the thesis narrative and conceptual figures remain under review.

Blocked

Confirmatory claims

Lookahead-headroom, learned-value, and population-coverage claims remain blocked until paired multi-scene evidence and an immutable experiment registry exist. The current rollout bundle is explicitly pilot evidence.

Next

Paired evidence to freeze

Run equal-budget baselines, train and evaluate the finite-horizon value model under actor-visible inputs, synthesize results and limitations, and reach the reproducible full-draft freeze recorded below.

Critical path

The strategic dependency is evidence-first: close the substrate, establish paired baseline headroom, evaluate the learned value model, synthesize the claims, and freeze the full draft. The December-to-January interval is a safety buffer, not planned scientific work.




Figure 9: Critical development gates from experiment substrate to the 30 November 2026 full-draft freeze.

Milestones

ID	Phase	Dates	State	Exit gate
P0	Pilot infrastructure	7 Jul–30 Aug 2026	done	Committed smoke and MVP artifacts compile and remain explicitly bounded as pilot evidence.
P1	Freeze experiment substrate	31 Aug–30 Sep 2026	in progress	Multi-scene immutable store, scene-level split, camera/depth and target-RRI fidelity, candidate-support admission, and throughput receipts.
P2	Paired policy baselines	1 Oct–20 Oct 2026	planned	Equal-budget random-valid, greedy, and bounded-lookahead comparisons report uncertainty, support, and oracle receipts.
P3	Learned value evaluation	21 Oct–10 Nov 2026	planned	Held-out ranking, calibration, and paired policy value are reported under actor-visible inputs.
P4	Results and discussion synthesis	11 Nov–22 Nov 2026	planned	Final tables, failure cases, coverage limits, and claim dispositions resolve to the immutable evidence registry.
P5	Reproducibility and full-draft freeze	23 Nov–30 Nov 2026	planned	Full draft, reproducible configurations, final figures, and release checks pass with no unresolved placeholders.
P6	Safety buffer and submission	1 Dec–7 Jan 2027	safety buffer	Advisor feedback, formal checks, final upload, and the authenticated submission receipt are closed without reopening the scientific core.

Table 17: Evidence-gated schedule from the completed pilot infrastructure to submission.

Dependencies are checked against the adjacent roadmap data.

Blockers

- **B1 — Multi-scene experiment substrate** (affects P1): The committed artifacts do not yet establish a fresh multi-scene store, scene-level split, camera/depth fidelity, target-RRI robustness, and representative throughput in one immutable receipt chain.
- **B2 — Candidate-support scale admission** (affects P1, P2): The positive two-scene orbit MVP is explicitly not scale-up admission; family survival, framing, diversity, and proposal regret still require broader evidence.
- **B3 — Confirmatory experiment registry** (affects P2, P3, P4): The available report bundle is a one-scene pilot. Final comparisons need immutable run/config identities, coverage accounting, uncertainty, and claim-to-artifact resolution.

The primary evidence pointers remain beside each record in `roadmap.toml`; the roadmap contract checks their existence, while scientific review decides whether they are sufficient for promotion.

Promotion queue

Promotion queue — candidate: Promote bounded-lookahead and learned-value results only after paired held-out evidence is immutable.

Source: P2-P3 experiment registry and rollout reports target: Results and Discussion gate: positive headroom, actor-visible evaluation, uncertainty, and support coverage

Promotion queue — blocked: Resolve candidate-support and substrate limits before claiming population coverage.

Source: P1 evidence receipts and candidate benchmark manifests target: Experimental Design and Results gate: multi-scene support, family survival, scene-level split, and representative throughput

Promotion queue — deferred: Keep online discrete and continuous target-then-pose extensions outside the core until the offline finite-candidate gates close.

Source: Discussion and roadmap P1-P3 target: Discussion gate: stable offline learned-value evidence and an explicit post-core scope decision

Maintenance contract

`roadmap.toml` is the single owner for the snapshot, dates, states, dependencies, blockers, release checks, and evidence pointers rendered above. Update it at least every 14 days and whenever a gate changes state. `make thesis-roadmap-contract` rejects stale review dates, broken dependencies, missing local evidence

pointers, divergence from the thesis submission date, or references to the retired M1 snapshot. Internal trackers and hosted issue state may inform a refresh but are not imported as public thesis truth.

The native table, cards, and critical-path strip are deliberately sufficient: they show exact milestones and the consequential dependency chain without a second diagram data model. A date-dense Gantt can be reconsidered only if this compact projection no longer answers the planning question.

HM/FK07 release gate

These checks are human-owned and apply to the exact final candidate and the author’s authenticated records; the repository cannot satisfy them by itself.

- **R1 — Applicable rules:** Confirm the individually assigned programme and SPO version in PRIMUSS; public reading versions establish only the current baseline.
- **R2 — Registration record:** Retain the PRIMUSS confirmation, official start date, individual deadline, assigned examiners, supervisor-specific form requirements, and requested attachments.
- **R3 — Exact upload set:** Before the individual deadline, verify the exact final PDF and every PRIMUSS-marked mandatory document or attachment. A repository build, digest, or upload alone is not submission evidence.
- **R4 — Conditional paper copy:** Provide an identical paper copy only if the examiner or authenticated PRIMUSS record requires one under the current FK07 process.
- **R5 — Final-PDF review:** Check the declaration and AI/tool disclosure in the exact PDF prepared for upload; confirm any mandatory signature or attachment against the authenticated record.
- **R6 — Submission closure:** Finalize with the PRIMUSS action “Abschlussarbeit abgeben”, retain the authenticated timestamp, and archive the exact submitted PDF; only that receipt closes the gate.

Official baseline checked 31 August 2026: HM examination regulations, FK07 thesis process, PRIMUSS registration guide, PRIMUSS submission guide. At the full-draft freeze, recheck the official versions and the individually assigned instructions in authenticated PRIMUSS records.

Freeze and submission

By 30 November 2026, the full draft, experiment registry, configurations, figures, and release checks must be reproducible and free of unresolved placeholders. The interval from 1 December to 7 January 2027 is reserved for advisor feedback, formal HM/FK07 checks, final upload, and recovery from unexpected defects. Only the authenticated submission receipt closes the final gate.

9 Development Pilots

9.1 Candidate-Family Support




Figure 10: Attempted candidate centres from a two-scene, two-state real-data method audit. Each state is centred on its factual expansion pose, yaw-aligned to the oracle-evaluation target with world Z-up, and divided by its current three-dimensional target distance. The cross marks the expansion/root pose, the stars mark the two oracle task target centres, colour denotes proposal family, and marker shape denotes actor-valid or selected status. The figure diagnoses proposal geometry under the `v0_gt_input` target protocol; it is neither a population estimate nor evidence of actor-visible target discovery or policy quality.

The target-orbit challenger adds attempted support on both target-relative sides in this bounded pilot, while pruning still leaves invalid rows and one configured family/state pair without actor-valid support. The associated immutable evidence manifest records the common-frame reducer, store-manifest hashes, expected state index, committed reduced candidate rows, resolved challenger profile, reproduction commands, metric denominators, and plot hashes. No frame-mixed side-balance or orbit-span values are used. The interactive evidence and Streamlit view can optionally add short arrows for the valid candidates' camera optical axes; the static thesis figure omits them to preserve legibility.

9.2 Target-Frame $\mathcal{S}^2$ Diagnostics

The following frozen figures are generated by the same Python reporting and plotting owners as the stored-rollout application. For the selected real-data shard, the report records these support counts: source samples 1, snippets 1, scenes 1, targets 1, retained rollout chains 1, and selected steps 5. It remains pilot evidence: its purpose is to validate the representation and reporting contract, not to estimate policy quality. The two point-direction views contain 5 factual displacements and 5 factual optical axes. Colour denotes rollout-chain identity, marker shape denotes decision-step identity, and the sphere heat map is computed from the complete equal-solid-angle count grid rather than from the bounded incidence overlay.





Figure 11: Target-frame $\mathcal{S}^2$ movement directions from the frozen pilot report. This view records how the camera moved; it is not a visibility measurement.

$\hat{V}[j,i]$ — target-frame camera +Z direction · s01 · 39089f593837



Persisted acquisition step · acquisition 1 (i=0) · acquisition 2 (i=1) · acquisition 3 (i=2) · acquisition 4 (i=3) · acquisition 5 (i=4)

Figure 12: Target-frame $\mathcal{S}^2$ selected-camera optical-axis directions from the same frozen pilot report. This view records where the camera pointed; it is distinct from both camera motion and calibrated surface support.

The calibrated proxy diagnostic integrates 5 selected frusta. Their mean intrinsic field-of-view solid angle is 2.49 sr. On the geometric-mean-scale proxy sphere, one view covers a mean fraction of 0.168, while their union covers 0.312. These fractions express geometric potential support under the proxy and front-facing tests; they exclude true target shape, depth ordering, and scene occlusion.




Figure 13: Calibrated selected-frustum support on the target-centred proxy sphere. The surface heat map is the complete equal-solid-angle coverage field; overlaid incidence samples preserve rollout and time heritage. This is not measured target-mesh visibility.

Development Register: Method Alternatives

This register preserves executable but unpromoted controls and future hypotheses outside the thesis-core Method chapter. Entry here means that an idea has a named comparison and promotion gate; it does not mean that the thesis evaluates or recommends it.

Implementation: exploratory Evidence: pending

The thesis-core selection remains $A1-H0-s_t^{S0-pose}$ with direct continuous regression. Executable controls, planned candidates, and exploratory ideas below remain available but are excluded from its central method claim.

Promotion gate: promote only a one-factor comparison with an immutable receipt and untouched scene-disjoint evaluation

Development source: [aria_nbv/aria_nbv/vin/](#); [aria_nbv/aria_nbv/lightning/](#); [docs/contents/theory/](#)

State and history alternatives

- $s_t^{S1-surface}$ **selected surfaces:** the executable privileged control backprojects only factual selected mesh depth, pools the causal point set, and initializes its residual at zero so a fresh S1 scorer matches H0. It lacks free/unknown evidence, is density weighted, and has no confirmatory receipt; it therefore cannot support a geometry-benefit or deployment claim.
- s_t^{S2-ray} **ray-aware memory:** a planned state would distinguish observed surface, observed free, unknown, support, uncertainty, source, and recency. It requires deterministic fusion and no-future-observation tests before a model comparison.
- **H1 ordered history:** an executable causal Transformer adds relative age to the same selected-pose tokens used by H0. It remains a pose-only capacity ablation until repeated-seed endpoint evidence justifies chronology.
- **Other carriers:** target-centred EFM3D re-lifting, appearance-on-point, sparse TSDF/SDF encoders, object-aware 3D Gaussian splats, and SceneScript context are hypotheses for diagnosed representation failures, not parallel thesis methods.

Interaction alternatives

The canonical feature-integration order and per-stage maturity are recorded in Table 11. This register does not duplicate that ladder. Candidate-relative relation embeddings are the planned next one-factor comparison; candidate-set summaries, masked candidate interaction, and recurrent reading remain possible extensions. A stage may be promoted only after A0/A1 are measured under the same state, target protocol, decoder, objective, and evaluation contract and a concrete failure identifies the missing interaction. Candidate-set models must additionally retain row equivariance, invalid-row isolation, duplicate tests, and absolute-value semantics.

Objective and estimator alternatives

- **CORAL decoding** is executable over the same continuous fitted-Q targets, but its ordinal loss and fixed representatives add support-provenance and saturation questions. It is excluded until direct regression supplies the frozen control.
- **Fixed-horizon or separate heads** alter parameter sharing and the public interface. They are ablations against the scalar requested-horizon family, not co-equal definitions of Q_H .
- **Behavior-return Monte Carlo regression** estimates the continuation policy that generated each chain rather than greedy finite-support value unless those policies coincide. It must retain that distinct estimand.
- **One-step plus residual prediction** can test whether calibrated immediate utility is a useful base, but the residual may not be mean-centred across a candidate table because unrelated or duplicate rows would then redefine an action’s absolute target.
- **Quantile regression** is justified only after stochastic return variation or state aliasing is measured and a distributional Bellman projection, support, decoding rule, and calibration protocol are frozen.

The promotion sequence is: establish population/action support, measure oracle headroom, learn actor-visible Q_1 , pass exact Q_2 , and only then introduce the smallest alternative that addresses the observed failure.

Ranked scientific-core work

Architecture and data TODOs use one extension of the thesis draft-marker contract rather than a separate backlog. Priority is ordinal and claim-centred: C0 protects a validity or information boundary shared by the central claims; C1 blocks the next principal-RQ inference; C2 discriminates a diagnosed alternative explanation; C3 is an exploratory extension outside the core claim. Readiness (ready, blocked, or contingent) is orthogonal, so a blocked C0 remains scientifically more critical than a ready C2 even though only the latter may be immediately executable. Every marker must name its affected claim, source, promotion gate, and blocker when blocked. The policy deliberately avoids weighted scores whose apparent precision would become stale as evidence changes.

Scientific-core TODO C0: Freeze and audit the actor-visible target corpus before interpreting any learned target-conditioned result.

Rank C0 (validity boundary) — domain data — readiness ready

Affected claim: actor/oracle separation and deployability of the target-conditioned value claim

Source: `v1_observed` target protocol and frozen corpus manifest

Gate: scene-disjoint corpus, explicit matching failures, source dropout, descriptor identity, and leakage audit

Scientific-core TODO C1: Implement the smallest observation-updated state that preserves surface, free, unknown, support, uncertainty, source, and recency.

Rank C1 (core inference) — domain architecture — readiness blocked

Affected claim: whether actor-visible state retains information needed for target-specific bounded return

Blocked by: an admitted actor-visible selected-observation data path

Source: $s_t^{S2\text{-ray}}$ candidate realization and selected-observation replay contract

Gate: deterministic fusion, no-future-observation test, source masks, and matched held-out comparison against $s_t^{S0\text{-pose}}$

Scientific-core TODO C1: Produce the immutable evidence chain that separates implementation parity, learned recursion accuracy, and policy recovery.

Rank C1 (core inference) — domain data — readiness blocked

Affected claim: learned lookahead and held-out endpoint recovery

Blocked by: validated rollout stores under the selected target, state, candidate, and reward protocols

Source: exact- Q_2 census, oracle-headroom table, and paired endpoint receipt

Gate: frozen eligible population, support floor, tolerance, positive oracle headroom, and untouched scene-disjoint evaluation

Scientific-core TODO C2: Promote exactly one control that tests the diagnosed limitation; do not bundle state, interaction, and objective changes.

Rank C2 (diagnostic discriminator) — domain architecture — readiness contingent

Affected claim: alternative explanation for residual value error after state and support gates pass

Source: feature-integration ladder F2–F5, H1, CORAL, residual, and separate-head controls

Gate: one measured failure mapped to one factor-matched comparison

Scientific-core TODO C3: Retain these ideas as scoped extensions; do not treat conceptual interest as evidence for inclusion in the core method.

Rank C3 (exploratory extension) — domain architecture — readiness contingent

Affected claim: post-core generalization beyond the finite-candidate offline study

Source: 3DGS memory, SceneScript context, continuous policy, and quantile-value ideas

Gate: completed core evidence chain plus a failure that requires the changed state, action, or estimand contract

Bibliography

- [Baj88] R. Bajcsy, “Active Perception,” *Proceedings of the IEEE*, vol. 76, no. 8, pp. 966–1005, 1988, doi: 10.1109/5.5968.
- [AWB88] J. Aloimonos, I. Weiss, and A. Bandyopadhyay, “Active Vision,” *International Journal of Computer Vision*, vol. 1, no. 4, pp. 333–356, 1988, doi: 10.1007/BF00133571.
- [SRR03] W. R. Scott, G. Roth, and J.-F. Rivest, “View Planning for Automated Three-Dimensional Object Reconstruction and Inspection,” *ACM Computing Surveys*, vol. 35, no. 1, pp. 64–96, 2003, doi: 10.1145/641865.641868.
- [Jia+25] Z. Jia, Y. Li, Q. Hao, and S. Zhang, “PB-NBV: Efficient Projection-Based Next-Best-View Planning Framework for Reconstruction of Unknown Objects,” *IEEE Robotics and Automation Letters*, vol. 10, no. 7, pp. 7444–7451, 2025, doi: 10.1109/LRA.2025.3573631.
- [GML22] A. Guédon, P. Monasse, and V. Lepetit, “SCONE: Surface Coverage Optimization in Unknown Environments by Volumetric Integration,” in *Advances in Neural Information Processing Systems*, 2022. [Online]. Available: <https://arxiv.org/abs/2208.10449>
- [JLD24] W. Jiang, B. Lei, and K. Daniilidis, “FisherRF: Active View Selection and Uncertainty Quantification for Radiance Fields using Fisher Information.” [Online]. Available: <https://arxiv.org/abs/2311.17874>
- [Fra+25] N. Frahm *et al.*, “VIN-NBV: A View Introspection Network for Next-Best-View Selection.” [Online]. Available: <https://arxiv.org/abs/2505.06219>
- [Jeo+26] S. Jeong, E. Lee, J. Kim, and A. Kim, “Informative Object-centric Next Best View for Object-aware 3D Gaussian Splatting in Cluttered Scenes.” [Online]. Available: <https://arxiv.org/abs/2602.08266>
- [Che+24] X. Chen, Q. Li, T. Wang, T. Xue, and J. Pang, “GenNBV: Generalizable Next-Best-View Policy for Active 3D Reconstruc-

- tion,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 16436–16445. [Online]. Available: https://openaccess.thecvf.com/content/CVPR2024/html/Chen_GenNBV_Generalizable_Next-Best-View_Policy_for_Active_3D_Reconstruction_CVPR_2024_paper.html
- [Lu+26] C.-Y. Lu *et al.*, “Hestia: Voxel-Face-Aware Hierarchical Next-Best-View Acquisition for Efficient 3D Reconstruction,” in *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, 2026. [Online]. Available: https://openaccess.thecvf.com/content/WACV2026/papers/Lu_Hestia_Voxel-Face-Aware_Hierarchical_Next-Best-View_Acquisition_for_Efficient_3D_Reconstruction_WACV_2026_paper.pdf
- [AKC20] K. Ashutosh, S. Kumar, and S. Chaudhuri, “3D-NVS: A 3D Supervision Approach for Next View Selection.” [Online]. Available: <https://arxiv.org/abs/2012.01743>
- [Gué+23] A. Guédon, T. Monnier, P. Monasse, and V. Lepetit, “MACARONS: Mapping And Coverage Anticipation with RGB Online Self-Supervision,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023. [Online]. Available: <https://arxiv.org/abs/2303.03315>
- [Str+24] M. Strong, B. Lei, A. Swann, W. Jiang, K. Daniilidis, and M. K. III, “Next Best Sense: Guiding Vision and Touch with FisherRF for 3D Gaussian Splatting.” [Online]. Available: <https://arxiv.org/abs/2410.04680>
- [Str+24] J. Straub, D. DeTone, T. Shen, N. Yang, C. Sweeney, and R. Newcombe, “EFM3D: A Benchmark for Measuring Progress Towards 3D Egocentric Foundation Models.” [Online]. Available: <https://arxiv.org/abs/2406.10224>

- [Eng+23] J. Engel *et al.*, “Project Aria: A New Tool for Egocentric Multi-Modal AI Research.” [Online]. Available: <https://arxiv.org/abs/2308.13561>
- [Met25] Meta Platforms Inc., “Aria Synthetic Environments Dataset.” [Online]. Available: https://facebookresearch.github.io/projectaria_tools/docs/open_datasets/aria_synthetic_environments_dataset
- [SB18] R. S. Sutton and A. G. Barto, *Reinforcement Learning: An Introduction*, 2nd ed. Cambridge, MA: The MIT Press, 2018. [Online]. Available: <http://incompleteideas.net/book/the-book-2nd.html>
- [LHP23] M. Lauri, D. Hsu, and J. Pajarinen, “Partially Observable Markov Decision Processes in Robotics: A Survey,” *IEEE Transactions on Robotics*, vol. 39, no. 1, pp. 21–40, 2023, doi: 10.1109/TRO.2022.3200138.
- [De +20] K. De Asis, A. Chan, S. Pitis, R. S. Sutton, and D. Graves, “Fixed-Horizon Temporal Difference Methods for Stable Reinforcement Learning,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, 2020, pp. 3741–3748. doi: 10.1609/aaai.v34i04.5784.
- [Sch+15] T. Schaul, D. Horgan, K. Gregor, and D. Silver, “Universal Value Function Approximators,” in *Proceedings of the 32nd International Conference on Machine Learning*, in Proceedings of Machine Learning Research, vol. 37. PMLR, 2015, pp. 1312–1320. [Online]. Available: <https://proceedings.mlr.press/v37/schaul15.html>
- [FMP19] S. Fujimoto, D. Meger, and D. Precup, “Off-Policy Deep Reinforcement Learning without Exploration,” in *Proceedings of the 36th International Conference on Machine Learning*, in Proceedings of Machine Learning Research, vol. 97. PMLR, 2019, pp. 2052–2062. [Online]. Available: <https://proceedings.mlr.press/v97/fujimoto19a.html>
- [Kum+20] A. Kumar, A. Zhou, G. Tucker, and S. Levine, “Conservative Q-Learning for Offline Reinforcement Learning,” in *Advances in Neural Information Processing Systems*, 2020, pp. 1179–1191. [Online].

Available: <https://papers.nips.cc/paper/2020/hash/0d2b2061826a5df3221116a5085a6052-Abstract.html>

- [Zho+23] Z. Zhou, J. Wang, Y.-H. Li, and Y.-K. Huang, “Query-Centric Trajectory Prediction,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [Bro+21] M. M. Bronstein, J. Bruna, T. Cohen, and P. Veličković, “Geometric Deep Learning: Grids, Groups, Graphs, Geodesics, and Gauges.” [Online]. Available: <https://arxiv.org/abs/2104.13478>
- [Zah+17] M. Zaheer, S. Kottur, S. Ravanbakhsh, B. Póczos, R. Salakhutdinov, and A. J. Smola, “Deep Sets,” in *Advances in Neural Information Processing Systems*, 2017. [Online]. Available: https://papers.nips.cc/paper_files/paper/2017/hash/f22e4747da1aa27e363d86d40ff442fe-Abstract.html
- [Lee+19] J. Lee, Y. Lee, J. Kim, A. R. Kosíorek, S. Choi, and Y. W. Teh, “Set Transformer: A Framework for Attention-based Permutation-Invariant Neural Networks,” in *Proceedings of the 36th International Conference on Machine Learning*, in Proceedings of Machine Learning Research, vol. 97. PMLR, 2019, pp. 3744–3753. [Online]. Available: <https://proceedings.mlr.press/v97/lee19d.html>
- [Zha+21] H. Zhao, L. Jiang, J. Jia, P. H. S. Torr, and V. Koltun, “Point Transformer,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021. [Online]. Available: https://openaccess.thecvf.com/content/ICCV2021/papers/Zhao_Point_Transformer_ICCV_2021_paper.pdf
- [CGS19] C. Choy, J. Gwak, and S. Savarese, “4D Spatio-Temporal ConvNets: Minkowski Convolutional Neural Networks,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2019. [Online]. Available: <https://arxiv.org/abs/1904.08755>
- [SHW21] V. G. Satorras, E. Hoogeboom, and M. Welling, “E(n) Equivariant Graph Neural Networks,” in *Proceedings of the 38th International*

- Conference on Machine Learning*, in Proceedings of Machine Learning Research, vol. 139. PMLR, 2021, pp. 9323–9332. [Online]. Available: <https://proceedings.mlr.press/v139/satorras21a.html>
- [DAT20] N. De Cao, W. Aziz, and I. Titov, “The Power Spherical distribution.” [Online]. Available: <https://arxiv.org/abs/2006.04437>
- [Ave+24] A. Avetisyan *et al.*, “SceneScript: Reconstructing Scenes With An Autoregressive Structured Language Model.” [Online]. Available: <https://arxiv.org/abs/2403.13064>
- [EGW05] D. Ernst, P. Geurts, and L. Wehenkel, “Tree-Based Batch Mode Reinforcement Learning,” *Journal of Machine Learning Research*, vol. 6, pp. 503–556, 2005, [Online]. Available: <https://jmlr.org/papers/v6/ernst05a.html>
- [HGS15] H. van Hasselt, A. Guez, and D. Silver, “Deep Reinforcement Learning with Double Q-learning.” [Online]. Available: <https://arxiv.org/abs/1509.06461>
- [CMR19] W. Cao, V. Mirjalili, and S. Raschka, “Rank consistent ordinal regression for neural networks with application to age estimation.” [Online]. Available: <https://arxiv.org/abs/1901.07884>

A Reproducibility Record

The report bundle is the single numerical interface between validated rollout artifacts and this manuscript. Its schema version, evidence status, store manifests, typed parameter rows, failure records, and sidecar hashes preserve provenance without duplicating analysis logic in Typst. The document loader rejects schema or status mismatches before any table is rendered. Compact labels and digest prefixes are used below; complete store identifiers, paths, and hashes remain in the machine-readable bundle.

The implementation anchors behind that interface are the immutable `vin_offline/` source store, the task-specific `rollouts.zarr/` replay store, and the derived `q_h/` training cache. Their exact schemas, normalized joins, chunking, codecs, receipt versions, and execution commands remain manifest-owned reproducibility metadata rather than scientific entities in Chapter 3. A reported row is admissible only when the bundle resolves these identities and preserves source role, missingness, and failure records.

Evidence class	Store	Manifest digest	Validated
Development fixture	store S1	not-a-real-m...	yes

Table 18: Loaded development-fixture provenance. The fixture verifies the document interface only and is not scientific evidence.

The bundled development fixture contains synthetic values solely to test typed loading and missingness. Its numerical rows are deliberately not rendered or interpreted.

B Development Diary

Archived source note: This marked diary is a development-only record. It is excluded from submission mode; only evidence-backed definitions, protocols, and results graduate into the thesis body.
Source: thesis mode contract

B.1 Decisions Retained

The thesis keeps the privileged oracle and actor-visible policy as separate contracts, evaluates finite candidate sets under a shared validity mask, and treats target-specific reconstruction improvement as the common comparison signal. Rollout stores retain selected transitions and provenance rather than becoming a second implementation of the oracle or training pipeline.

Open decision: Freeze the scene-disjoint analysis population and the minimum evidence scale only when the resolved manifests and exclusion ledger are available.
Source: experimental-design contract
Gate: confirmatory bundle freeze

B.2 Failures That Changed the Plan

Early rollout attempts reached the mesh-rendering path but exceeded available accelerator memory when candidate views were rendered without a bounded batch. This failure narrowed the immediate claim to implementation readiness and made peak memory, throughput, failure provenance, and completed-store validation explicit scale gates.

Validation TODO: Require completed stores, matching manifests, resource summaries, and recorded failures before extrapolating rollout bandwidth or storage demand.
Source: rollout attempt artifacts
Gate: cluster-readiness evidence

B.3 Rejected Scope

Continuous actions, online simulators, additional scene representations, and alternative reinforcement-learning families are excluded from the core study while the finite-candidate target-conditioned comparison remains unvalidated. They would

change both the action contract and the supervision regime, so adding them now would weaken rather than extend the primary experiment.

Remove or rewrite before submission: Delete bridge material that cannot be tied to a concrete limitation or a matched follow-up experiment after the confirmatory analysis is complete.

Source: scope review

Gate: discussion freeze

B.4 Directional-Memory Hypothesis


Research TODO: Treat target-centred directional memory as an ablation hypothesis, not an actor feature or contribution. The representation must be implemented, mask-audited, and compared against pose distance and overlap before any predictive claim enters the manuscript.

Source: RQ3 representation hypothesis

Gate: paired held-out ablation

A. Target-centred directions on $\mathbb{S}^2$

actual logged camera centres, expressed in the target-object frame




● logged history direction ● example query

Scene 81286, sample ASE_81286_Atek_000024, target instance 128 frames 0, 2, 4, 6, 8, 10 and query frame 15.

B. Complete-domain view and prospective compression

Mollweide is equal-area; no smooth field or SH lobe is implied



$+x_e$ at map centre · $+z_e$ north · wrap seam at $\pm 180^\circ$

$$\mathbf{M}_{\text{dir}(v)} = \sum_{k < t} w_{k(v)} \mathbf{d}_{k(v)} \mathbf{d}_{k(v)}^\top$$

$$\nu_{i(v)} = 1 - \frac{\mathbf{d}_i^\top \mathbf{M}_{\text{dir}} \mathbf{d}_i}{\text{tr} \mathbf{M}_{\text{dir}} + \epsilon}$$

This fixture uses uniform weights and a logged future pose only to make the second-moment query inspectable. It does not claim that the learner currently stores $\mathbf{M}_{\text{dir}}$ or that this novelty predicts target RRI.

HYPOTHESIS / ABLATION MATERIAL. Keep outside submission claims until the representation exists and paired evaluation tests whether it adds information beyond pose distance and overlap.

Figure 14: Development-only directional-memory fixture. The sphere and Mollweide panels show the same logged ASE camera directions in a target-object frame; the orange query is another logged pose used to inspect one prospective second-moment novelty definition. No smooth field, spherical-harmonic representation, or learned benefit is claimed.

B.5 Next Evidence Gates

The next admissible evidence is a validated report bundle that fixes the study population, records target and candidate exclusions, supplies paired endpoint outcomes with uncertainty, and joins runtime and storage statistics to the same manifests. Only then can the manuscript replace the present availability statements with policy comparisons.

Implementation TODO: Export the confirmatory report bundle, compile both thesis modes against it, and inspect the resulting tables and failure cases before freezing claims.

Source: reporting contract

Gate: claim freeze

Ehrenwoertliche Erklaerung

Ich versichere, dass ich diese Masterarbeit selbstständig verfasst, noch nicht anderweitig für Prüfungszwecke vorgelegt, keine anderen als die angegebenen Quellen oder Hilfsmittel benutzt sowie wörtliche oder sinngemäße Zitate als solche gekennzeichnet habe.

Muenchen, 7. Jan 2027

Jan Duchscherer